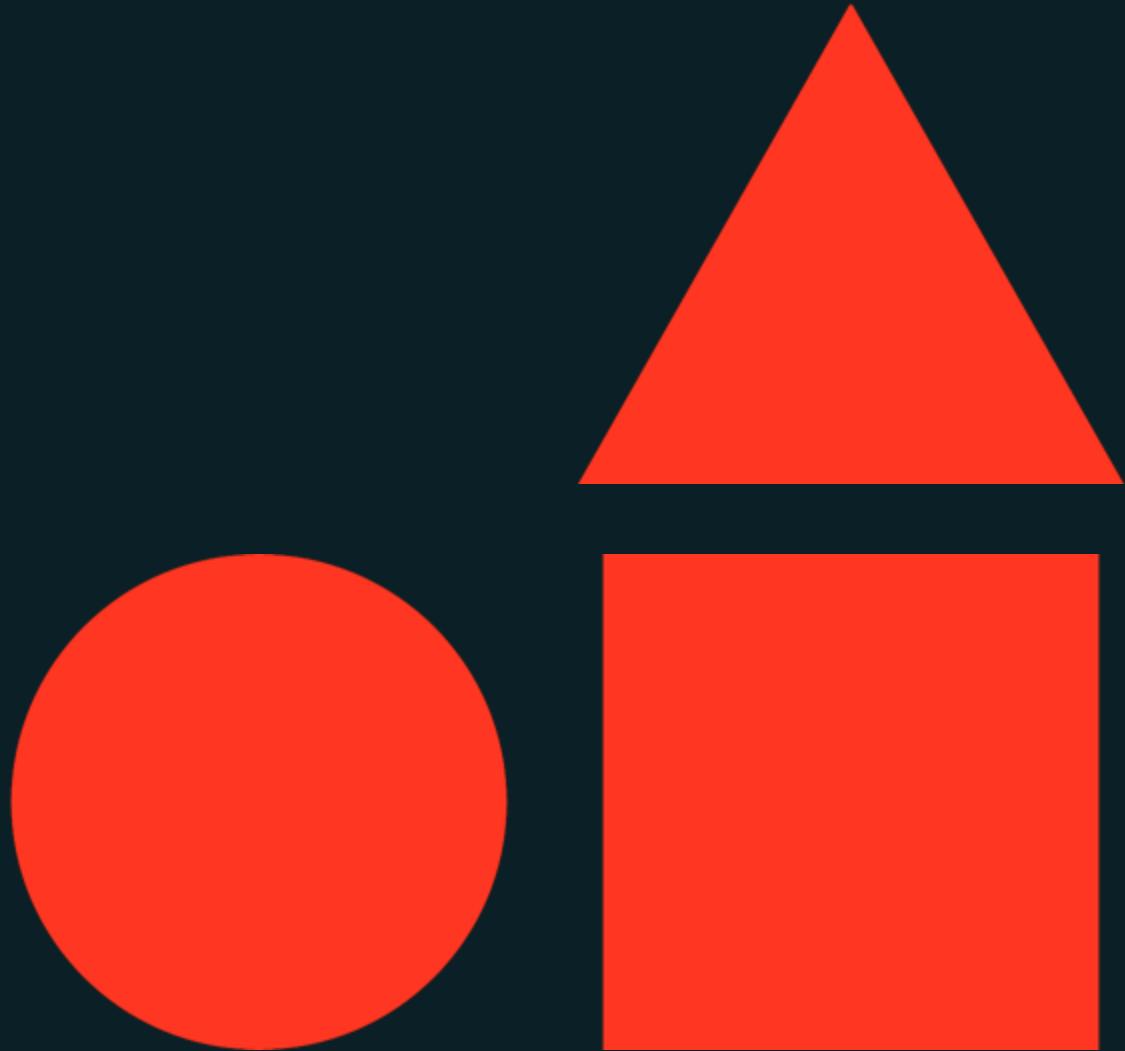




# Data Ingestion with Lakeflow Connect



---

**Databricks Academy**



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Course Learning Objectives

- Describe Lakeflow Connect as a scalable and simplified solution for data ingestion into Databricks from a variety of sources.
- Review the benefits of Delta tables and the Medallion architecture.
- Demonstrate how to ingest data from cloud object storage into Delta tables using CREATE TABLE AS, COPY INTO, and Auto Loader, including capturing input file metadata in Bronze layer tables.
- Explain how rescued columns are used during ingestion to manage malformed records.
- Illustrate techniques for ingesting and flattening semi structured JSON data from cloud storage.
- Describe available options for ingesting data from enterprise systems using Lakeflow Connect Managed Connectors.
- Discuss alternative ingestion methods such as MERGE INTO, Delta Sharing and Databricks Marketplace.



# Agenda

## Course Sections

- **Introduction to Data Engineering in Databricks**
- **Cloud Storage Ingestion with LakeFlow Connect Standard Connectors**
- **Enterprise Data Ingestion with LakeFlow Connect Managed Connectors**
- **Ingestion Alternatives**



# Course Prerequisites (REQUIRED)



## Fundamental Knowledge of the Databricks Platform

- Course: **Get Started with Databricks for Data Engineering**

OR

- Knowledge of Databricks Workspaces, Apache Spark, Delta Lake and the Medallion architecture, Unity Catalog Data Objects



# Course Prerequisites (REQUIRED)

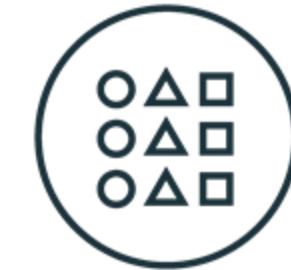


## Fundamental Knowledge of the Databricks Platform

- Course: Get Started with Databricks for Data Engineering

OR

- Knowledge of Databricks Workspaces, Apache Spark, Delta Lake and the Medallion architecture, Unity Catalog Data Objects



## Experience working with a variety of file types

- Parquet
- CSV
- JSON
- TXT and others



# Course Prerequisites (REQUIRED)



## Fundamental Knowledge of the Databricks Platform

- Course: **Get Started with Databricks for Data Engineering**

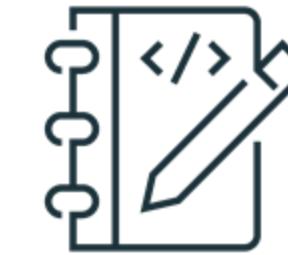
OR

- Knowledge of Databricks Workspaces, Apache Spark, Delta Lake and the Medallion architecture, Unity Catalog Data Objects



## Experience working with a variety of file types

- Parquet
- CSV
- JSON
- TXT and others



## Proficiency in SQL/Python and Databricks Notebooks

- Experience coding with SQL and Python
- Familiarity with executing code in Databricks notebooks



# Lab Exercise Environment



## Technical Details

- Your lab environment is provided by Vocareum.
- It will open in a new tab.
- It has been configured with the permissions and resources required to accomplish the tasks outlined in the lab exercise.
- Third party cookies must be enabled in your browser for Vocareum's user experience to work properly.
- Make sure to enable pop ups!

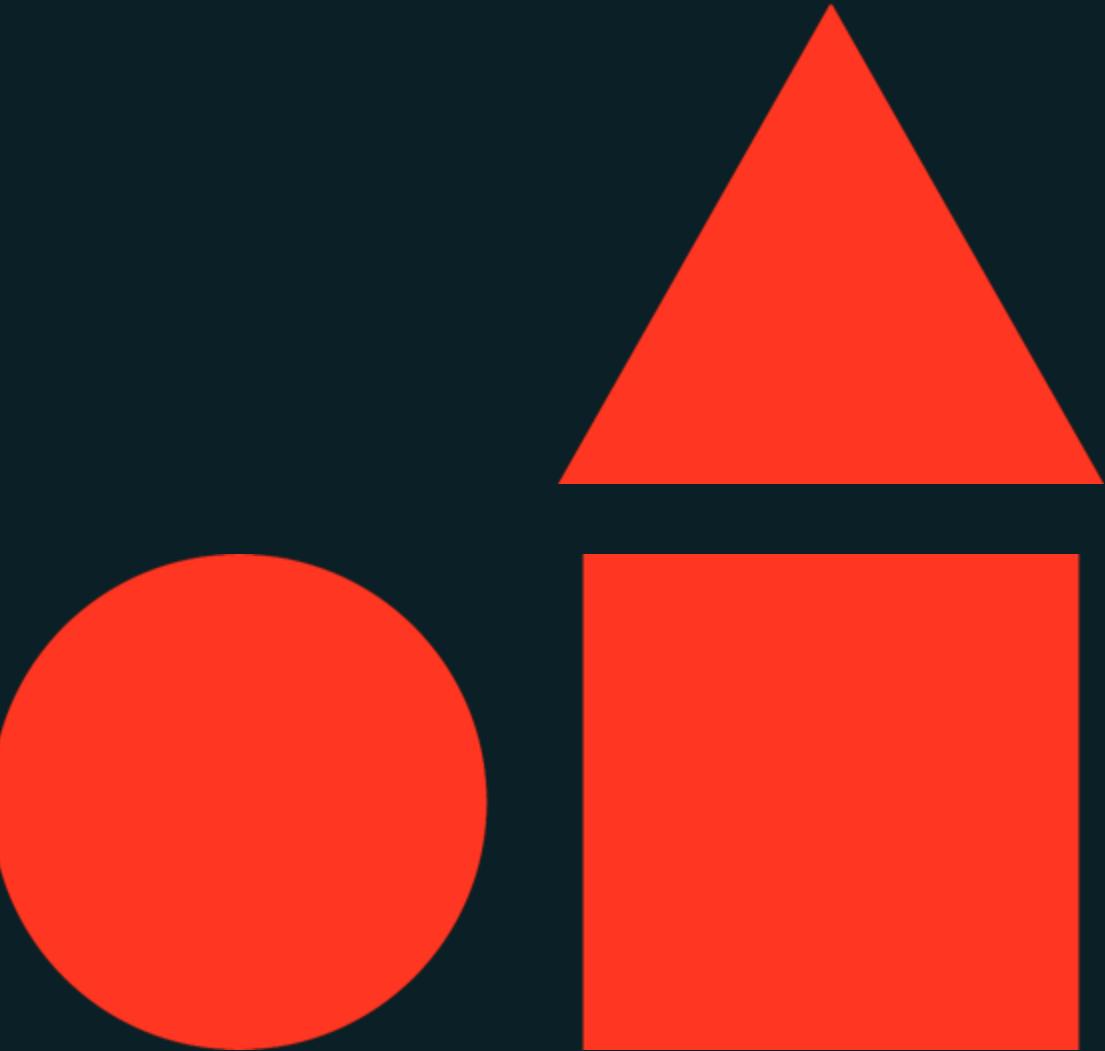




# Introduction to Data Engineering in Databricks

---

## Data Ingestion with Lakeflow Connect



# Agenda

## Section Overview – Introduction to Data Engineering in Databricks

- **Data Engineering in Databricks**
- **What is Lakeflow Connect?**
- **Delta Lake Review**
- **Exploring the Lab Environment**





Introduction to Data Engineering in Databricks

LECTURE

# Data Engineering in Databricks



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Data Engineering in Databricks



databricks **LAKEFLOW**

UNIFIED DATA ENGINEERING FOR THE DATA INTELLIGENCE PLATFORM

## CONNECT



Efficient  
ingestion  
connectors

## SPARK DECLARATIVE PIPELINES (Previously DLT)



Accelerated ETL  
development

## JOBS (Previously Workflows)



Reliable  
orchestration for  
analytics and AI

## INDUSTRY LEADING DATA PROCESSING ENGINE

(Apache Spark + Structured Streaming)

## UNIFIED GOVERNANCE

 **Unity Catalog**

## OPTIMIZED STORAGE

 **DELTA LAKE**  **Parquet**  **ICEBERG**



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Data Engineering in Databricks



databricks

LAKEFLOW

UNIFIED DATA ENGINEERING FOR THE DATA INTELLIGENCE PLATFORM

This course will focus on data ingestion with **LakeFlow Connect**

CONNECT



Efficient  
ingestion  
connectors

DECLARATIVE PIPELINES  
(Previously DLT)



Accelerated ETL  
development

JOBS  
(Previously Workflows)



Reliable  
orchestration for  
analytics and AI

INDUSTRY LEADING DATA PROCESSING ENGINE

(Apache Spark + Structured Streaming)

UNIFIED GOVERNANCE

Unity Catalog

OPTIMIZED STORAGE

DELTA LAKE

Parquet

ICEBERG



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Introduction to Data Engineering in Databricks

LECTURE

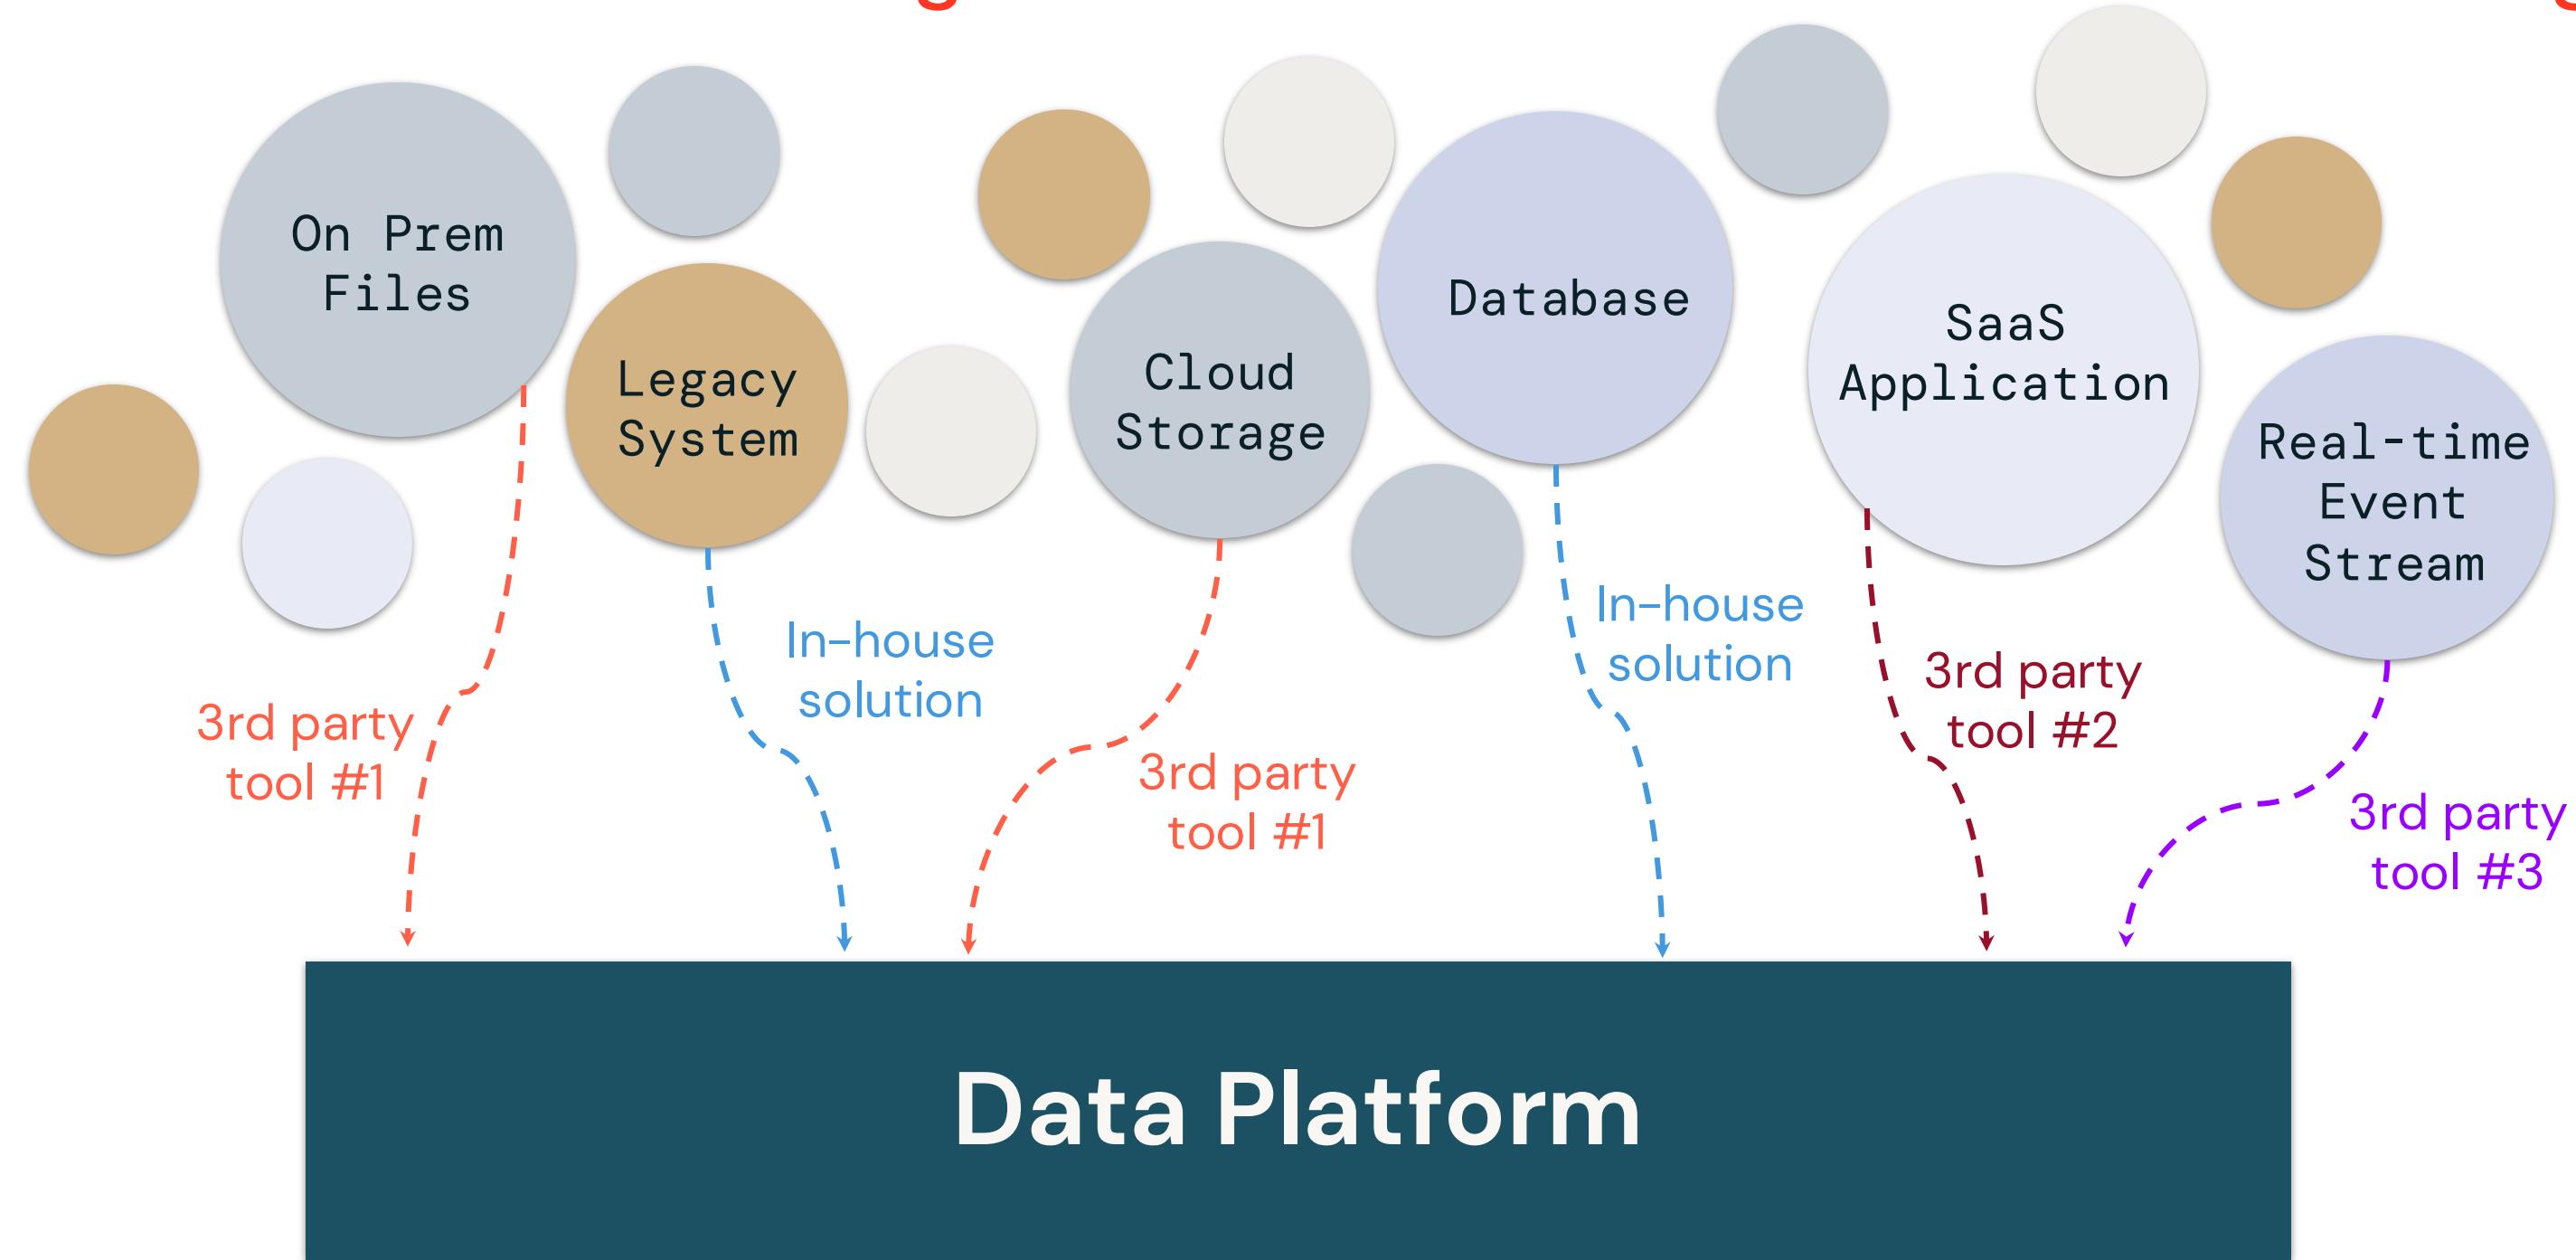
# What is Lakeflow Connect?



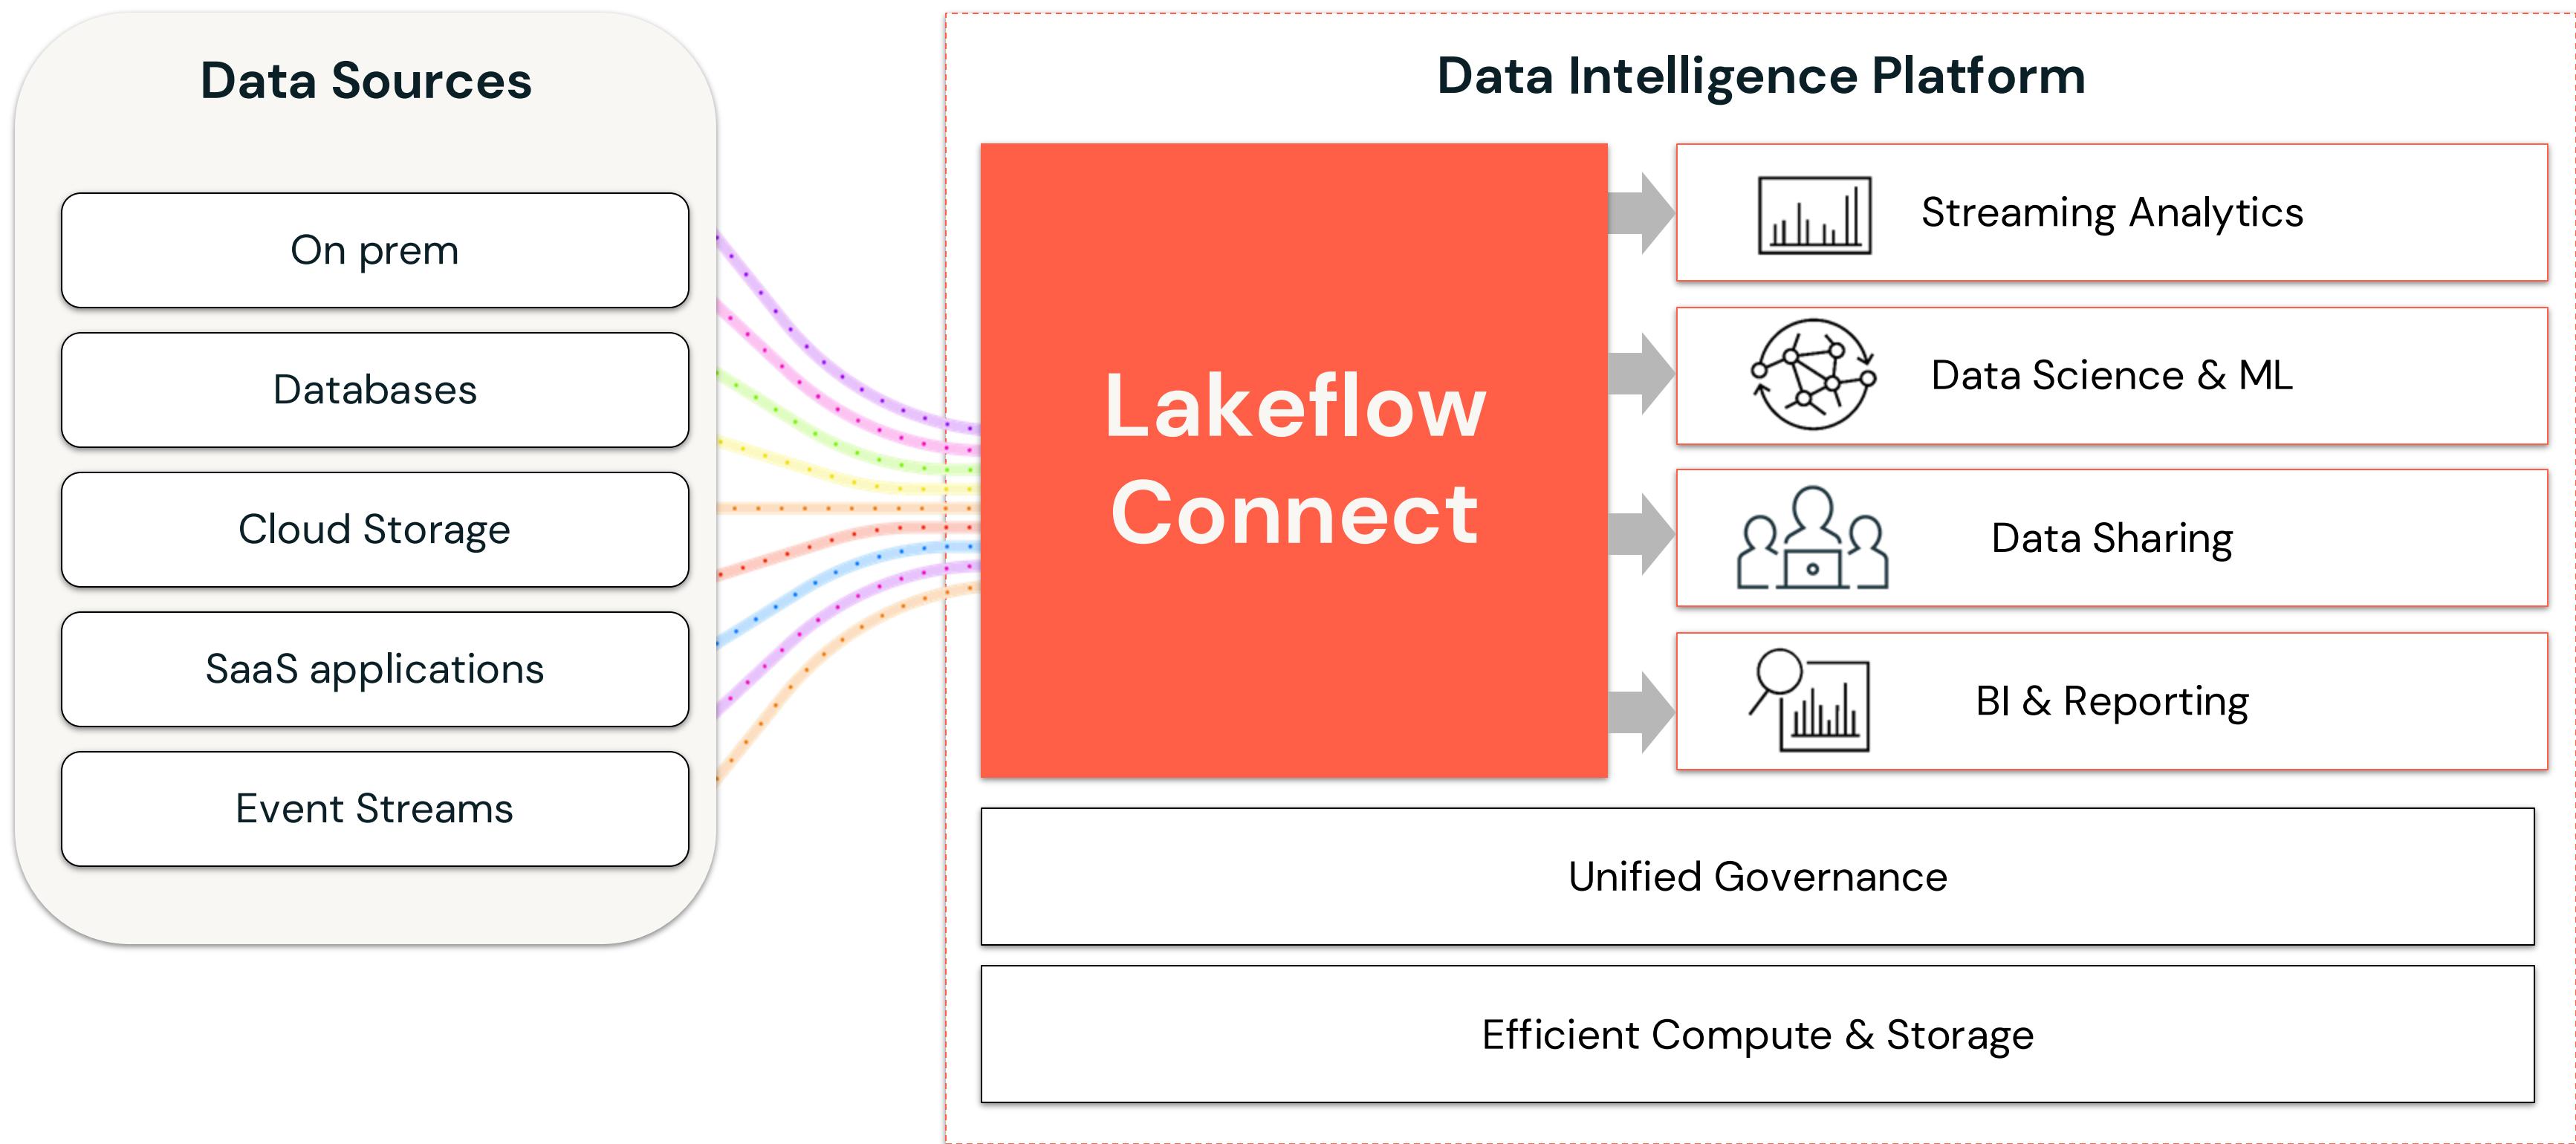
© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# What is Lakeflow Connect?

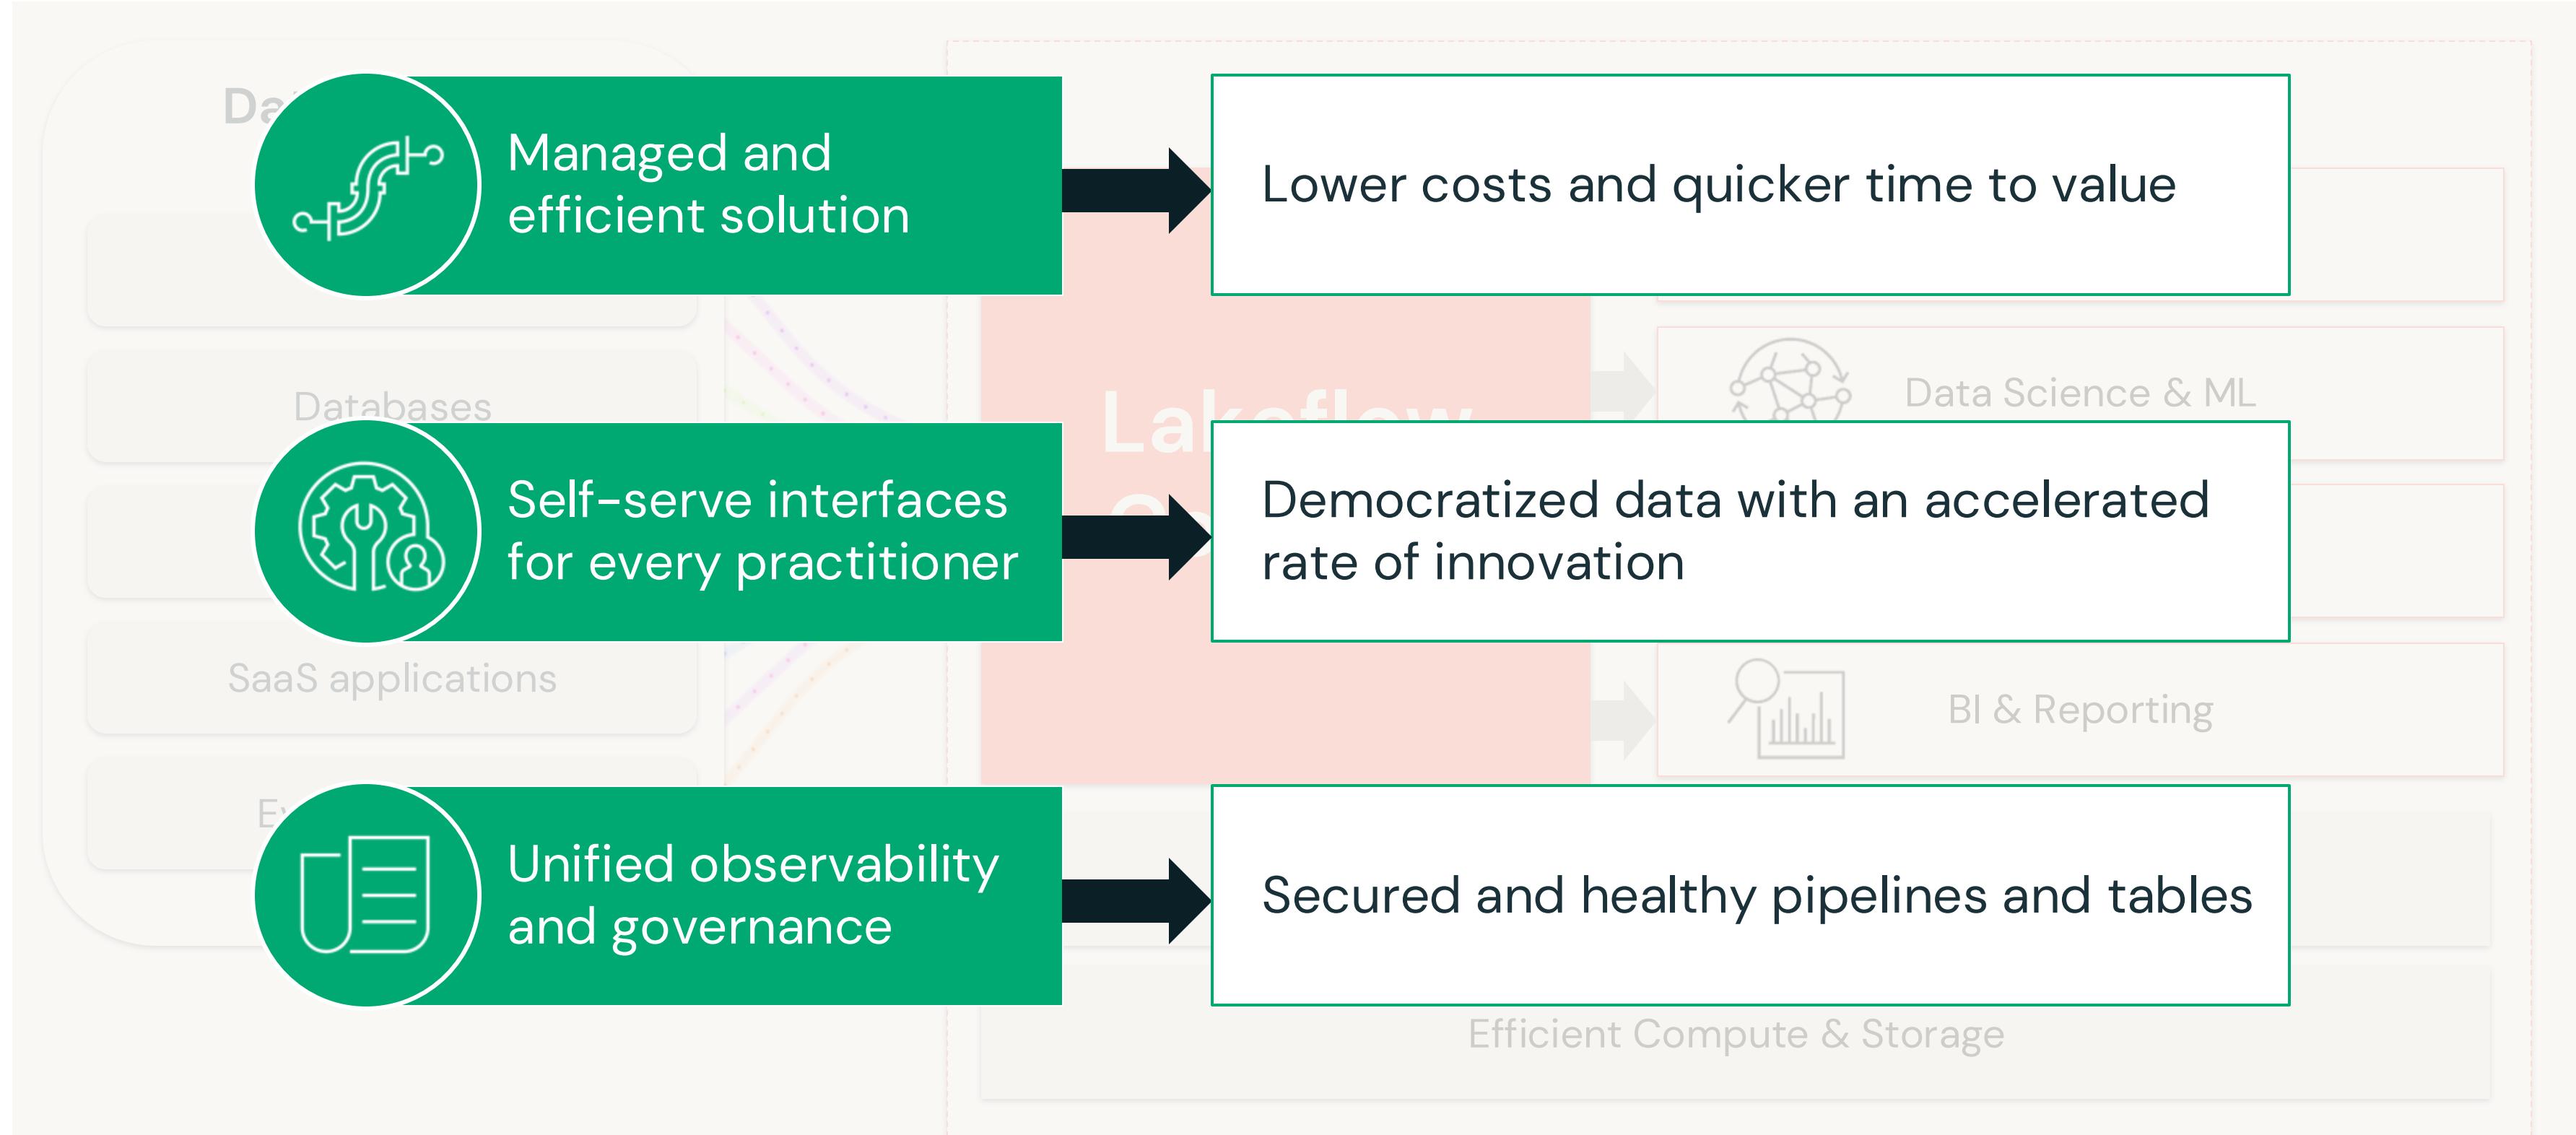
Organizations are Resorting to a Patchwork of Solutions for Data Ingestion



# Lakeflow Connect is all Ingestion



# Built-in connectors for the Data Intelligence Platform



# What is Lakeflow Connect?

## Lakeflow Connect – Connectors Overview



### Upload Files

- Uploading local files to Databricks
  - Upload a file to a volume
  - Create a table from a local file



### Standard Connectors

Ingest data into the lakehouse using various sources and methods:

#### Supported Sources:

- Cloud Object Storage
- Kafka
- Other Sources

#### Ingestion Methods:

- Batch
- Incremental Batch
- Streaming



### Managed Connectors

Ingesting data into the lakehouse from:

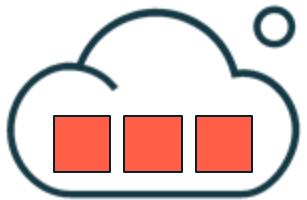
- Software as a Service (SaaS) applications
- Databases

Leverage efficient **incremental reads** and **writes** to make data ingestion faster, scalable, and more cost-efficient

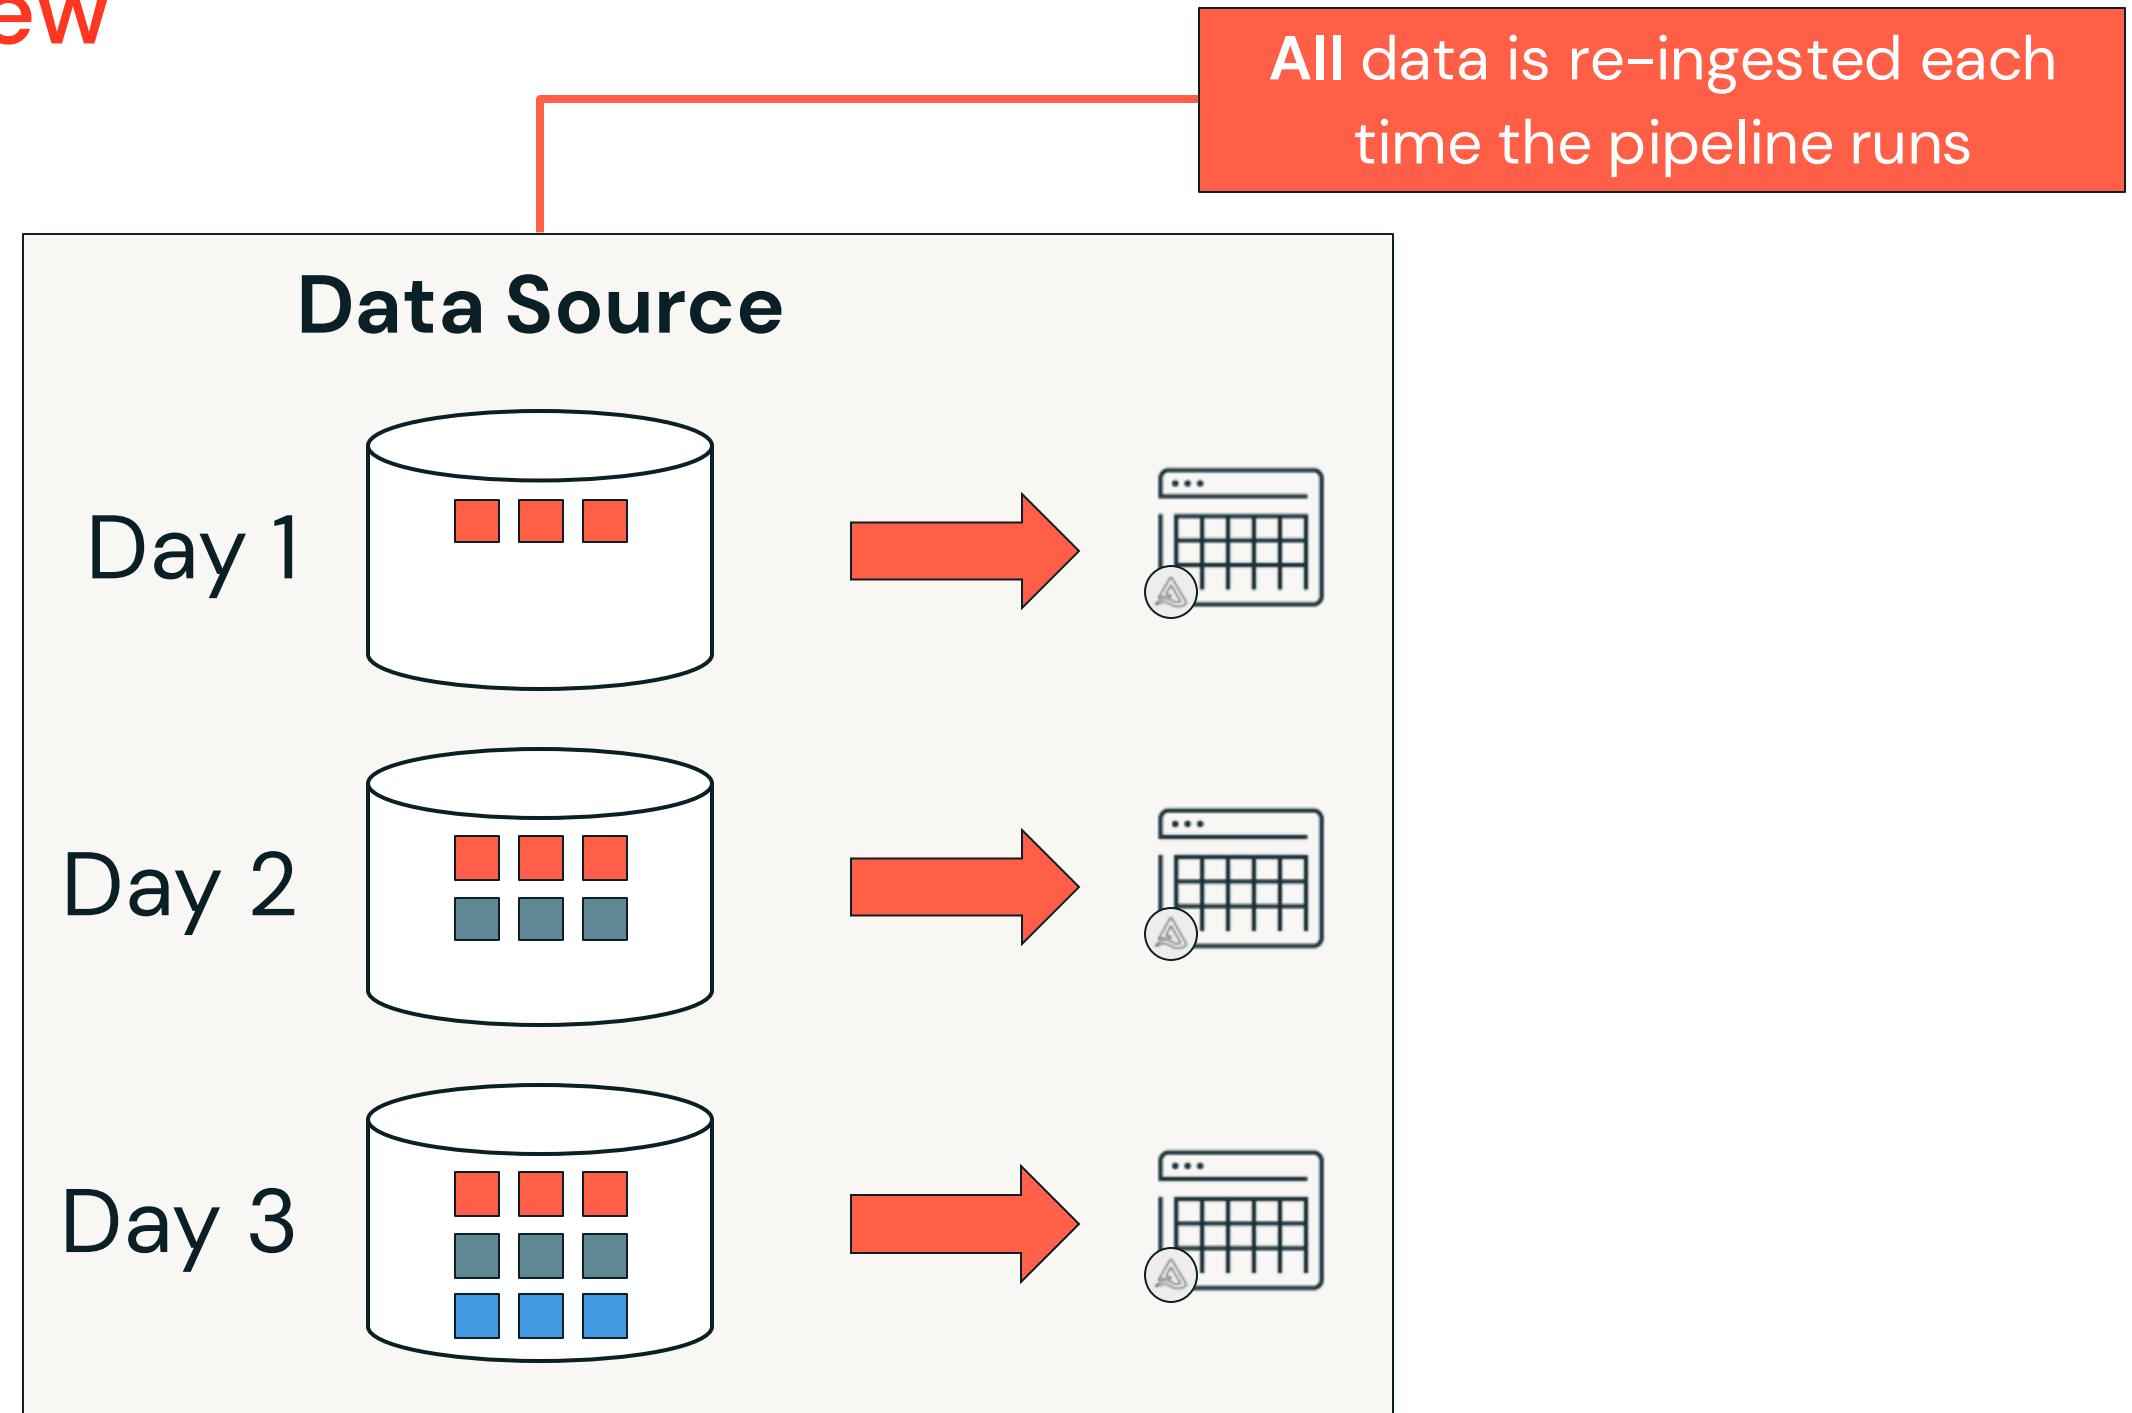


# What is Lakeflow Connect?

## Ingestion Methods Overview

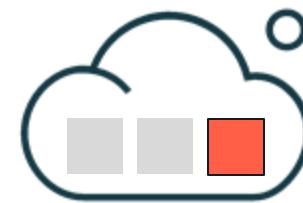
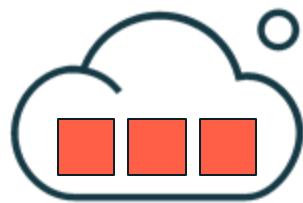


- **Batch**
  - Load data as **batches of rows** into **Databricks**, often based on a schedule
  - Traditional batch ingestion **processes all records** each time it runs
    - `CREATE TABLE AS (CTAS)`
    - `spark.read.load()`



# What is Lakeflow Connect?

## Ingestion Methods Overview

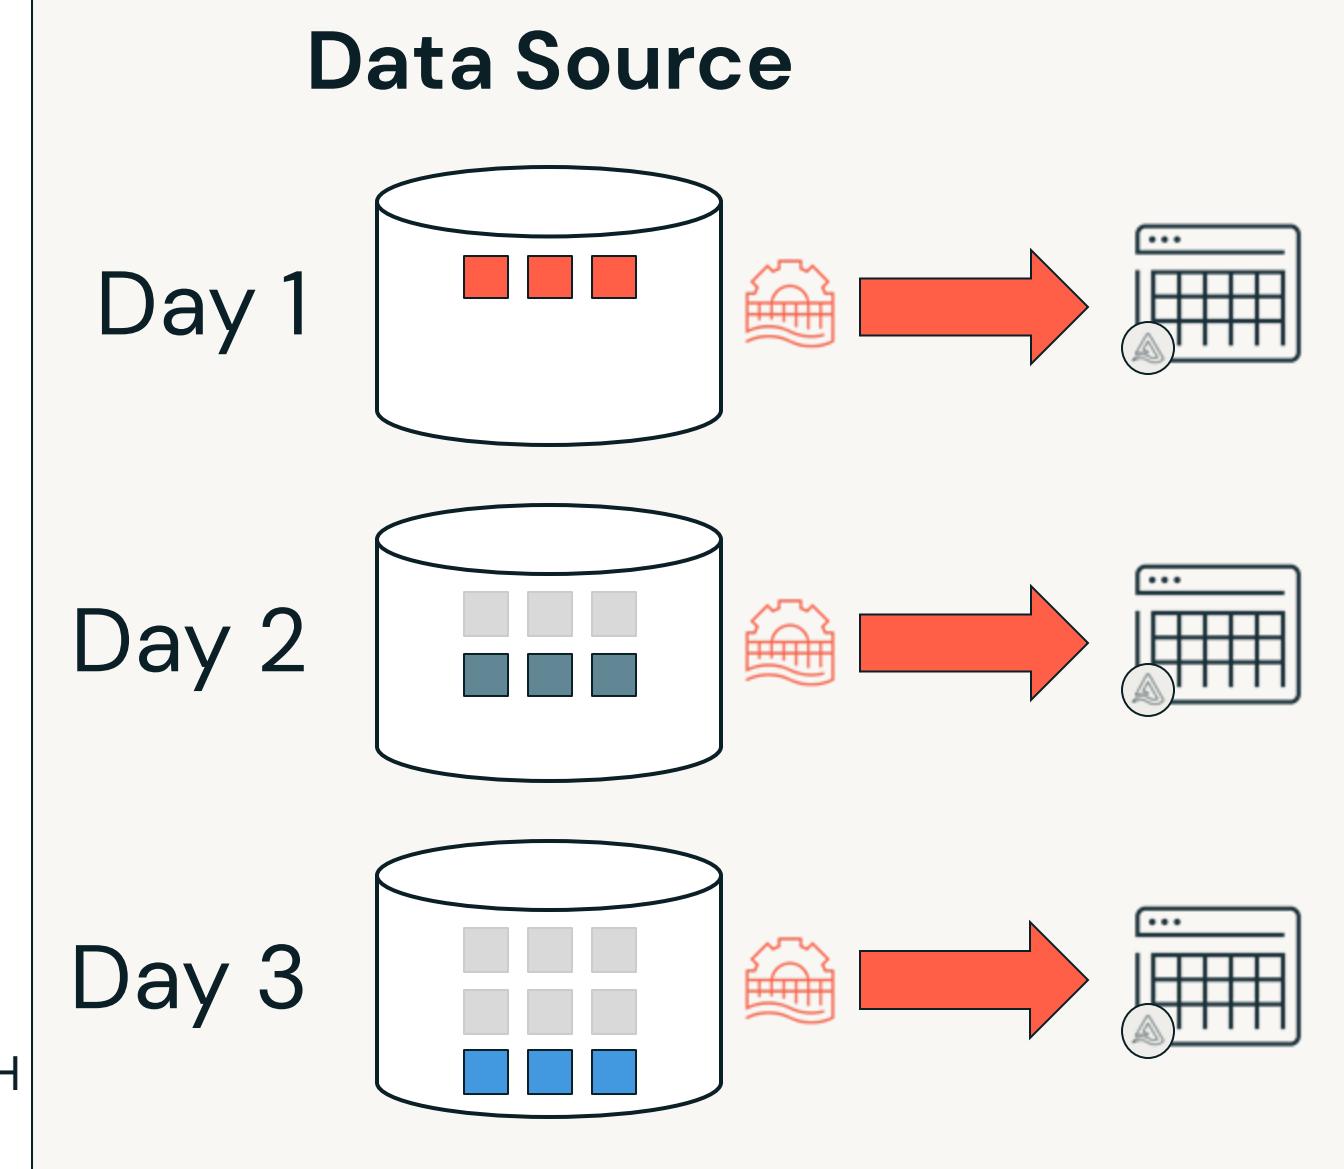


- Load data as **batches** of rows into Databricks, often based on a schedule
- Traditional batch ingestion processes all records each time it runs
  - `CREATE TABLE AS (CTAS)`
  - `spark.read.load()`

### Batch

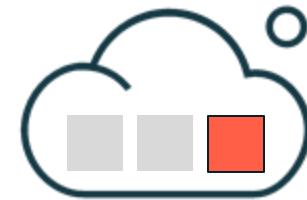
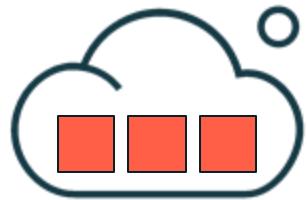
- **Incremental Batch**
  - Only new data is ingested, previously loaded records are **skipped automatically**
  - Provides **faster** and more **resource efficient** ingestion by processing less data
    - `COPY INTO`
    - `spark.readStream(Auto Loader with timed trigger)`
    - Declarative Pipelines (`CREATE OR REFRESH STREAMING TABLE`)

Ingests (appends) new data only, skipping previously loaded records



# What is Lakeflow Connect?

## Ingestion Methods Overview



### Batch

- Load data as **batches** of rows into Databricks, often based on a schedule
- Traditional batch ingestion **processes all records** each time it runs
  - `CREATE TABLE AS (CTAS)`
  - `spark.read.load()`

### Incremental Batch

- Only new data is ingested, previously loaded records are **skipped automatically**
- **Faster** ingestion and better **resource efficiency** by processing less data
  - `COPY INTO`
  - `spark.readStream` (Auto Loader with timed trigger)
  - Declarative Pipelines(`CREATE OR REFRESH STREAMING TABLE`)

### Streaming

- **Continuously load data** rows or batches of data rows as it is generated so you can query it as it **arrives in near real-time**
- **Micro-batch** processes small batches a **very short, frequent intervals**
  - `spark.readStream` (Auto Loader with continuous trigger)
  - Declarative Pipelines (trigger mode continuous)





Introduction to Data Engineering in Databricks

LECTURE

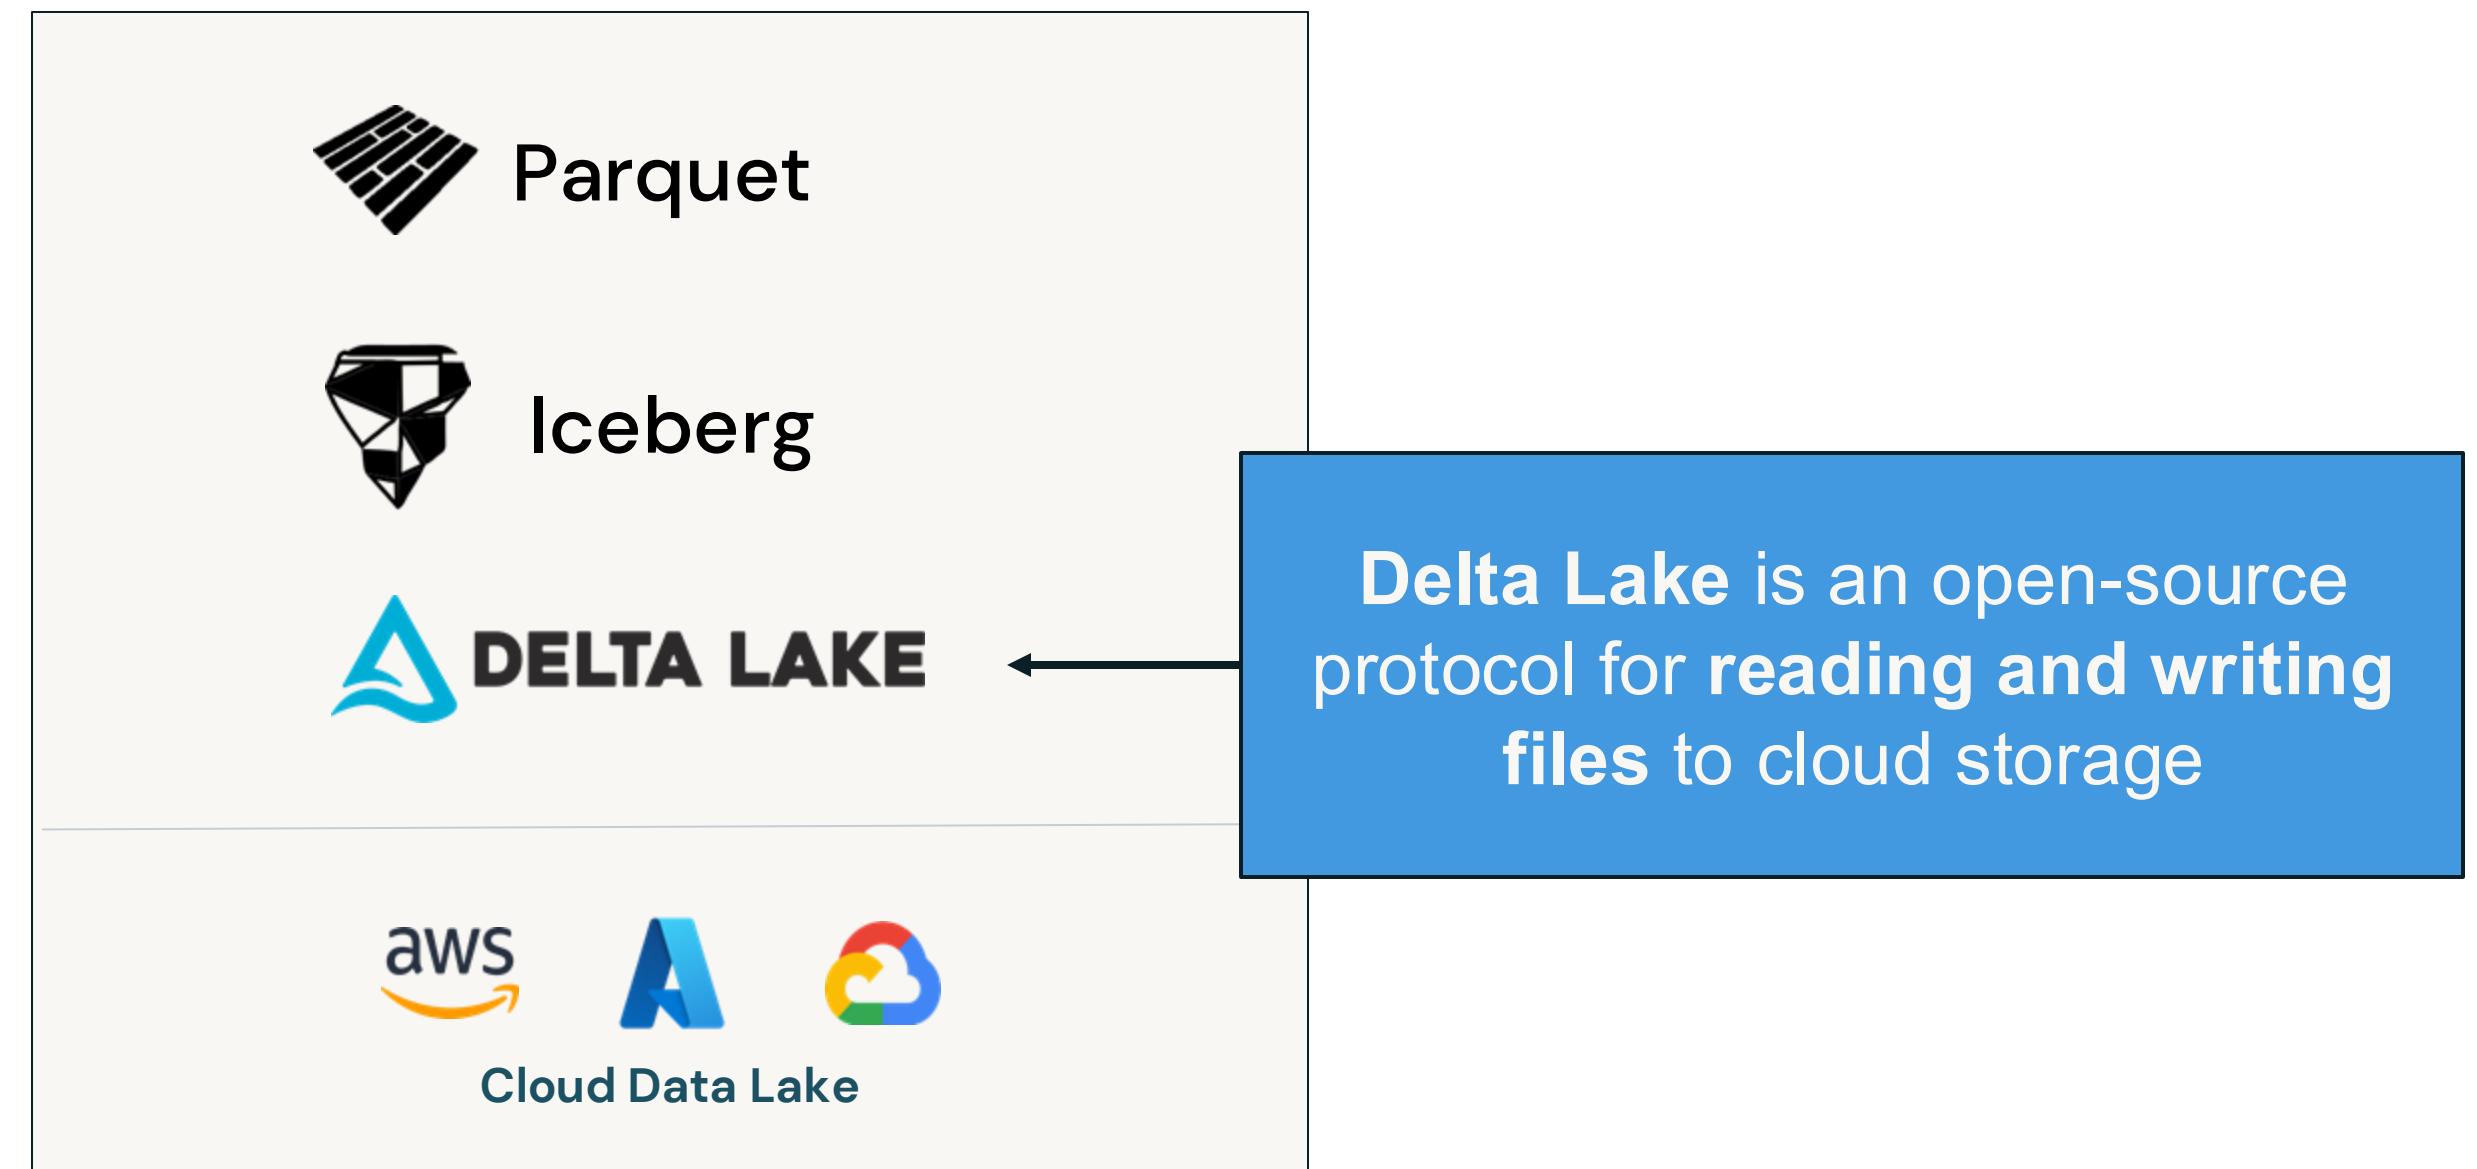
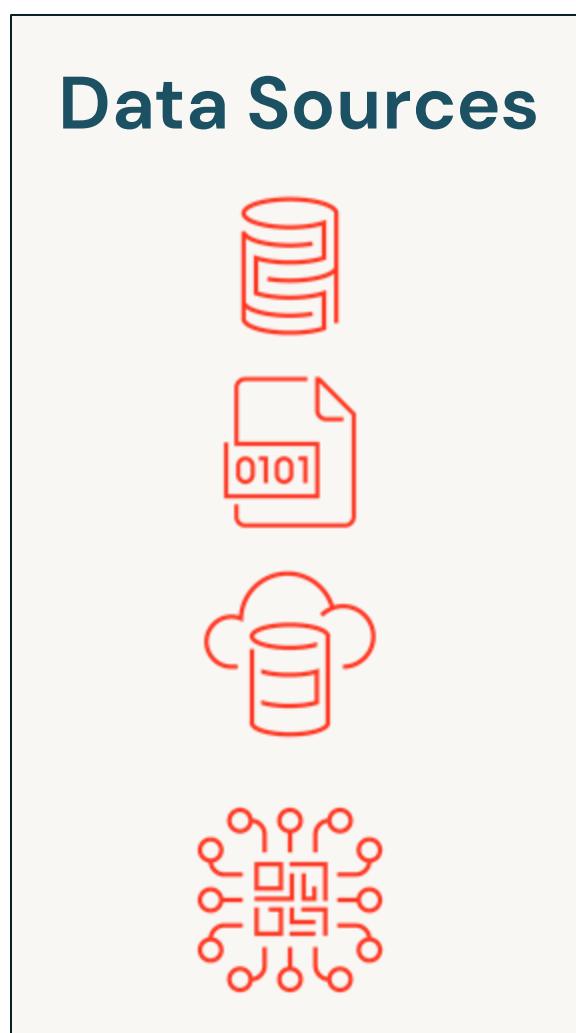
# Delta Lake Review



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

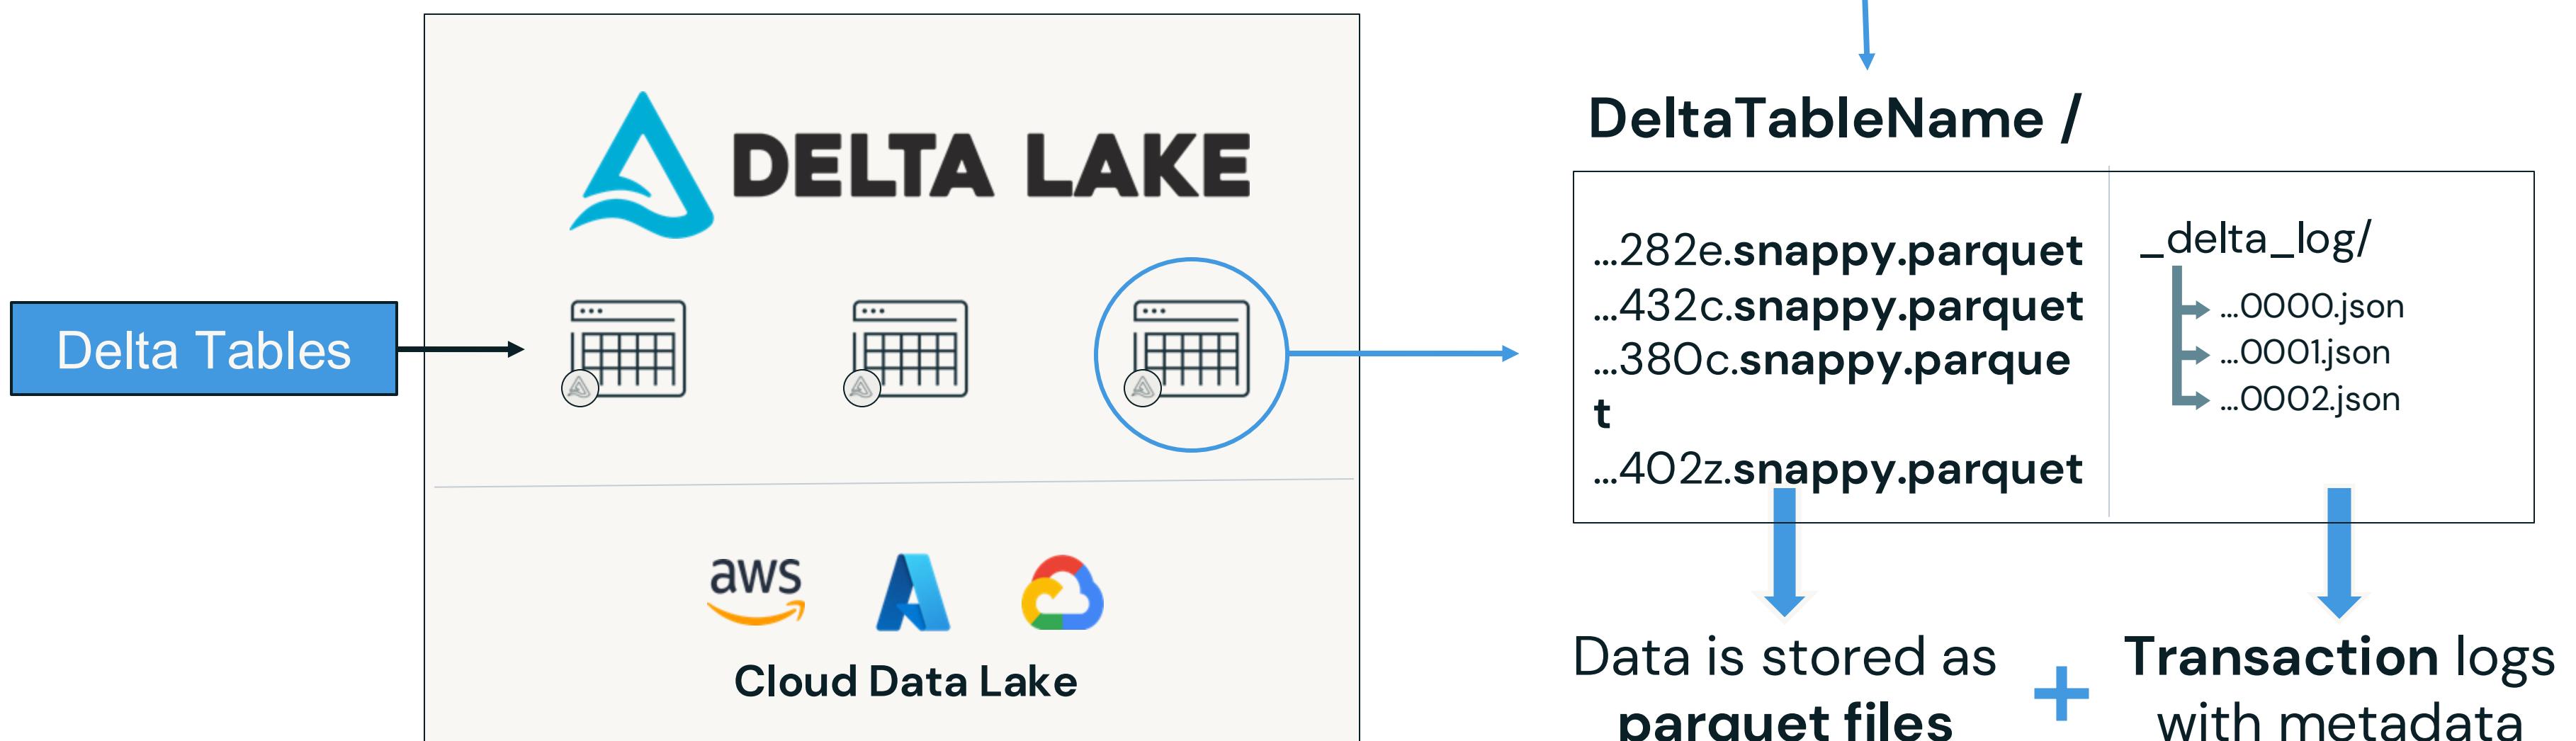
# Delta Lake Review

## Ingesting Data Into Delta Lake



# Delta Lake Review

## Delta Table Components Overview



# Delta Lake Review

## Delta Tables Key Features Review



ACID Transactions



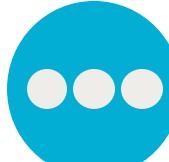
Data Manipulation Language (DML)



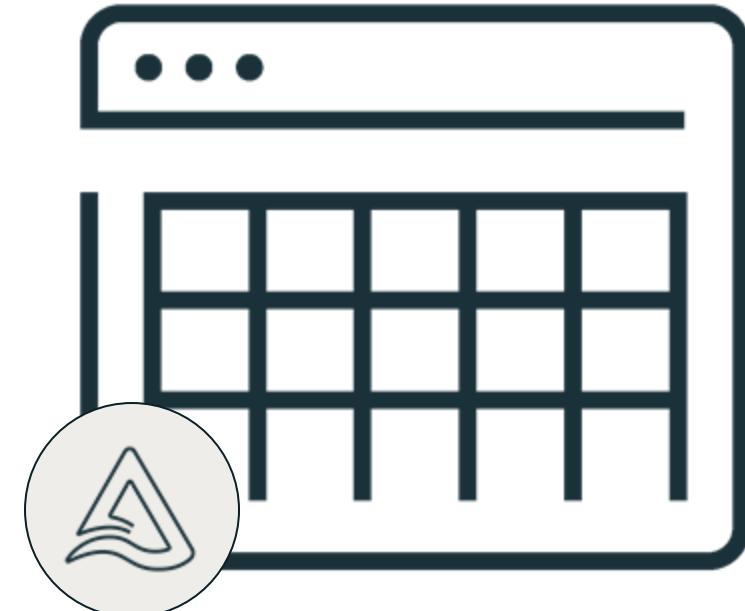
Time Travel



Schema Evolution and  
Enforcement

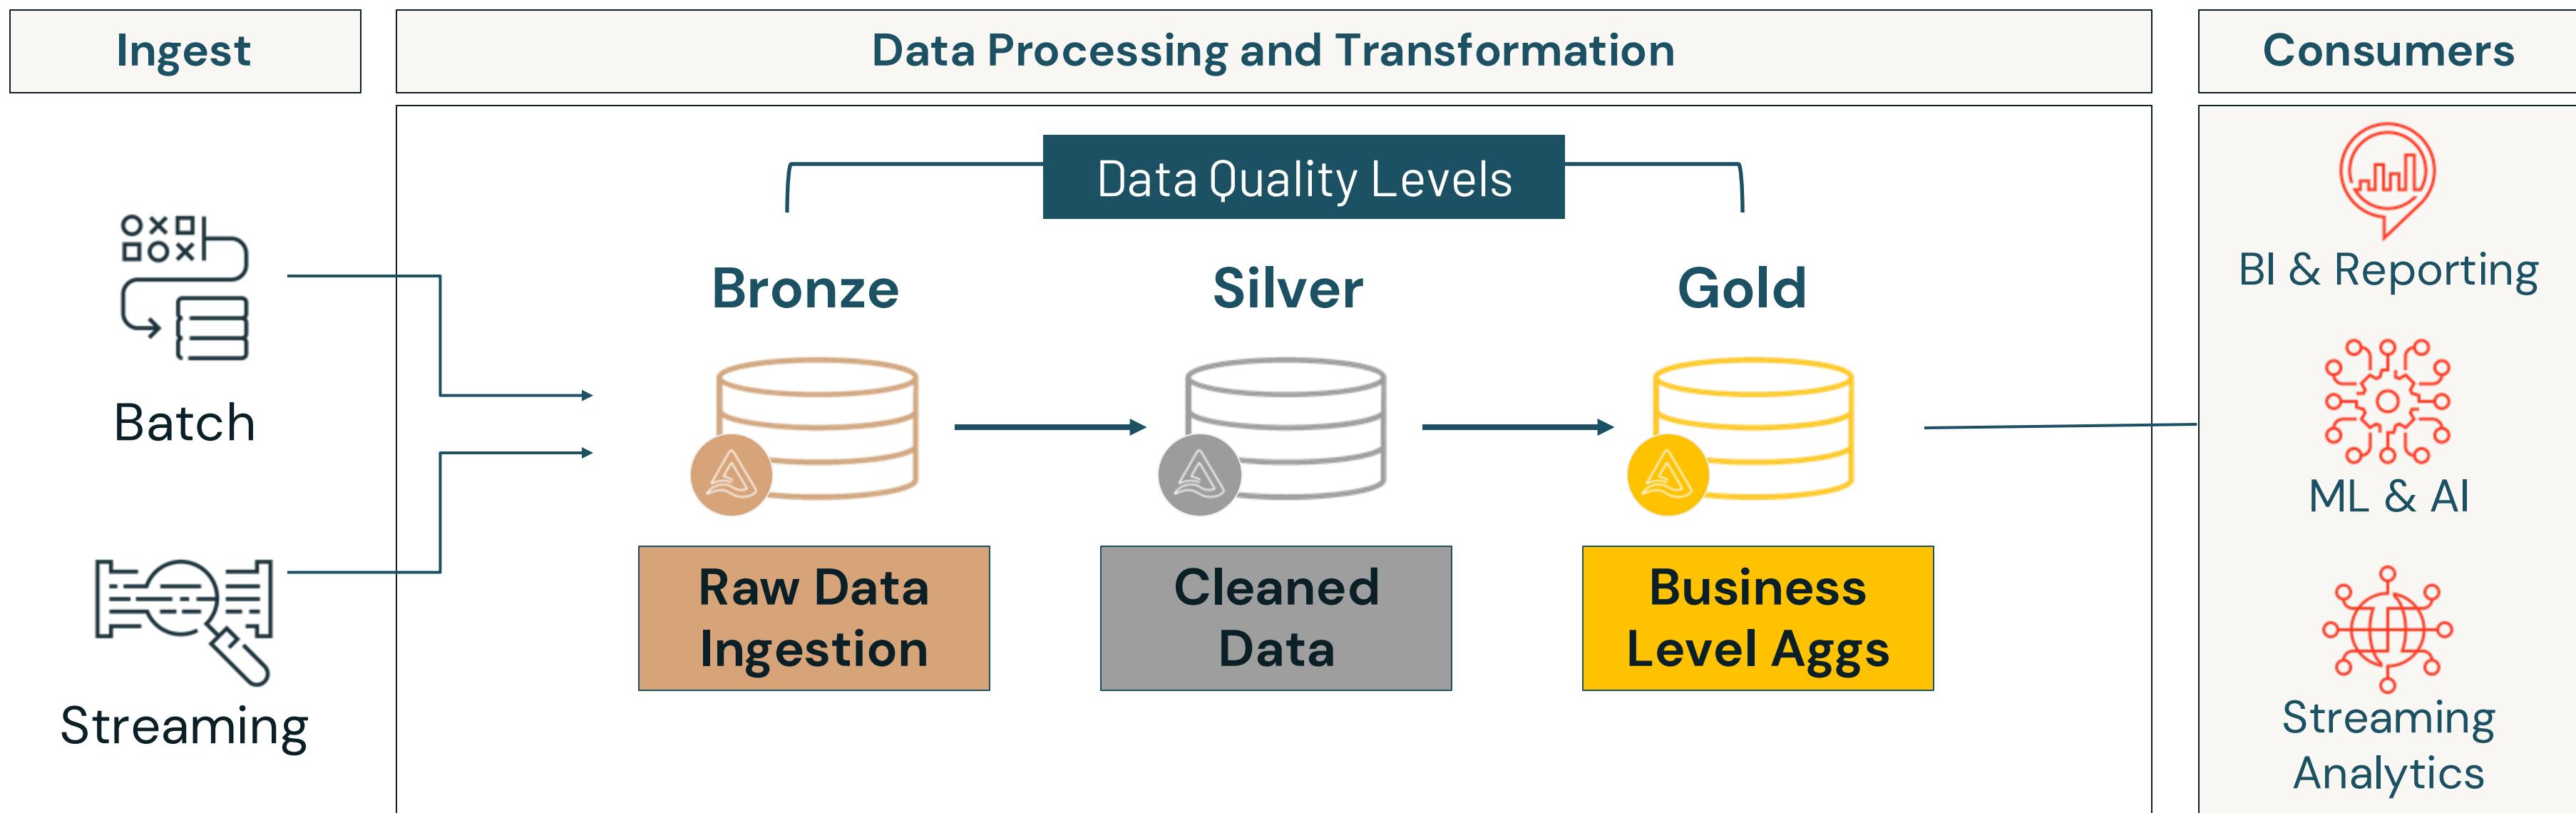


Many more!



# Delta Lake Review

## Medallion Architecture (Multi Hop) Review





## Introduction to Data Engineering in Databricks

### DEMONSTRATION

# Exploring the Lab Environment



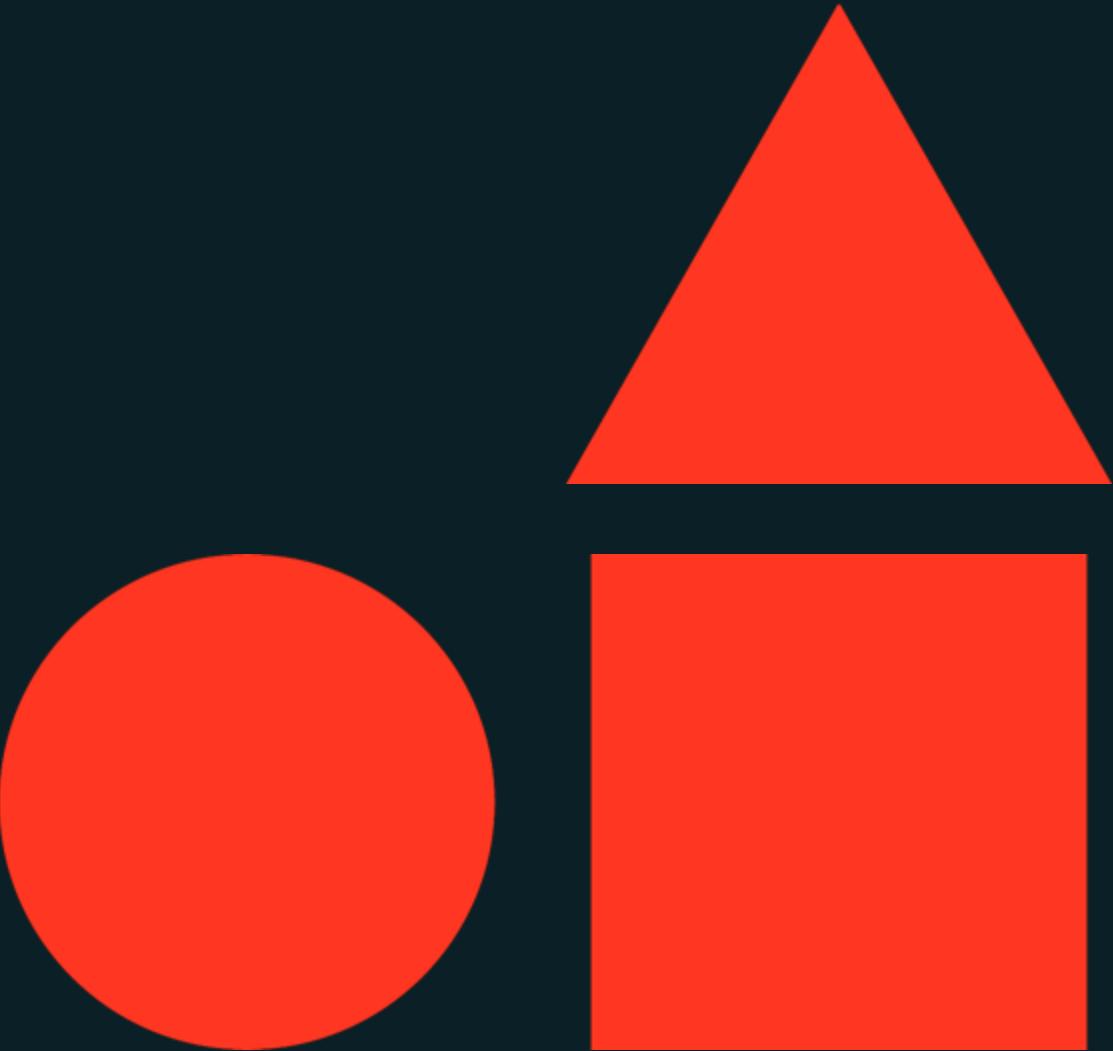
Notebook: 01 – Exploring the Lab Environment



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



# Cloud Storage Ingestion with LakeFlow Connect Standard Connectors



## Data Ingestion with Lakeflow Connect



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Agenda

## Section Overview – Cloud Storage Ingestion with LakeFlow Connect Standard Connectors

- **Introduction to Data Ingestion from Cloud Storage**
- **Appending Metadata Columns on Ingest**
- **Working with the Rescued Data Column**
- **Ingesting Semi-Structured Data: JSON**





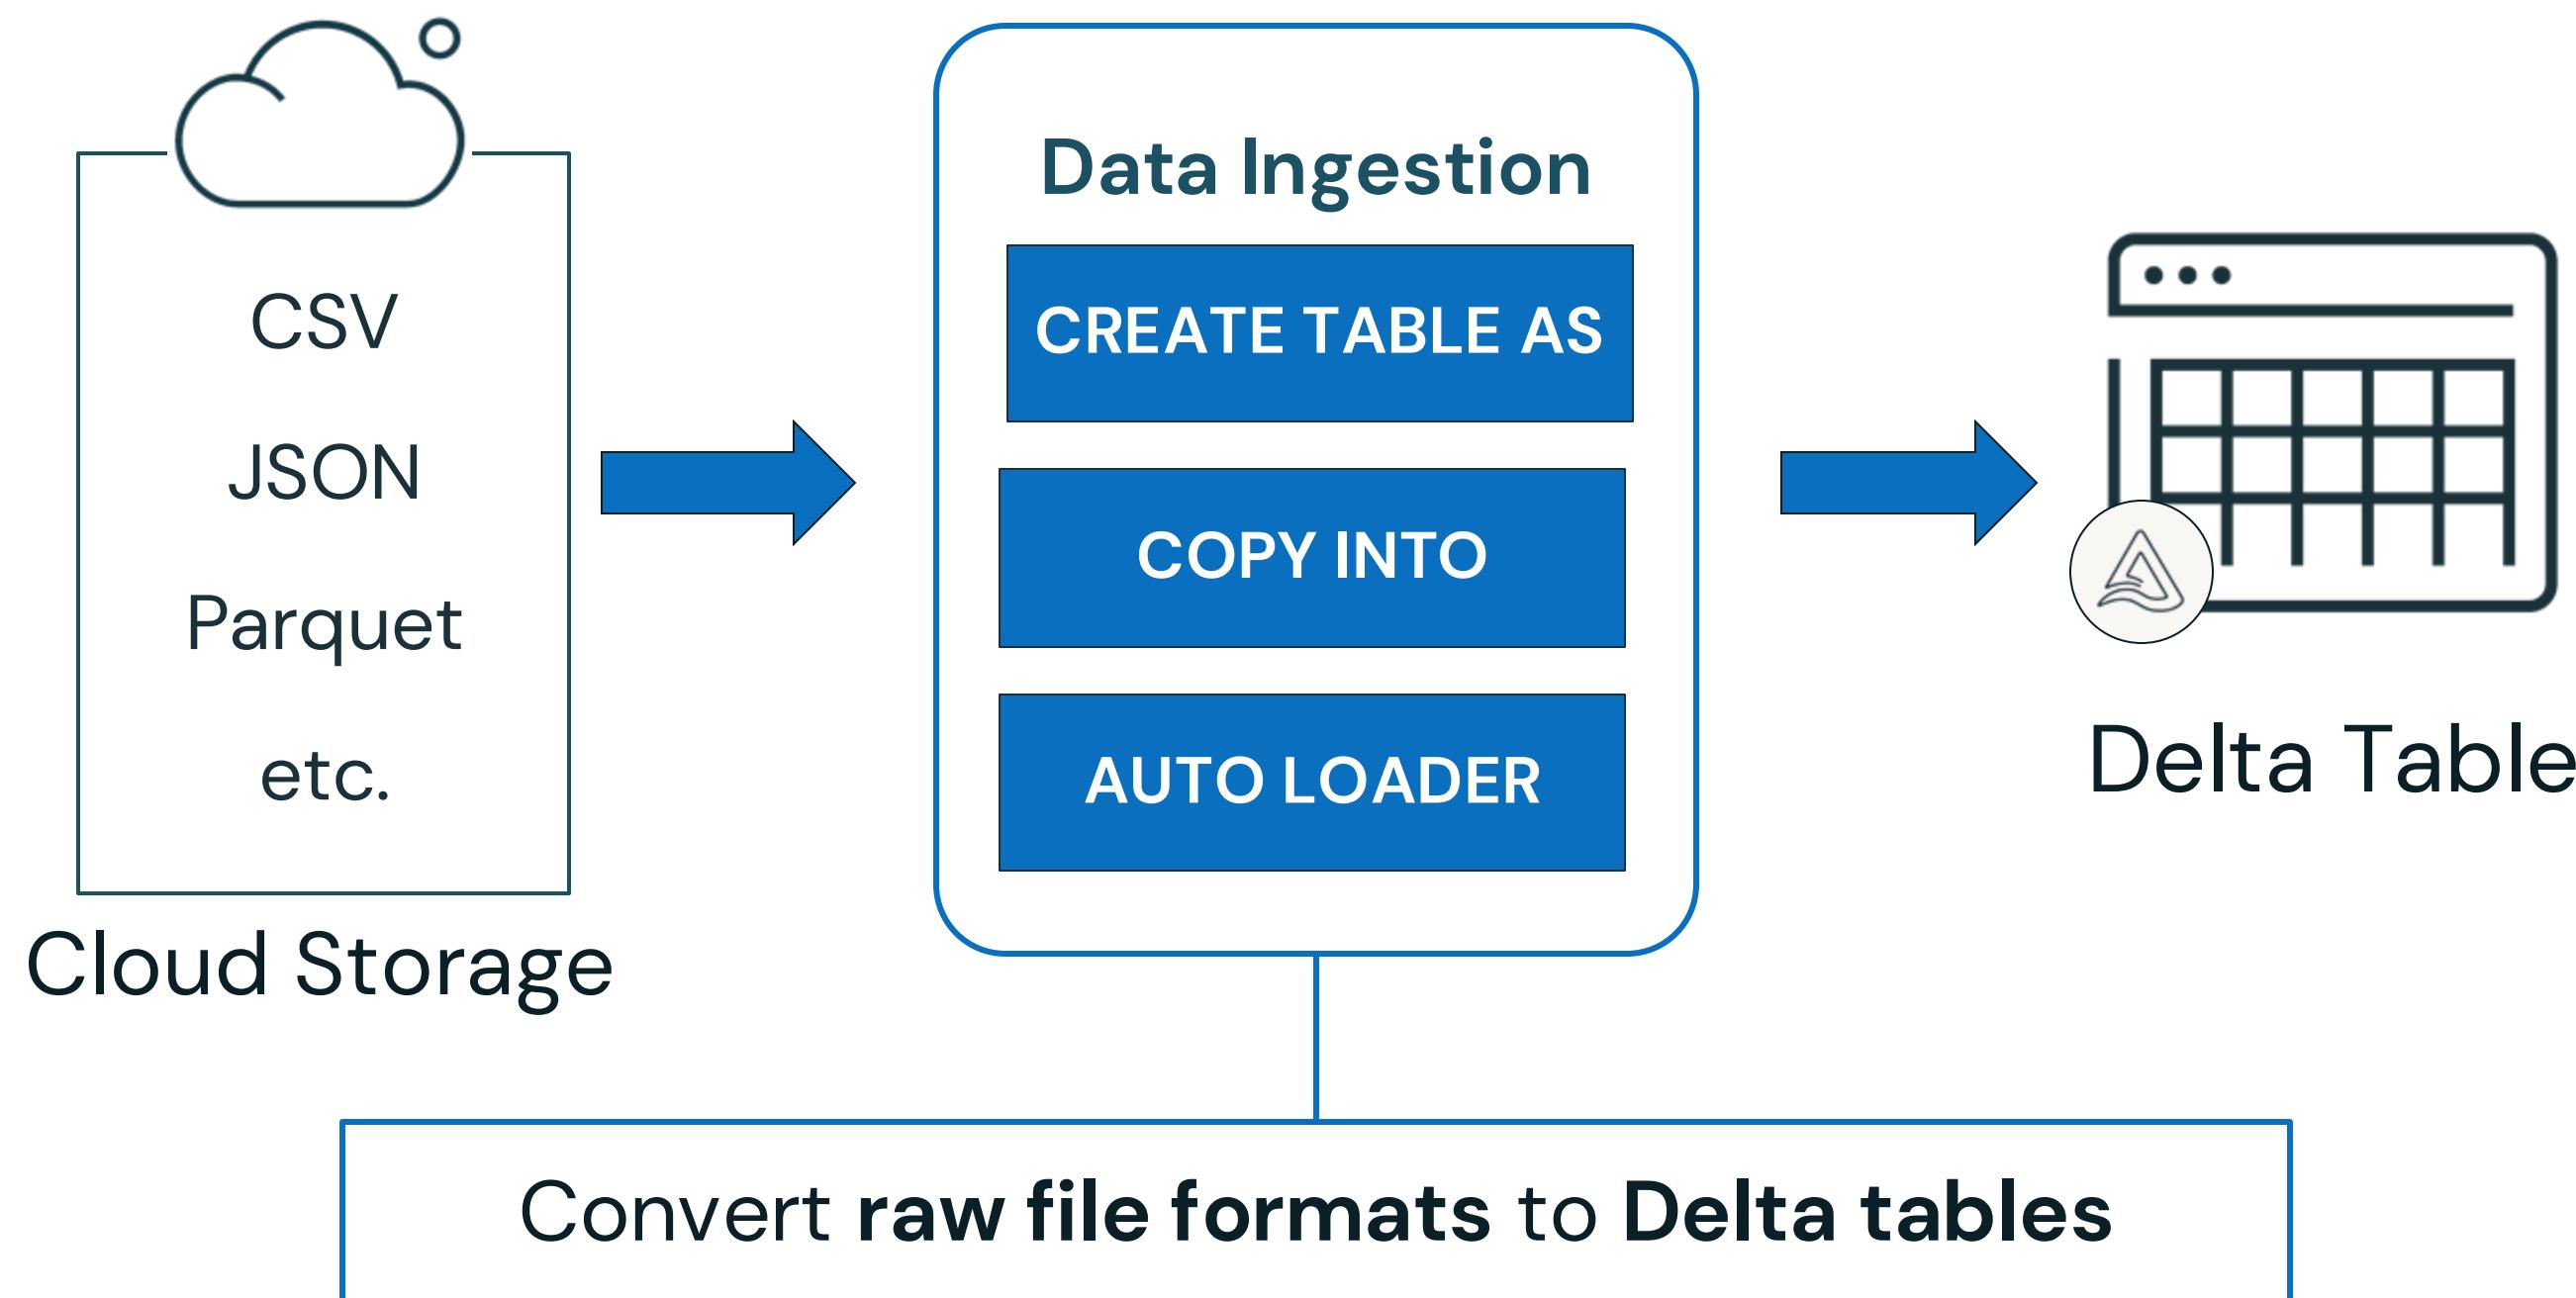
Cloud Storage with LakeFlow Connect  
Standard Connectors  
**LECTURE**

# Introduction to Data Ingestion from Cloud Storage



# Data Ingestion from Cloud Storage

## Data Ingestion Patterns From Cloud Object Storage



# (1 of 3) CREATE TABLE AS (CTAS)



# Data Ingestion from Cloud Storage

## Method 1 – Batch – CREATE TABLE AS (CTAS)



**CREATE TABLE AS (CTAS)** creates a Delta table **by default** from files in cloud object storage

```
CREATE TABLE new_table AS
SELECT *
FROM read_files(
    <path_to_file(s)>,
    format => '<file_type>',
    <other_format_specific_options>
);
```

The **read\_files()** function reads files under a provided location and returns the data in **tabular form**

- Supports reading **file formats** like JSON, CSV, XML, TEXT, BINARYFILE, PARQUET, AVRO, and ORC file formats.
- Can detect the **file format automatically and infer a unified schema** across all files.
- Specify **specific file format options** to read in the data based on the source file format
- Can be used in **streaming tables** to **incrementally ingest** files into Delta Lake using Auto Loader



# (2 of 3) COPY INTO



# Data Ingestion from Cloud Storage

## Method 2 – Incremental Batch – COPY INTO (legacy)



```
CREATE TABLE new_table;  
  
COPY INTO new_table  
FROM '<dir_path>'  
FILEFORMAT=<file_type>  
FORMAT_OPTIONS(<options>)  
COPY_OPTIONS(<options>)
```

1. **Create an empty table to copy data into**
  - You can create an empty table **without** a schema
  - You also can **explicitly** create the table with a schema



# Data Ingestion from Cloud Storage

## Method 2 – Incremental Batch – COPY INTO (legacy)



```
CREATE TABLE new_table;  
  
COPY INTO new_table  
FROM '<dir_path>'  
FILEFORMAT=<file_type>  
FORMAT_OPTIONS(<options>)  
COPY_OPTIONS(<options>)
```

### 2. COPY INTO for incremental batch ingestion

- Is a **retriable and idempotent operation** and will **skip files** that have already been loaded (**incremental**)
- Supports various common file types like parquet, JSON, XML, etc
- **FROM** specifies the path of the cloud storage location continuously adding files
- **FORMAT\_OPTIONS()** control how the source files are parsed and interpreted. The available options depends on the file format
- **COPY\_OPTIONS()** controls the behavior of the COPY INTO operation itself, such as schema evolution (**mergeSchema**) and idempotency (**force**)



# (3 of 3) AUTO LOADER



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Data Ingestion from Cloud Storage

## Method 3 – Incremental Batch or Streaming – Auto Loader

- **Incrementally** and **efficiently** processes new data files (in batch or streaming) as they arrive in cloud storage **without any additional setup**
- Auto Loader has support for both **Python** and **SQL (leveraging Declarative Pipelines)**
- You can use Auto Loader to process billions of files
- Auto Loader is built upon **Spark Structured Streaming**
- A deep dive into Auto Loader is **out of scope for this course**, please refer to these links and courses for more in depth information resources:
  - [Documentation](#) and [Tutorials](#)
    - [Stream Processing and Analysis with Apache Spark™ Course](#)
    - [Build Data Pipelines with Lakeflow Declarative Pipelines Course](#)
    - [Databricks Streaming and DLT \(Advanced\) Course](#)



# Data Ingestion from Cloud Storage

## Method 3 – Incremental Batch or Streaming – Auto Loader

### Python Auto Loader

```
(spark  
.readStream  
    .format("cloudFiles")  
    .option("cloudFiles.format", "json")  
    .option("cloudFiles.schemaLocation", "<checkpoint_path>")  
    .load("/Volumes/catalog/schema/files")  
  
.writeStream  
    .option("checkpointLocation", "<checkpoint_path>")  
    .trigger(processingTime="5 seconds")  
    .toTable("catalog.database.table")  
)
```

### Auto Loader with SQL (Declarative Pipelines)

```
CREATE OR REFRESH STREAMING TABLE catalog.schema.table  
SCHEDULE EVERY 1 HOUR  
AS  
SELECT *  
FROM STREAM read_files(  
    '<dir_path>',  
    format => '<file_type>'  
)
```



# Data Ingestion from Cloud Storage

FEATURE	CREATE TABLE AS (CTAS) + spark.read	COPY INTO	Auto Loader
Ingestion Type	Batch	Incremental Batch	Incremental (Batch or Streaming)
Use Cases	Best for smaller datasets	Ideal for thousands of files	Scale to millions+ of files per hour, backfills with billions of files
Syntax/Interface	<ul style="list-style-type: none"><li>• Python (spark.read)</li><li>• SQL (CTAS)</li></ul>	SQL	<ul style="list-style-type: none"><li>• Python (spark.readStream)</li><li>• SQL with Declarative Pipelines (CREATE OR REFRESH STREAMING TABLES)<ul style="list-style-type: none"><li>○ <a href="#">Use streaming tables in Databricks SQL</a></li></ul></li></ul>
Idempotency	No	Yes	Yes
Schema Evolution	Manual or inferred during read	Supported with options	Auto Loader automatically detects and evolves schemas. It supports loading data without predefined schemas and handles new columns as they appear.
Latency	High	Moderate (scheduled)	Low or high depending configuration
Easy of Use	Simple	Simple and SQL-based	Intermediate to advanced depending on the implementation (Python or SQL, incremental batch or streaming)
Summary	Best for one time, ad hoc ingestion. Can be scheduled to always read and process all data.	Simple and repeatable for incremental file ingestion. Great for schedule jobs or pipelines.	Best for near real-time streaming or incremental ingestion, with high automation and scalability.





## Cloud Storage with LakeFlow Connect Standard Connectors **DEMONSTRATION**

- **Data Ingestion with CREATE TABLE AS and COPY INTO**
- **Create Streaming Tables with SQL using Auto Loader**



Notebook: 02A - Data Ingestion with CREATE TABLE AS and COPY INTO

Notebook: 02B - Create Streaming Tables with SQL using Auto Loader





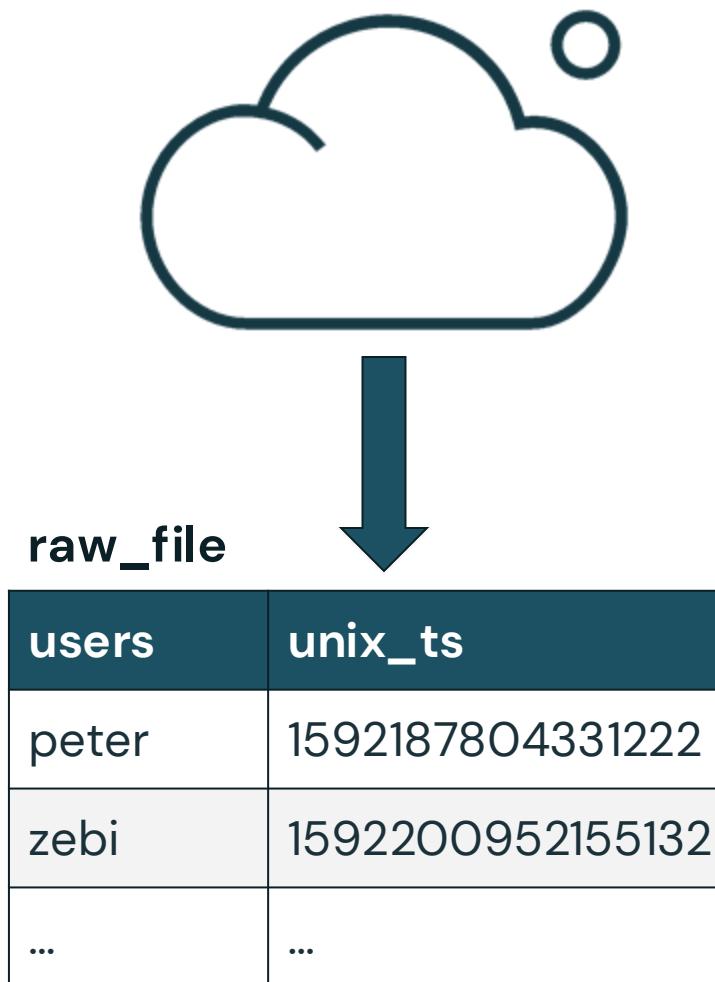
Cloud Storage with LakeFlow Connect  
Standard Connectors  
**LECTURE**

# Appending Metadata Columns on Ingest



# Appending Metadata Columns on Ingest

## Adding a Metadata Column



Create a Delta Table



Bronze Table

Bronze

users	unix_ts	last_mod_time	source
peter	1592187804331222	2024-10-01T18:04...	raw_file
zebi	1592200952155132	2024-10-01T18:04...	raw_file
...	...	...	...

Metadata columns

Add the last file modification time

Add the source file name



# Appending Metadata Columns on Ingest

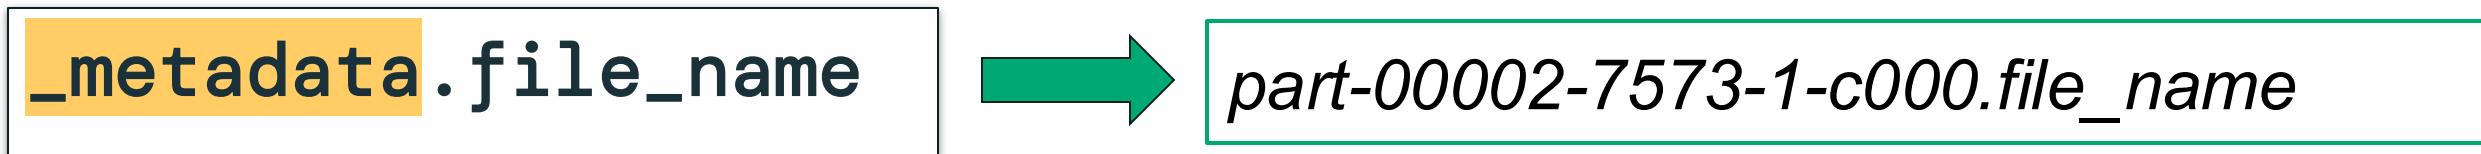
Common File Metadata Information From the Input Files



Add last file modification timestamp



Add input file name





Cloud Storage with LakeFlow Connect  
Standard Connectors  
**DEMONSTRATION**

# Adding Metadata Columns During Ingestion



Notebook: 03 – Adding Metadata Columns During Ingestion



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Cloud Storage with LakeFlow Connect  
Standard Connectors  
**LECTURE**

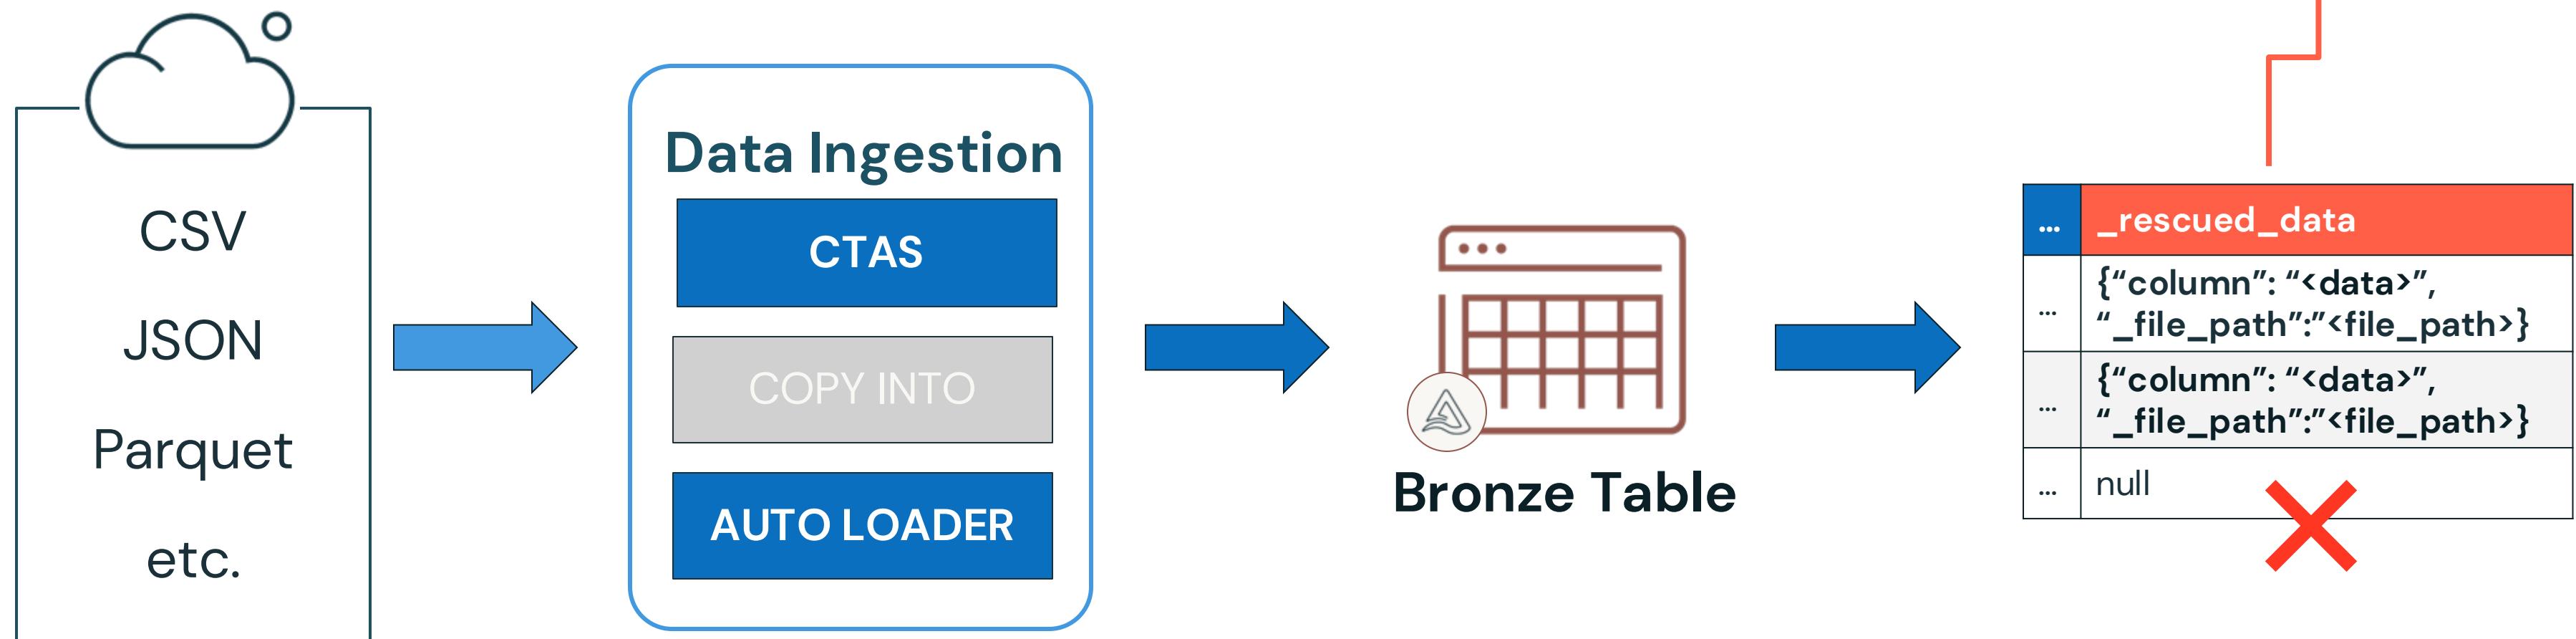
# Working with the Rescued Data Column



# Working with the Rescued Data Column

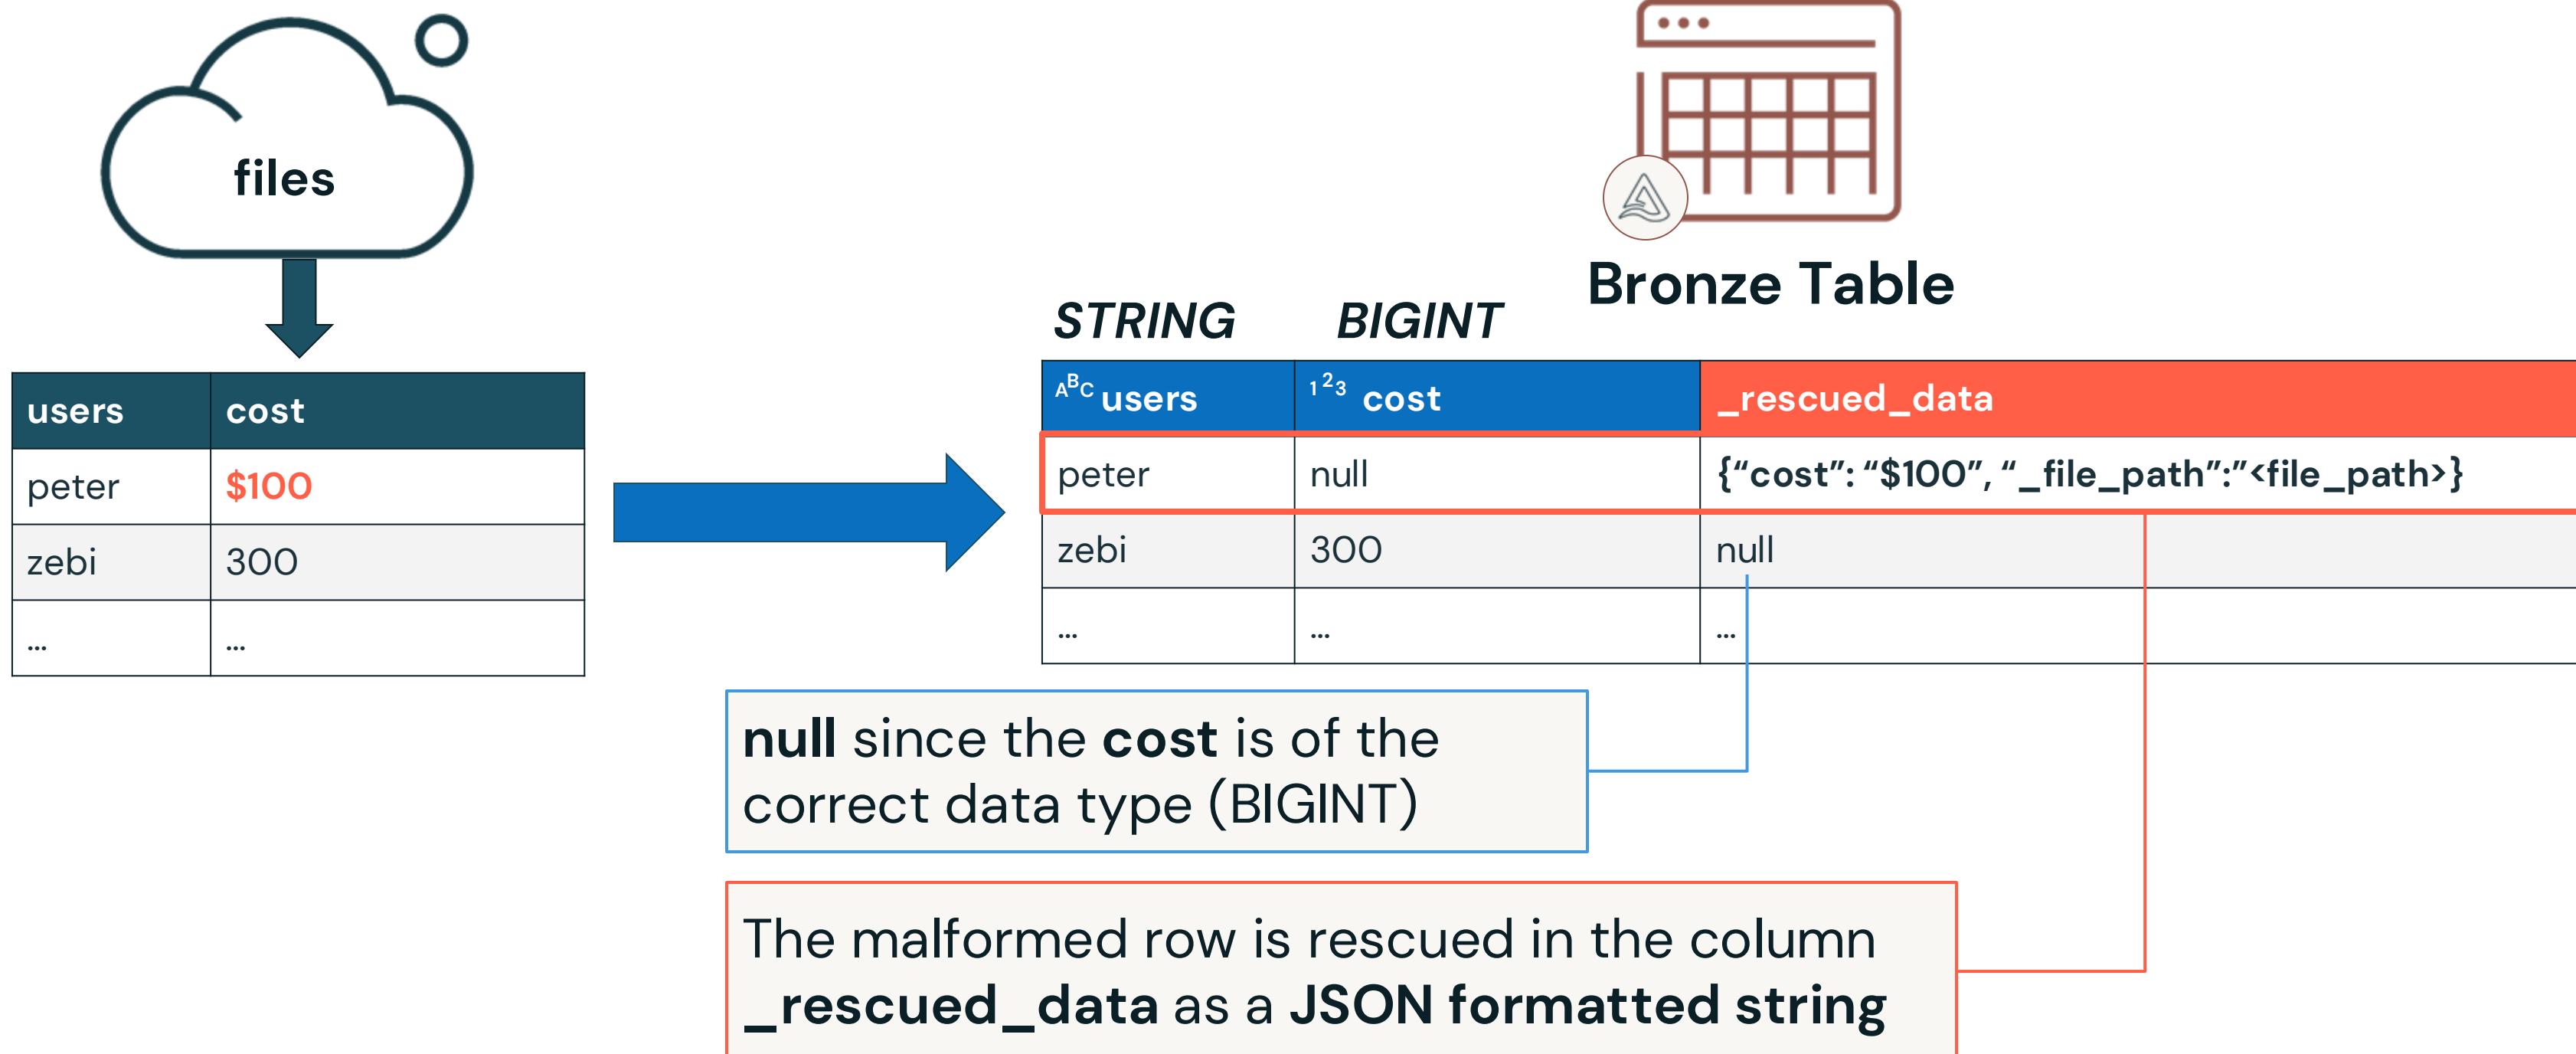
## Rescuing Malformed Rows on Ingestion

`read_files()`, `spark.read` or `Auto Loader` provides a **rescued data column** if the raw data does not match the schema



# Working with the Rescued Data Column

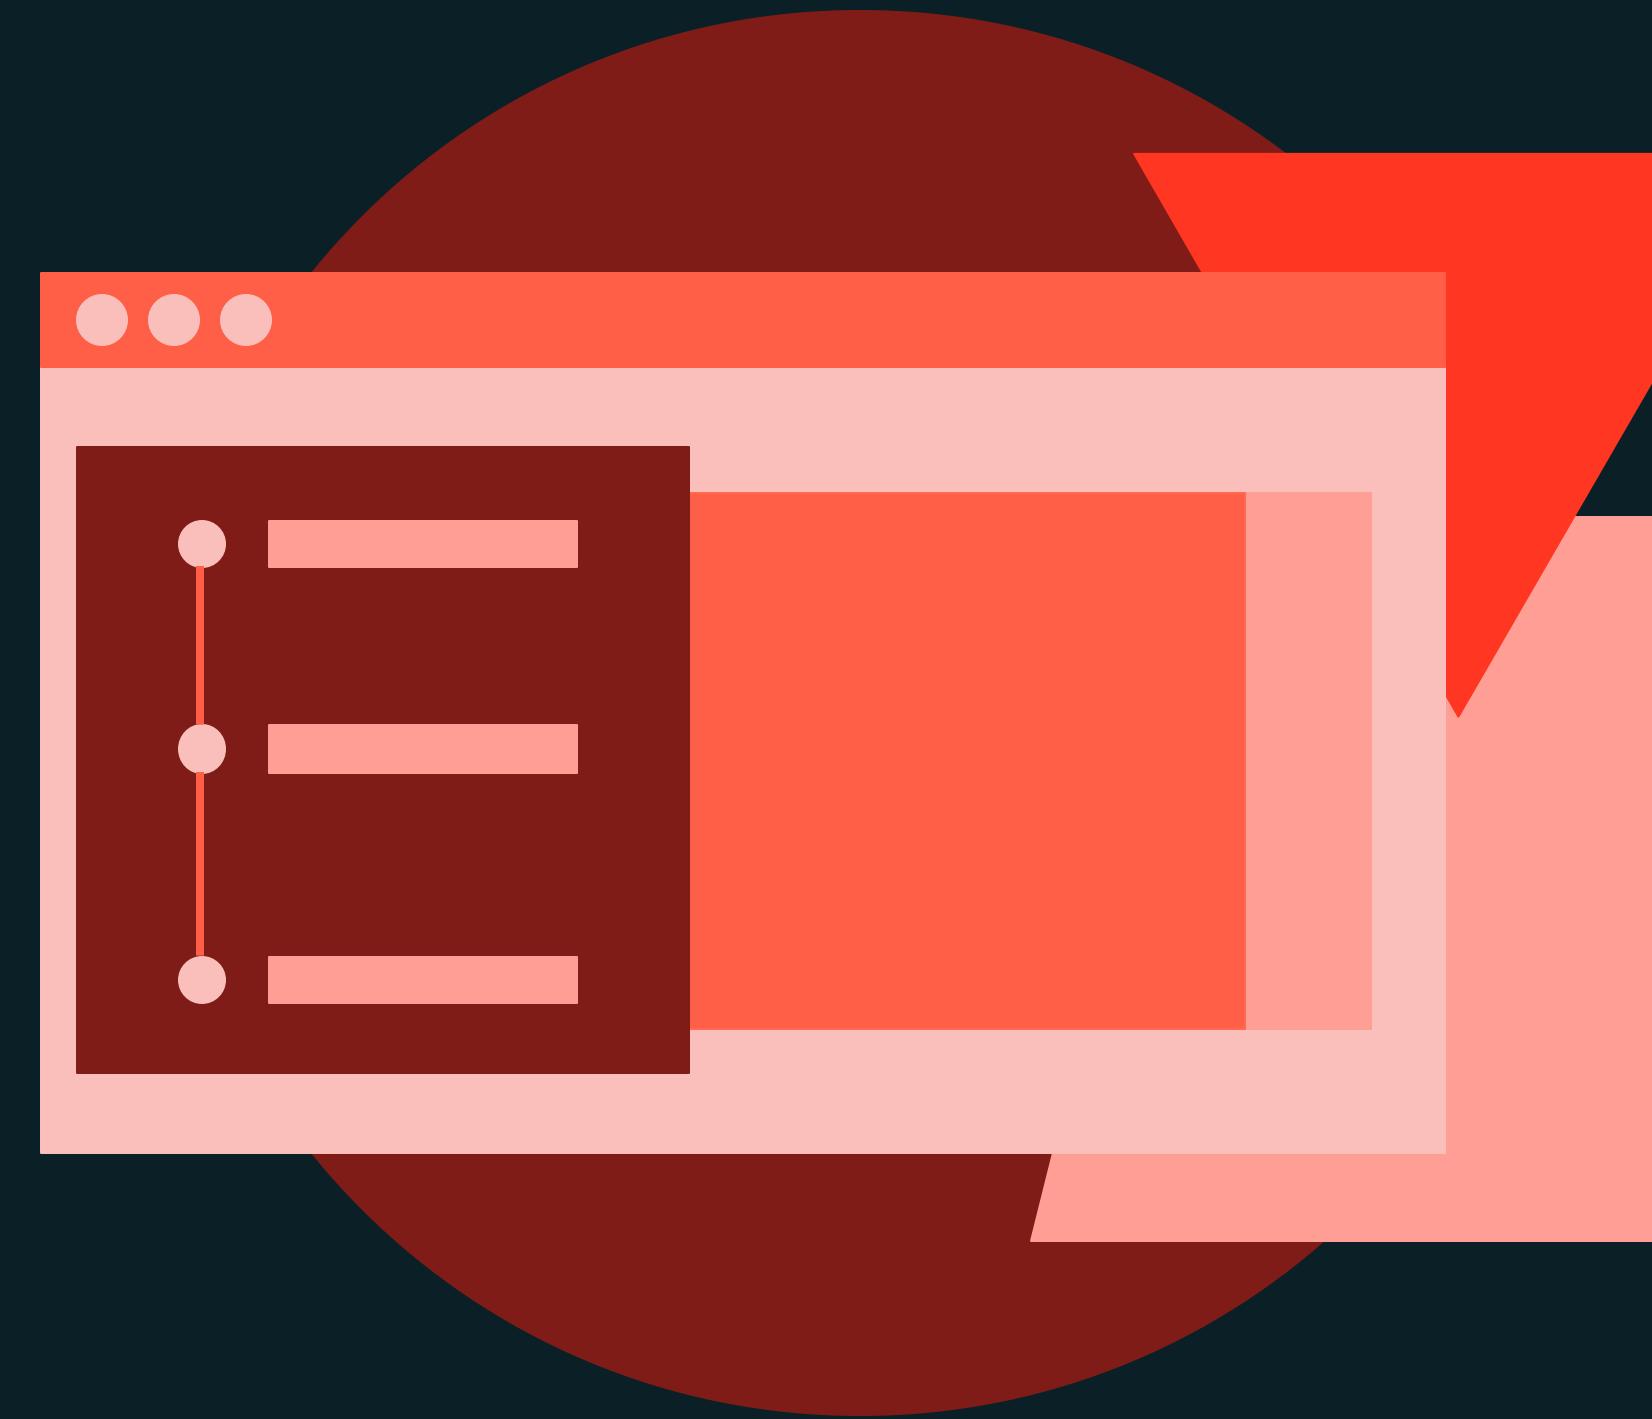
## Rescued Data Column Example





Cloud Storage with LakeFlow Connect  
Standard Connectors  
**DEMONSTRATION**

# Handling CSV Ingestion with the Rescued Data Column



Notebook: 04 – Handling CSV Ingestion with the Rescued Data Column



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Cloud Storage with LakeFlow Connect  
Standard Connectors  
**LAB EXERCISE**

# Creating Bronze Tables from CSV Files



Notebook: 05L - Creating Bronze Tables from CSV Files



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Cloud Storage with LakeFlow Connect  
Standard Connectors  
**LECTURE**

# Ingesting Semi- Structured Data: JSON



# Ingesting Semi-Structured Data: JSON

## JSON Overview

JSON objects are enclosed in brackets

The diagram illustrates a JSON object structure. A blue bracket on the left side of the slide points to the opening brace '{' at the top of the JSON code. Another blue bracket on the right side points to the closing brace '}' at the bottom. Within the JSON code, certain key-value pairs are highlighted in light blue: 'name': "John Doe", 'age': 35, 'address': { 'city': "Anytown", 'state': "CA"}, and the children array [{}]. The 'children' key is followed by a blue brace {}, which is itself followed by another blue brace {}, indicating that each child is an object. The entire JSON structure is enclosed in a grey rectangular box.

```
{  
  "name": "John Doe",  
  "age": 35,  
  "address": {  
    "city": "Anytown",  
    "state": "CA"  
  },  
  "children": [  
    {  
      "name": "Owen",  
      "age": 10  
    },  
    {  
      "name": "Eva",  
      "age": 8  
    }  
  ]  
}
```



# Ingesting Semi-Structured Data: JSON

## JSON Overview

Keys  
enclosed in  
quotation marks

```
{  
  "name": "John Doe",  
  "age": 35,  
  "address": {  
    "city": "Anytown",  
    "state": "CA"  
  },  
  "children": [  
    {  
      "name": "Owen",  
      "age": 10  
    },  
    {  
      "name": "Eva",  
      "age": 8  
    }  
  ]  
}
```

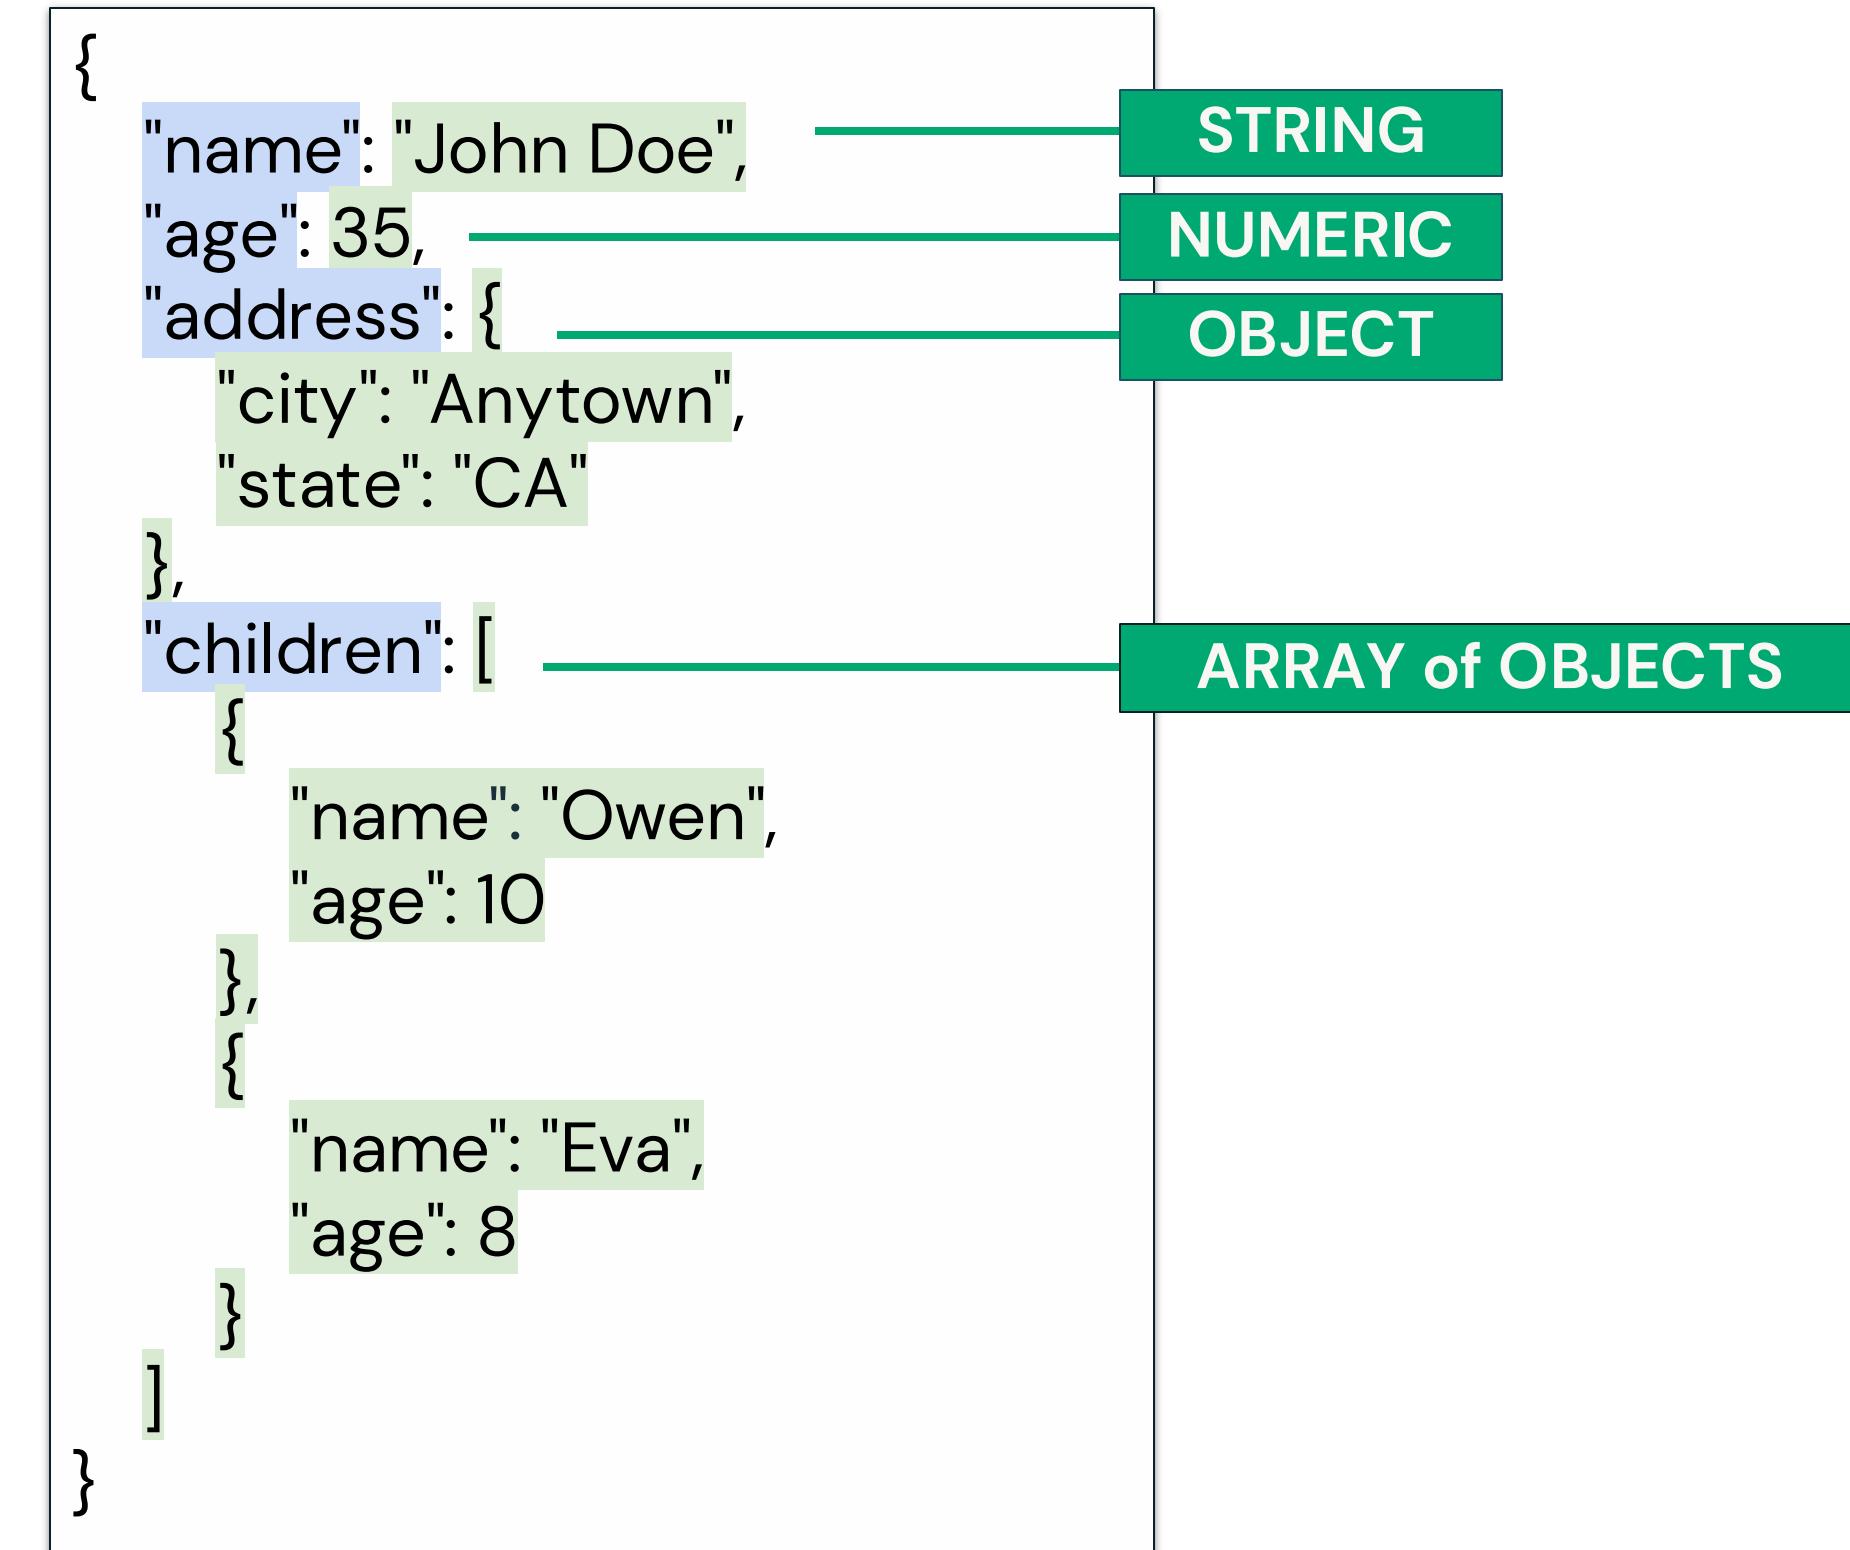


# Ingesting Semi-Structured Data: JSON

## JSON Overview

### Values

- String
- Number
- Boolean
- Array
- Object



# Ingesting Semi-Structured Data: JSON

## Working with a JSON-Formatted STRING Column

json_column
'{"name": "John Doe", "age": 35, "address": {"city": "Anytown", "state": "CA"}, "children": [{"name": "Owen", "age": 10}, {"name": "Eva", "age": 8}]}'
'{"name": "Kristi Doe", "age": 40, "address": {"city": "Anytown", "state": "CA"}, "children": [{"name": "Steve", "age": 10}]}'
...



Columns in tables can hold **JSON formatted strings** as values



# Ingesting Semi-Structured Data: JSON

## Working with a JSON-Formatted Column as a STRING

json_column
'{"name": "John Doe", "age": 35, "address": {"city": "Anytown", "state": "CA"}, "children": [{"name": "Owen", "age": 10}, {"name": "Eva", "age": 8}]}'
'{"name": "Kristi Doe", "age": 40, "address": {"city": "Anytown", "state": "CA"}, "children": [{"name": "Steve", "age": 10}]}'
...

*Working with JSON data can be done using different **column data types***

### 1. STRING Data Type

- JSON can be stored as a simple STRING
- Can hold **any JSON STRING** without constraints
- Less **performant**

Use : (colon) syntax to access subfields in JSON formatting strings

`SELECT json_column:name` → *John Doe*

`SELECT json_column:address:city` → *Anytown*



# Ingesting Semi-Structured Data: JSON

## Working with a JSON-Formatted Column as a STRING

json_column
'{"name": "John Doe", "age": 35, "address": {"city": "Anytown", "state": "CA"}, "children": [{"name": "Owen", "age": 10}, {"name": "Eva", "age": 8}]}'
'{"name": "Kristi Doe", "age": 40, "address": {"city": "Anytown", "state": "CA"}, "children": [{"name": "Steve", "age": 10}]}'
...

*Working with JSON data can be done using different **column data types***

### 1. STRING Data Type

- JSON can be stored as a simple STRING
- Can hold **any** JSON STRING without constraints
- Less performant

### 2. STRUCT Data Type

- You can parse JSON data into a STRUCT type, with a **defined schema**
- STRUCT **enforces** the JSON schema
- Is **more efficient** for querying than a JSON formatted STRING



# Ingesting Semi-Structured Data: JSON

Converting JSON Formatted Strings as STRUCTS

JSON String Types	Databricks SQL Data Type
String	<b>STRING</b>
Number	<b>INT/FLOAT/DOUBLE</b>
Boolean	<b>BOOLEAN</b>
Object	<b>STRUCT &lt;&gt;</b>
Array	<b>ARRAY &lt;&gt;</b>



# Ingesting Semi-Structured Data: JSON

## Converting JSON Formatted Strings as STRUCTS

```
{  
  "name": "John Doe",  
  "age": 35,  
  "address": {  
    "city": "Anytown",  
    "state": "CA"  
  },  
  "children": [  
    {  
      "name": "Owen",  
      "age": 10  
    },  
    {  
      "name": "Eva",  
      "age": 8  
    }  
  ]  
}
```



Define the **schema** of the  
**JSON formatted string**

```
STRUCT<  
  name: STRING,  
  age: INT,  
  address: STRUCT<  
    city: STRING,  
    state: STRING  
>,  
  children: ARRAY<  
    STRUCT<  
      name: STRING,  
      age: INT  
>  
>
```



# Ingesting Semi-Structured Data: JSON

## Converting JSON Formatted Strings as STRUCTS

```
{  
  "name": "John Doe",  
  "age": 35,  
  "address": {  
    "city": "Anytown",  
    "state": "CA"  
  },  
  "children": [  
    {  
      "name": "Owen",  
      "age": 10  
    },  
    {  
      "name": "Eva",  
      "age": 8  
    }  
  ]  
}
```



Specify the **STRUCT** data type to hold the JSON formatted string

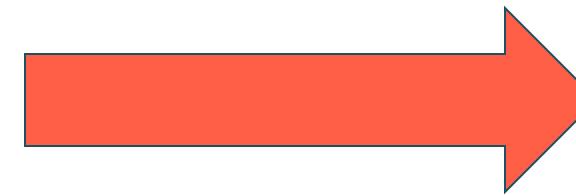
```
STRUCT<  
  name: STRING,  
  age: INT,  
  address: STRUCT<  
    city: STRING,  
    state: STRING  
  >,  
  children: ARRAY<  
    STRUCT<  
      name: STRING,  
      age: INT  
    >  
  >
```



# Ingesting Semi-Structured Data: JSON

## Converting JSON Formatted Strings as STRUCTS

```
{  
  "name": "John Doe",  
  "age": 35,  
  "address": {  
    "city": "Anytown",  
    "state": "CA"  
  },  
  "children": [  
    {  
      "name": "Owen",  
      "age": 10  
    },  
    {  
      "name": "Eva",  
      "age": 8  
    }  
  ]  
}
```



Specify the STRING and INT data types for the **name** and **age** keys

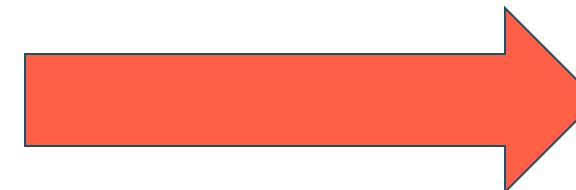
```
STRUCT<  
  name: STRING,  
  age: INT,  
  address: STRUCT<  
    city: STRING,  
    state: STRING  
>,  
  children: ARRAY<  
    STRUCT<  
      name: STRING,  
      age: INT  
>  
>
```



# Ingesting Semi-Structured Data: JSON

## Converting JSON Formatted Strings as STRUCTS

```
{  
  "name": "John Doe",  
  "age": 35,  
  "address": {  
    "city": "Anytown",  
    "state": "CA"  
  },  
  "children": [  
    {  
      "name": "Owen",  
      "age": 10  
    },  
    {  
      "name": "Eva",  
      "age": 8  
    }  
  ]  
}
```



The **address** key holds a **STRUCT** data type with the keys **city** and **state**

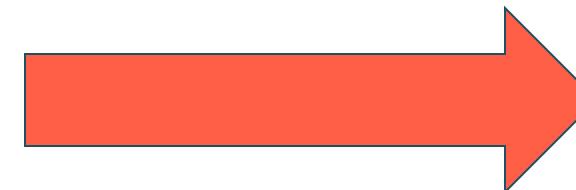
```
STRUCT<  
  name: STRING,  
  age: INT,  
  address: STRUCT<  
    city: STRING,  
    state: STRING  
>,  
  children: ARRAY<  
    STRUCT<  
      name: STRING,  
      age: INT  
>  
>
```



# Ingesting Semi-Structured Data: JSON

## Converting JSON Formatted Strings as STRUCTS

```
{  
  "name": "John Doe",  
  "age": 35,  
  "address": {  
    "city": "Anytown",  
    "state": "CA"  
  },  
  "children": [  
    {  
      "name": "Owen",  
      "age": 10  
    },  
    {  
      "name": "Eva",  
      "age": 8  
    }  
  ]  
}
```



The **children** key holds an  
**ARRAY** of **STRUCTS**

```
STRUCT<  
  name: STRING,  
  age: INT,  
  address: STRUCT<  
    city: STRING,  
    state: STRING  
>,  
  children: ARRAY<  
    STRUCT<  
      name: STRING,  
      age: INT  
>  
>
```



# Ingesting Semi-Structured Data: JSON

## Converting JSON Formatted Strings as STRUCTS

```
{  
  "name": "John Doe",  
  "age": 35,  
  "address": {  
    "city": "Anytown",  
    "state": "CA"  
  },  
  "children": [  
    {  
      "name": "Owen",  
      "age": 10  
    },  
    {  
      "name": "Eva",  
      "age": 8  
    }  
  ]  
}
```

We can **derive the schema** of the JSON-formatted STRING column with `schema_of_json`

1

```
SELECT schema_of_json('sample-json-string')
```

```
STRUCT<  
  name: STRING,  
  age: INT,  
  address: STRUCT<  
    city: STRING,  
    state: STRING  
>,  
  children: ARRAY<  
    STRUCT<  
      name: STRING,  
      age: INT  
>  
>
```

The function returns the **structure** of the JSON formatted string



# Ingesting Semi-Structured Data: JSON

## Converting JSON Formatted Strings as STRUCTS

```
STRUCT<
  name: STRING,
  age: INT,
  address: STRUCT<
    city: STRING,
    state: STRING
  >,
  children: ARRAY<
    STRUCT<
      name: STRING,
      age: INT
    >
  >
>
```

2

```
SELECT from_json(json_col, 'json-struct-schema') AS struct_column
FROM table
```

The `from_json` function returns a **struct column** using the JSON string and specified schema



# Ingesting Semi-Structured Data: JSON

## Working with a JSON-Formatted Column as a STRING

json_column
'{"name": "John Doe", "age": 35, "address": {"city": "Anytown", "state": "CA"}, "children": [{"name": "Owen", "age": 10}, {"name": "Eva", "age": 8}]}'
'{"name": "Kristi Doe", "age": 40, "address": {"city": "Anytown", "state": "CA"}, "children": [{"name": "Steve", "age": 10}]}'
...

*Working with JSON data can be done using different **column data types***

### 1. STRING Data Type

- JSON can be stored as a simple STRING
- Can hold **any JSON STRING** without constraints
- Less performant

### 2. STRUCT Data Type

- You can parse JSON data into a STRUCT type, with a **defined schema**
- STRUCT **enforces** the JSON schema
- Is **more efficient** for querying than a JSON formatted STRING

### 3. VARIANT Data Type

- Can store **any type of data**, including JSON, and is **ideal for semi-structured data**
- Highly **flexible**
- Improved **performance** over existing methods

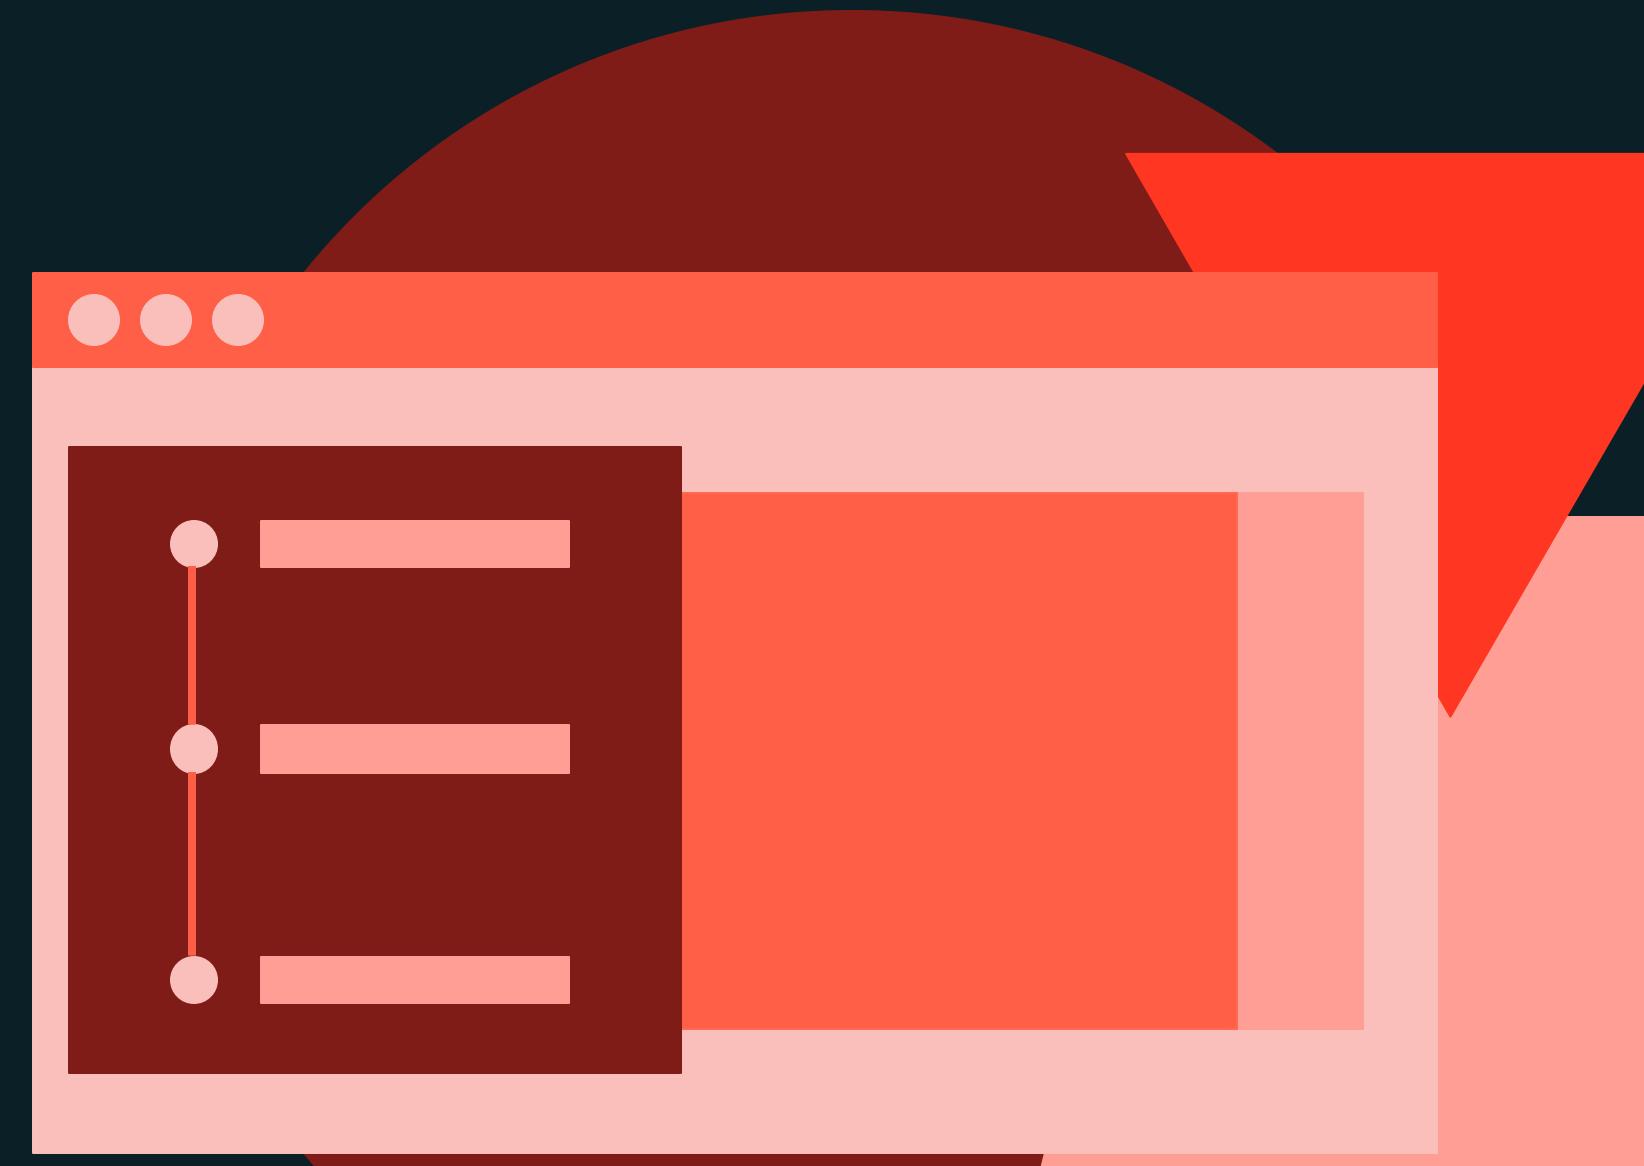
Public Preview as of 2025 Q2





Cloud Storage with LakeFlow Connect  
Standard Connectors  
**DEMONSTRATION**

# Ingesting JSON files with Databricks



Notebook: 06 – Ingesting JSON Files with Databricks



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Cloud Storage with LakeFlow Connect  
Standard Connectors  
**LAB EXERCISE**

# Creating Bronze Tables from JSON Files



Notebook: 07L - Creating Bronze Table from JSON Files



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



# Enterprise Data with LakeFlow Connect Managed Connectors



## Data Ingestion with Lakeflow Connect



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Agenda

## Section Overview – Enterprise Data with LakeFlow Connect Managed Connectors

- **Ingesting Enterprise Data into Databricks Overview**
- **Enterprise Data Ingestion with Lakeflow Connect**

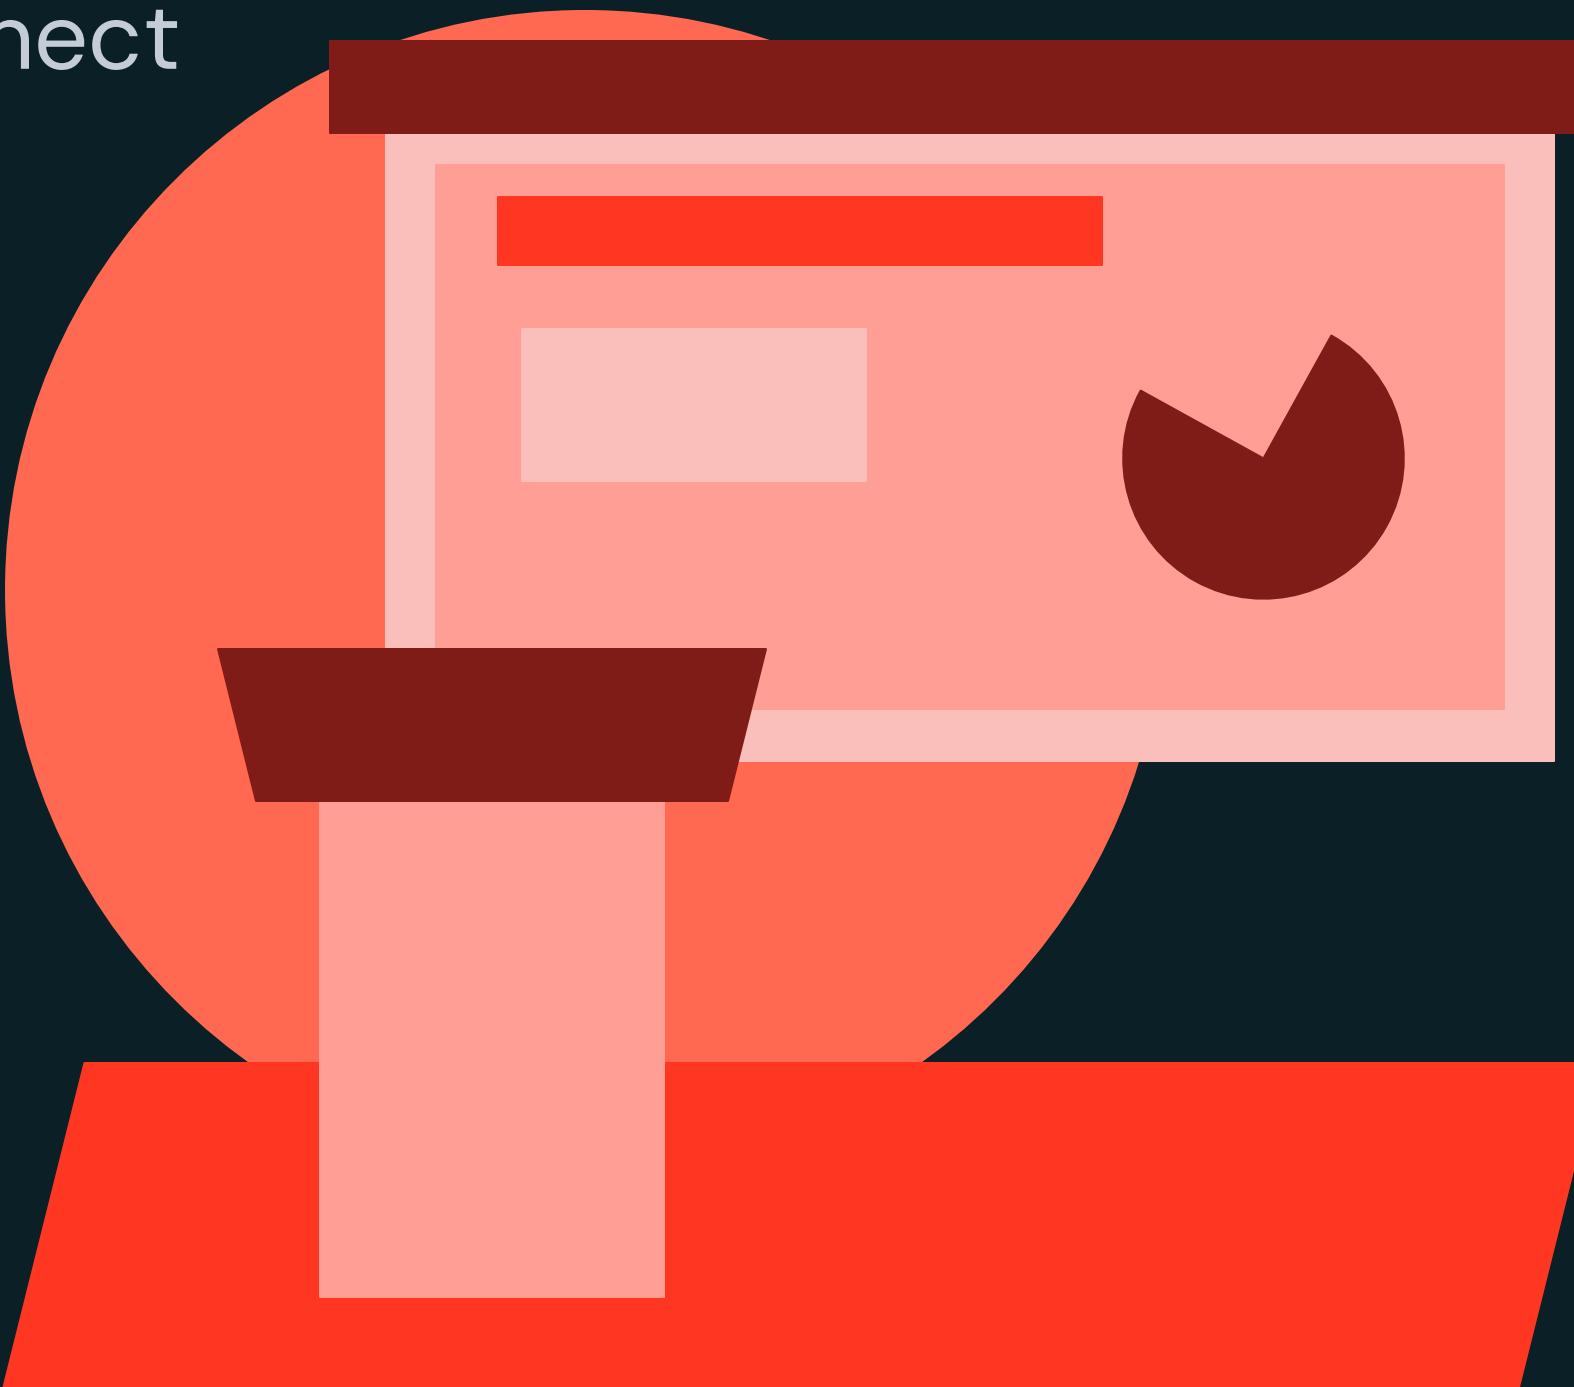




Ingesting Enterprise Data with LakeFlow Connect

LECTURE

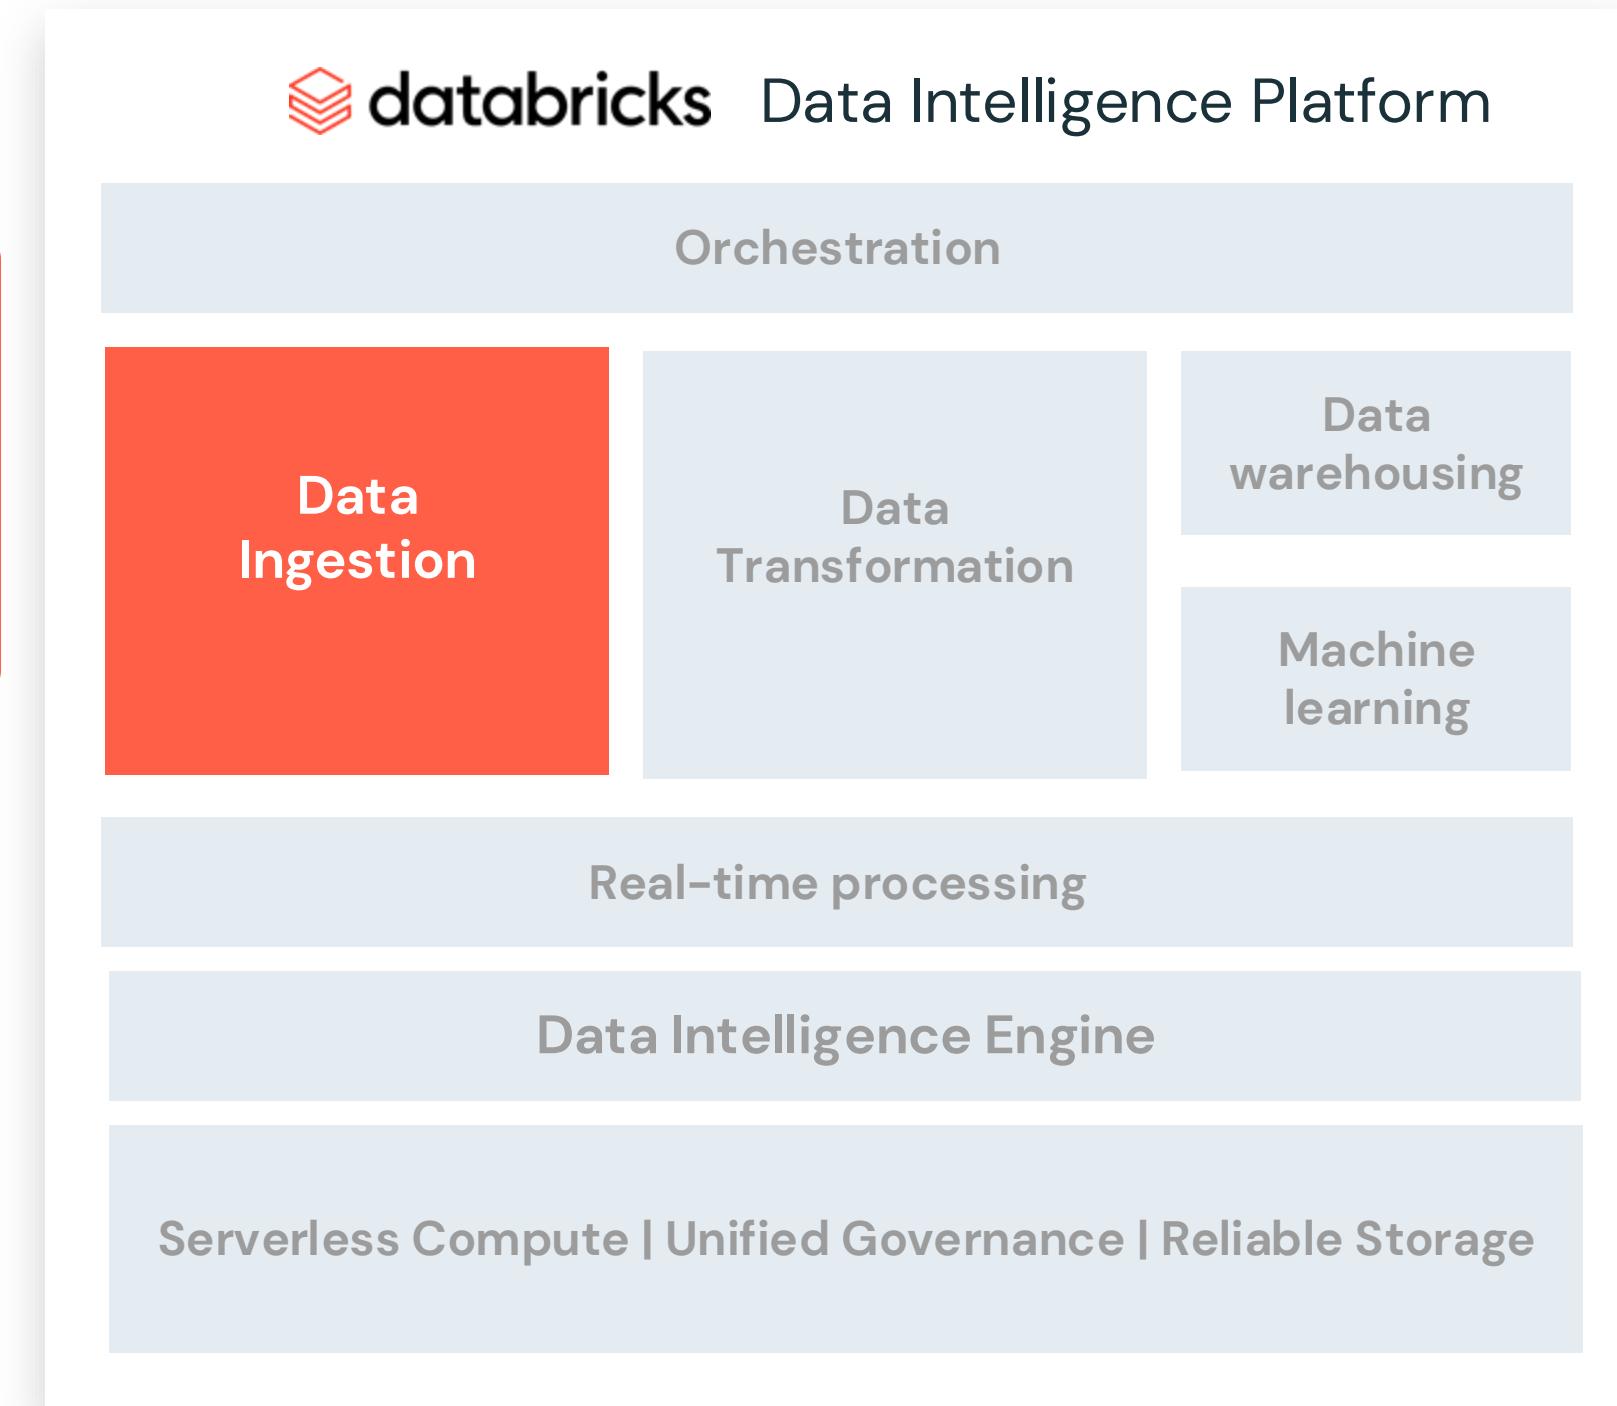
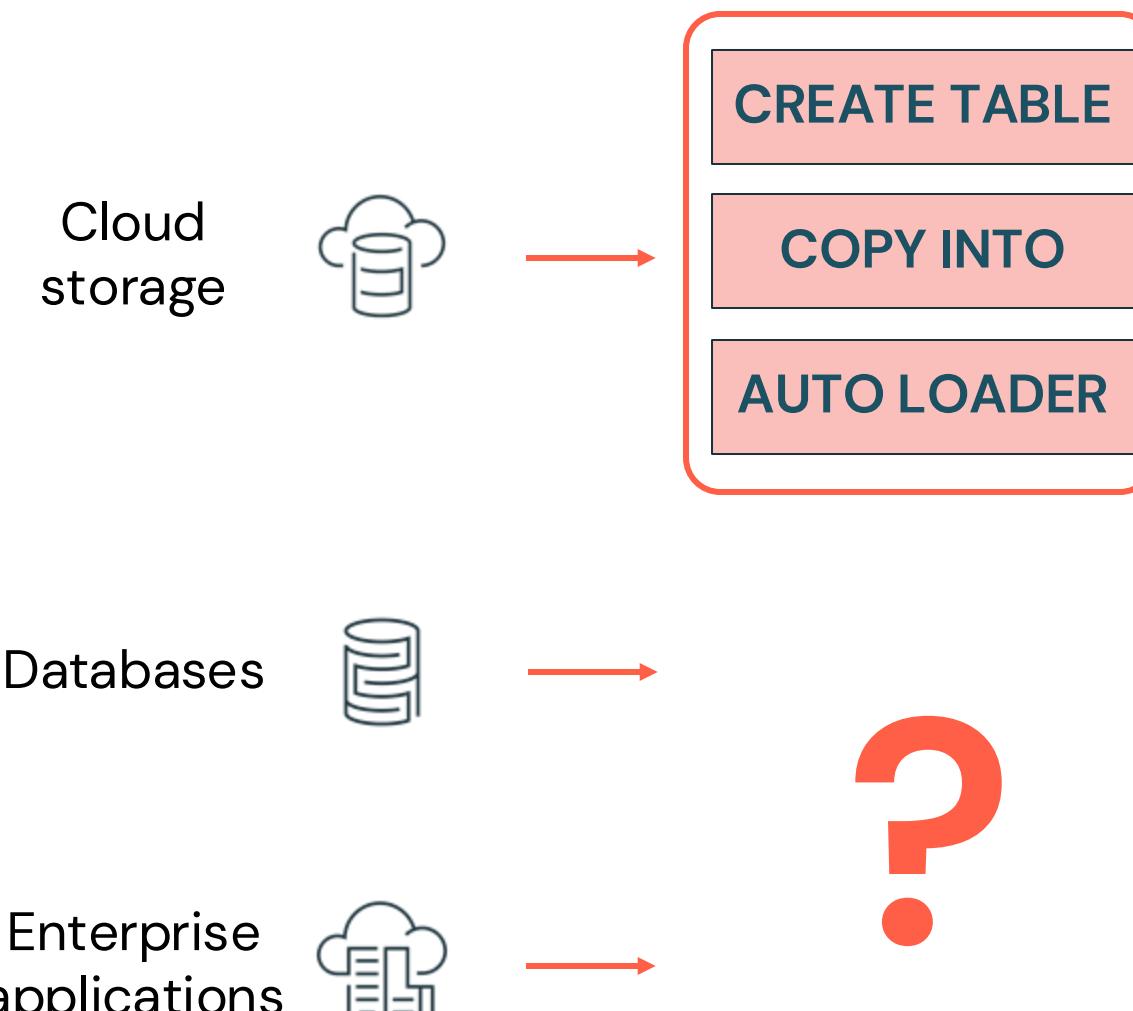
# Ingesting Enterprise Data Overview



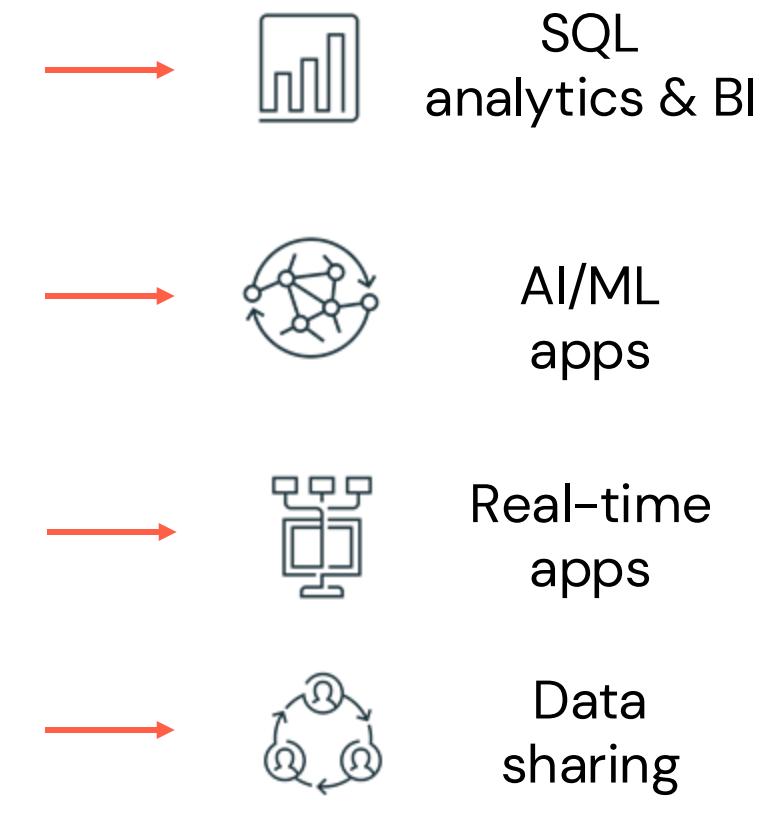
# Ingesting Enterprise Data Overview

## Data Ingestion to Databricks Overview

### Data sources



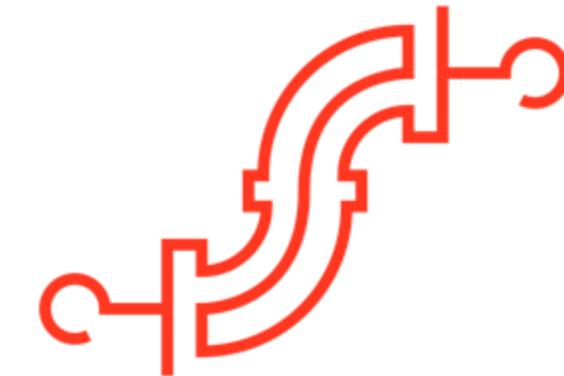
### Use cases



# Ingesting Enterprise Data Overview

## Lakeflow Connect Managed Connectors

Lakeflow Connect  
Managed Connectors



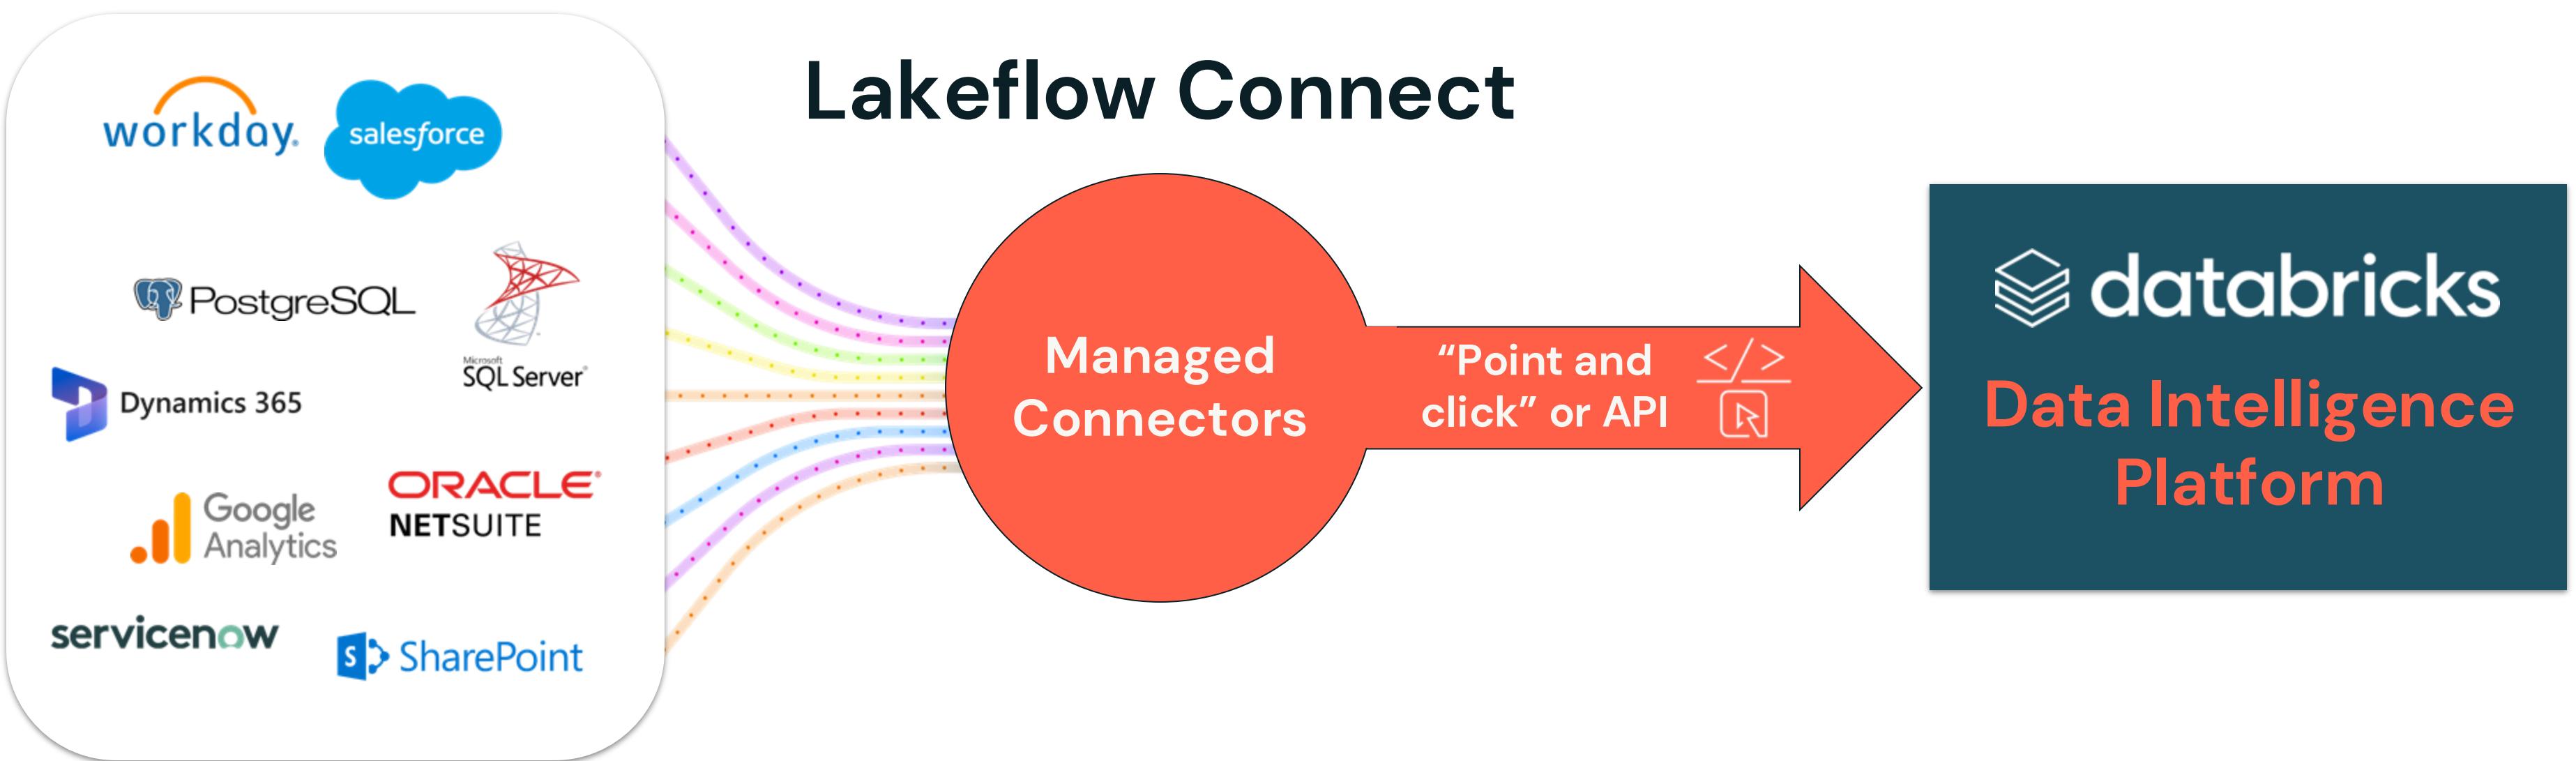
### BENEFITS

- Simplify the process of ingestion data from a wide variety of enterprise databases and applications
- Provide an easy to use **user interface (UI)** (or you can use the API)
- Fully managed by **Databricks**, reducing the need for manual configuration or custom code



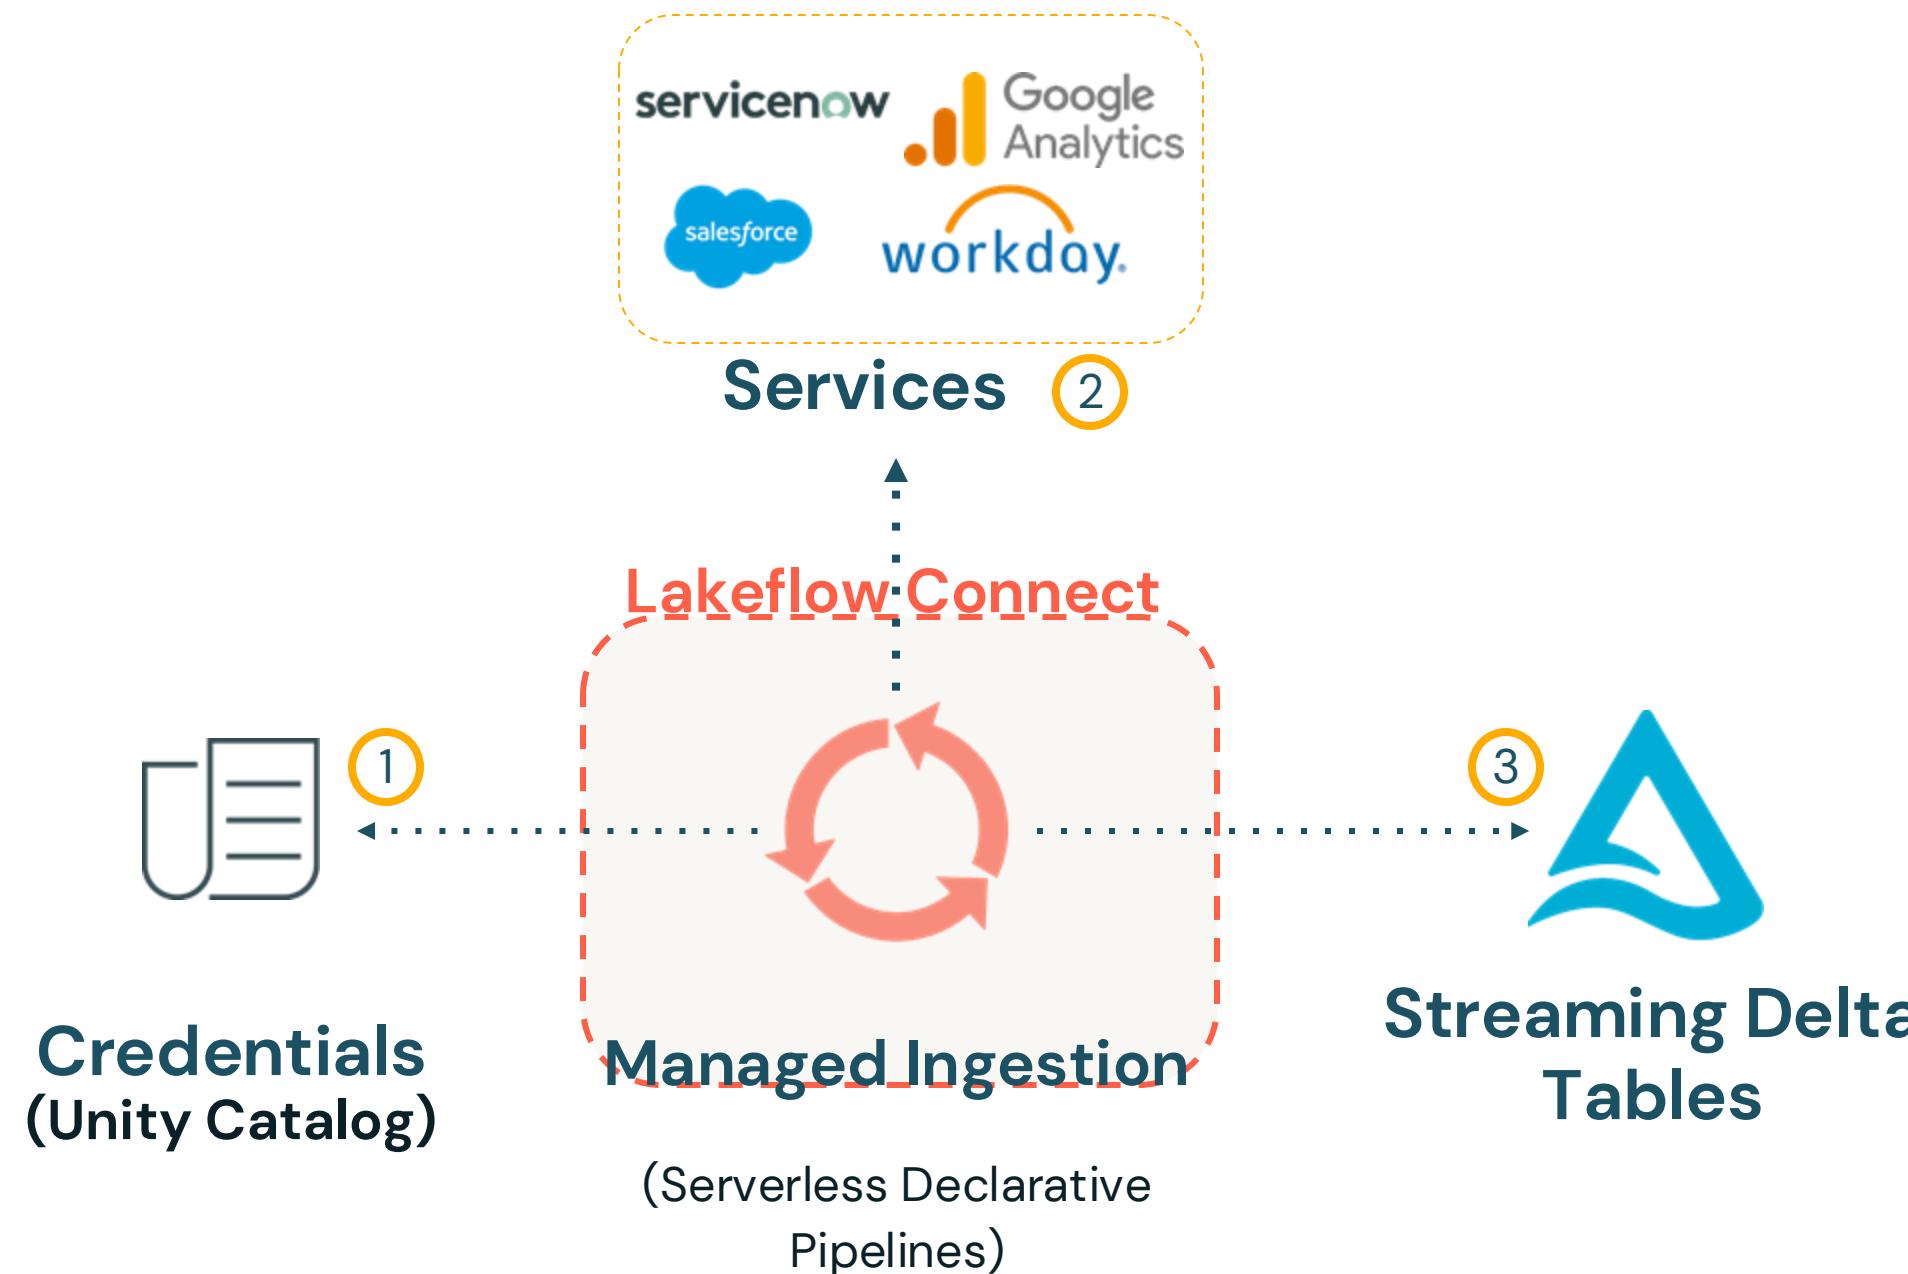
# Ingesting Enterprise Data Overview

## Data Ingestion with Lakeflow Connect Managed Connectors



# Ingesting Enterprise Data Overview

## Lakeflow Connect Managed Connectors: SaaS Ingestion



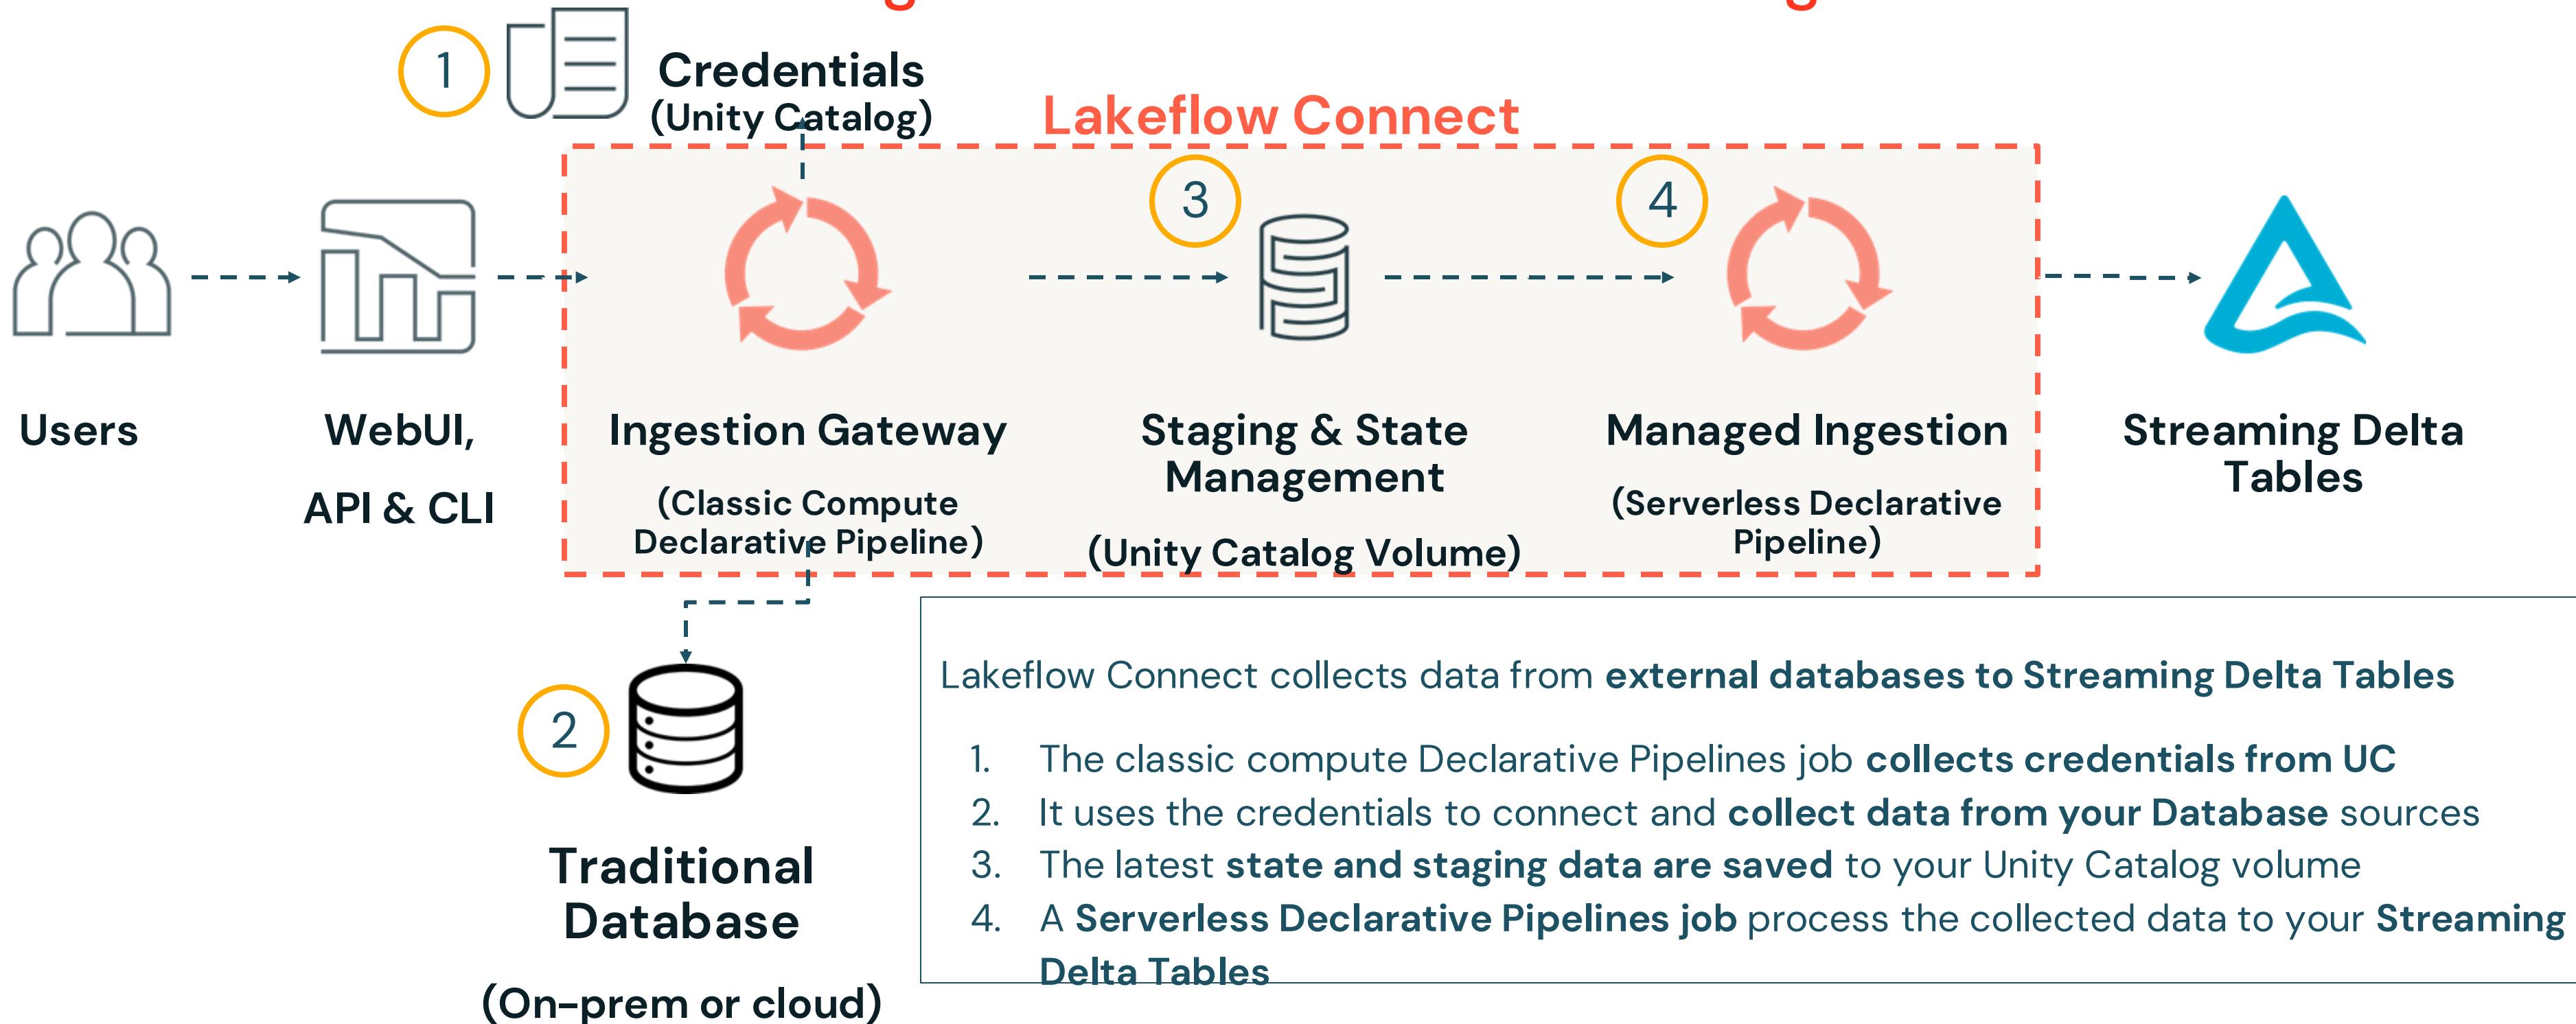
Lakeflow Connect collects data from **external sources** to **Streaming Delta Tables** using a **serverless compute Declarative Pipelines pipeline**:

1. A Lakeflow Serverless Declarative Pipelines job **collects credentials** from Unity Catalog.
2. The job **reaches out** to the publicly accessible data source (e.g., API, open OLAP port, etc.).
3. The service transforms the data and stores it to a **Streaming Delta Table**.



# Ingesting Enterprise Data Overview

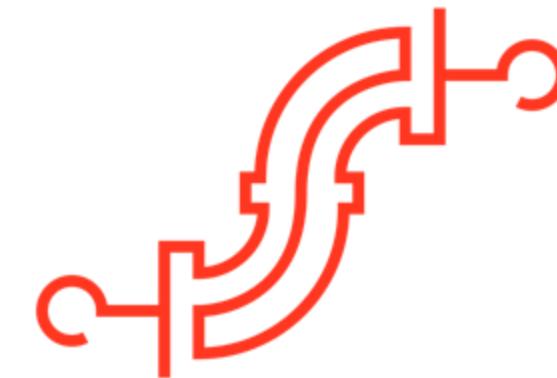
## Lakeflow Connect Managed Connectors: Database Ingestion



# Ingesting Enterprise Data Overview

## Data Ingestion with Partner Connect

Lakeflow Connect  
Managed Connectors



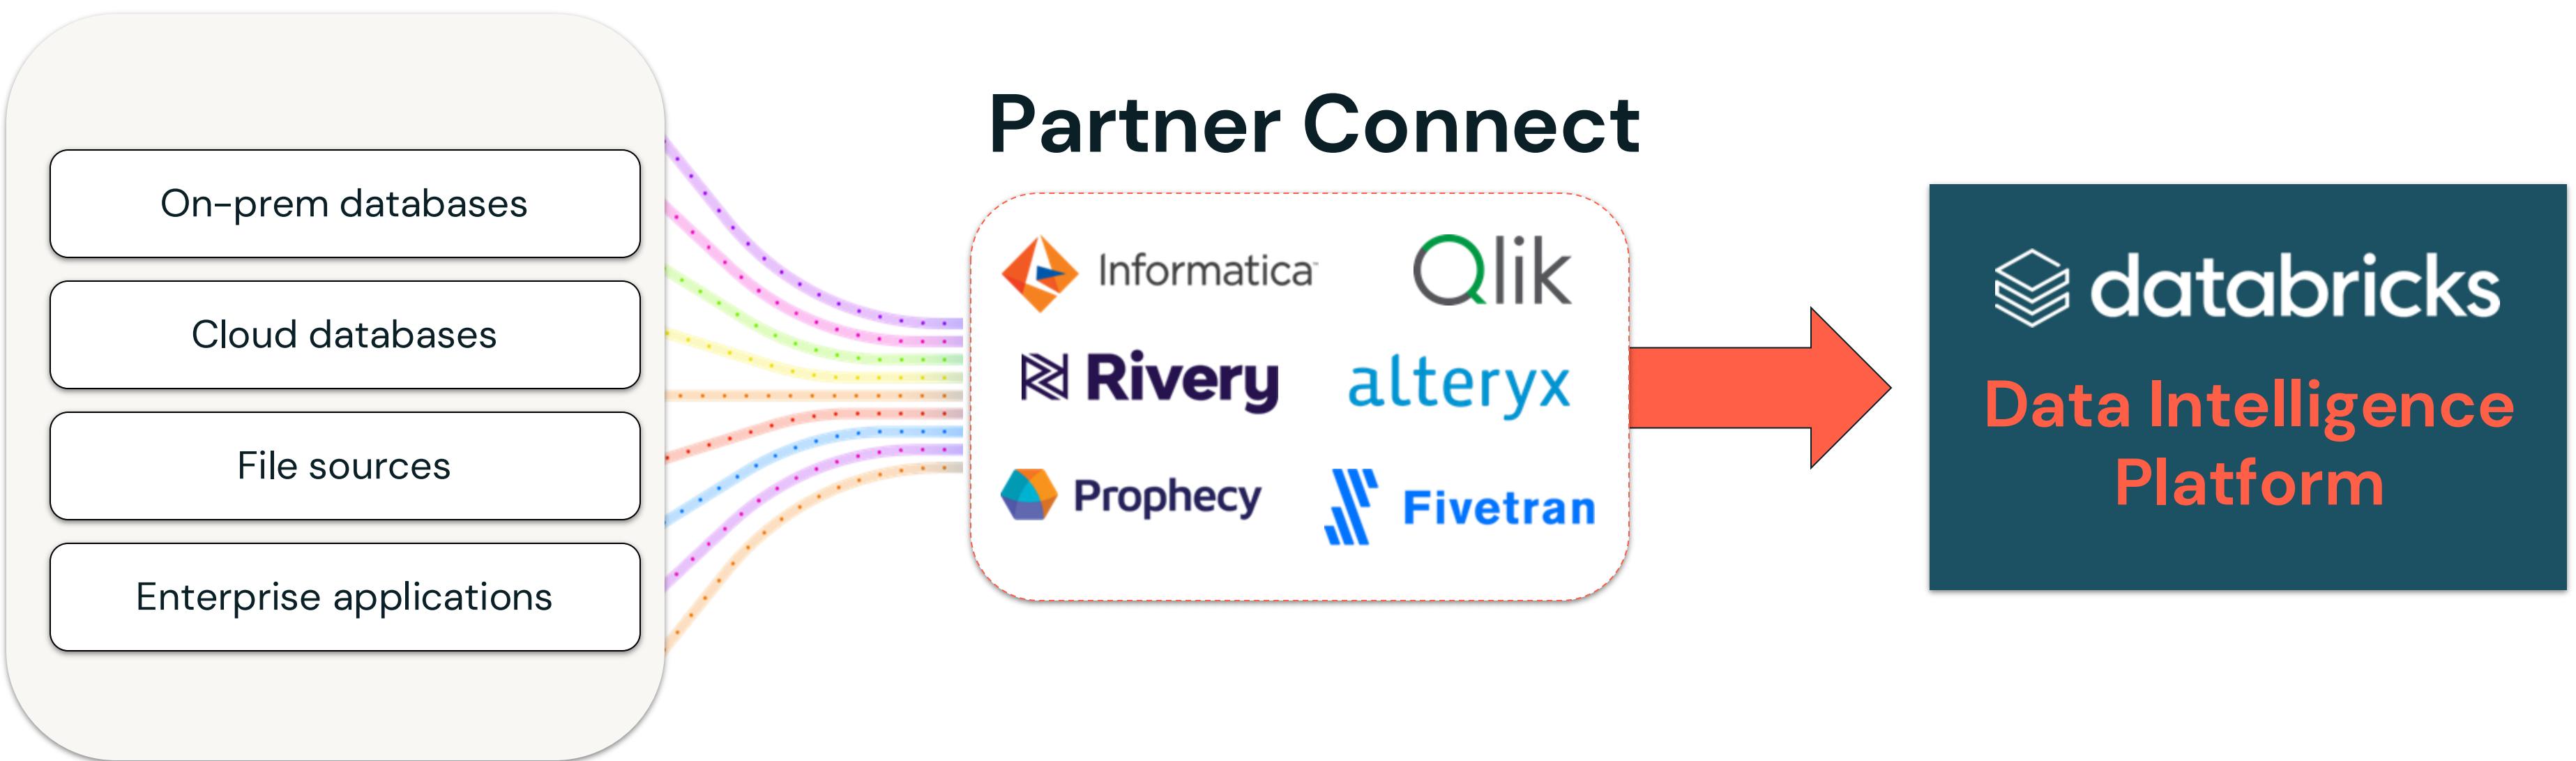
 **databricks**

### Partner Connect



# Ingesting Enterprise Data Overview

## Data Ingestion with Partner Connect



Utilize a rich ecosystem of partner solutions for data ingestion

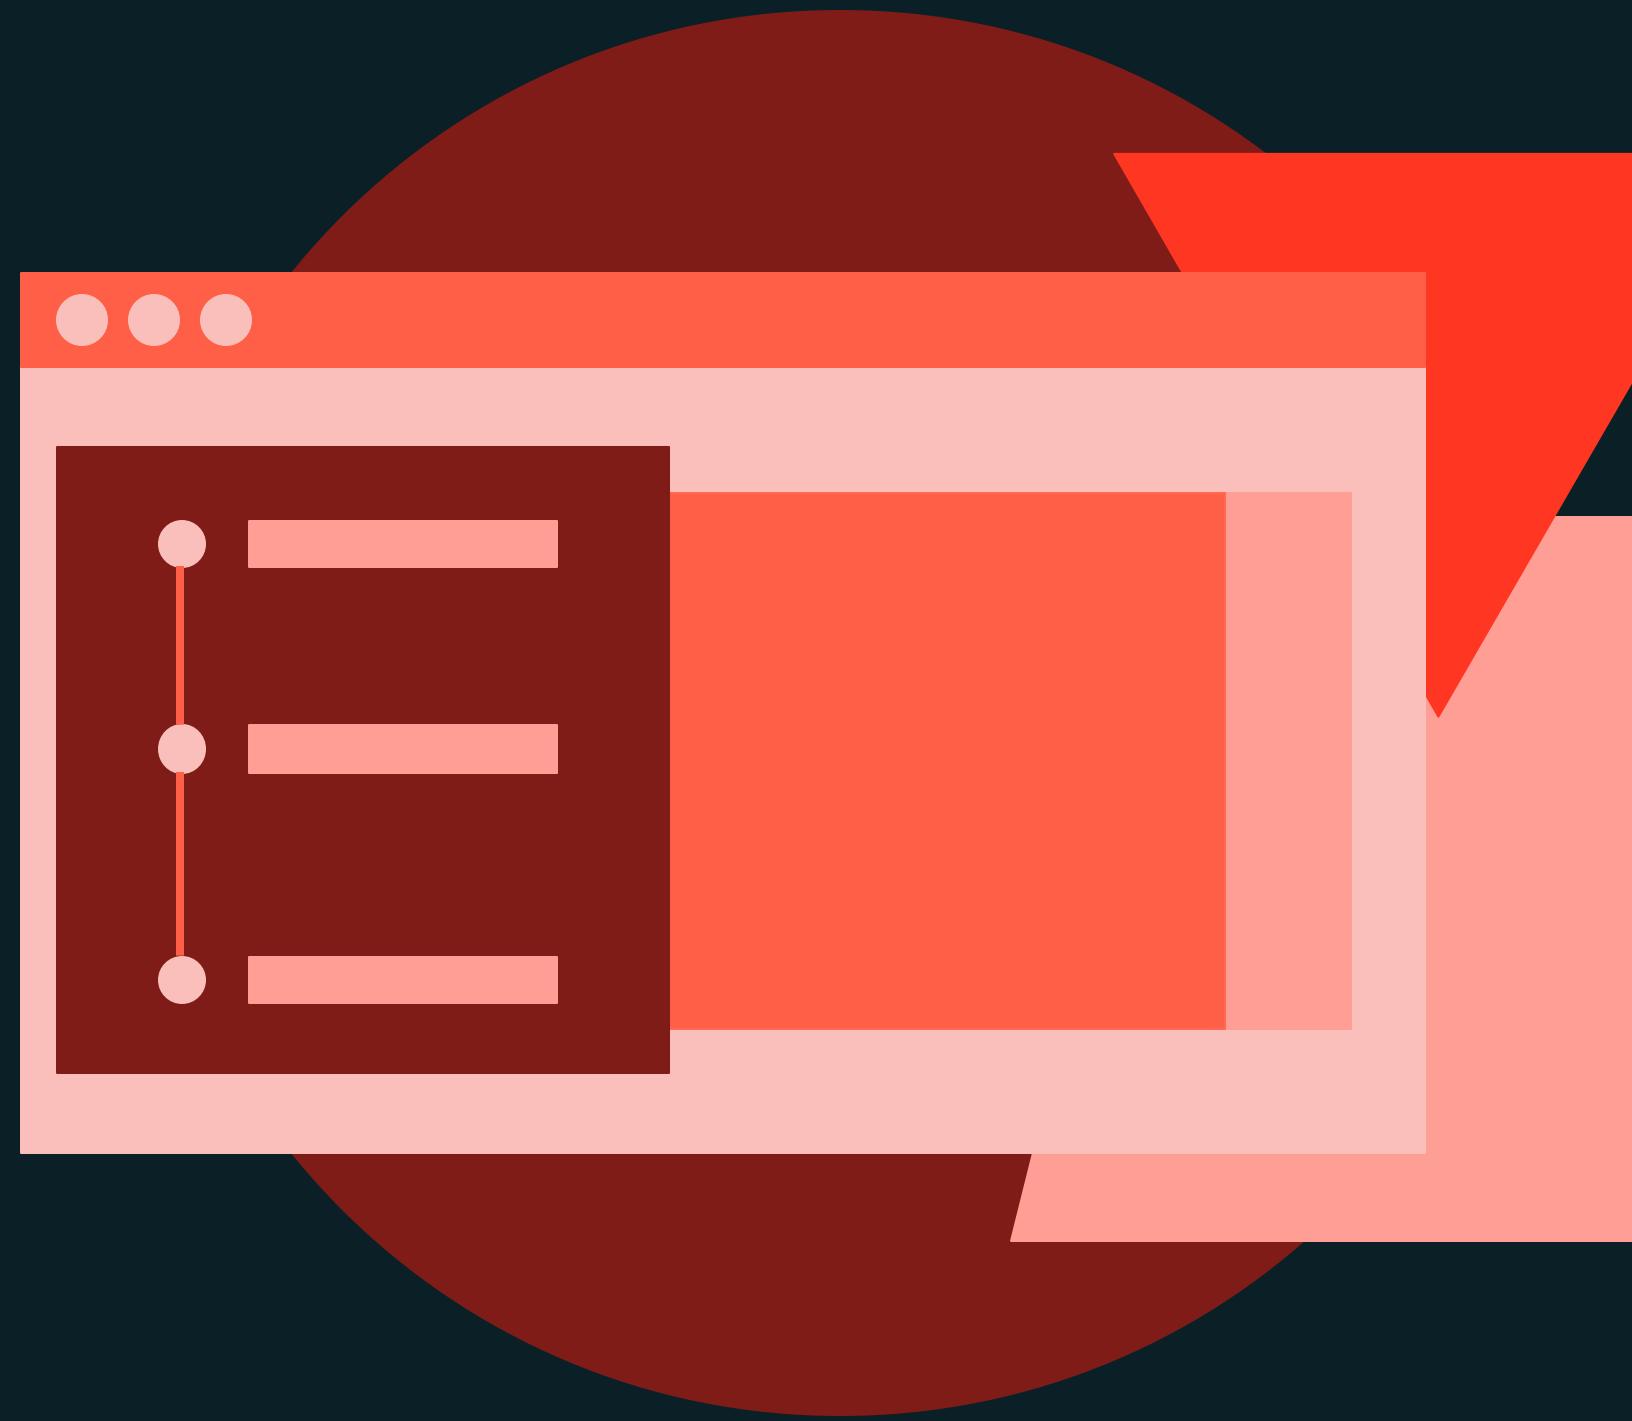




Data Ingestion and Transformation

## DEMONSTRATION

# Enterprise Data Ingestion with Lakeflow Connect



Notebook: 08 – Enterprise Data Ingestion with Lakeflow Connect



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



# Ingestion Alternatives (BONUS)



## Data Ingestion with Lakeflow Connect



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Agenda

## Section Overview – Ingestion Alternatives

- Additional Ingestion Features Overview
- Ingesting into Existing Delta Tables
- Data Ingestion with MERGE INTO





Building Pipelines

LECTURE

# Additional Features Overview



# Additional Features Overview

## What's Next: Features Outside This Course

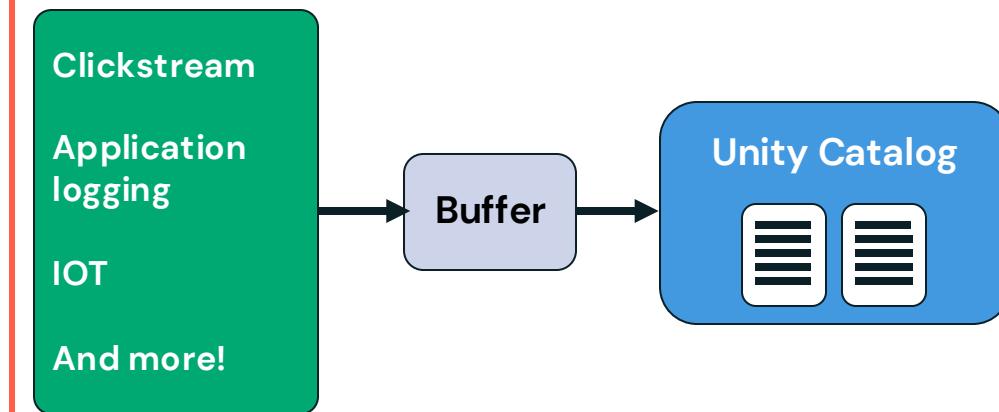


### Lakehouse Federation

Allows you to **query external data sources** without moving your data

Especially useful for:

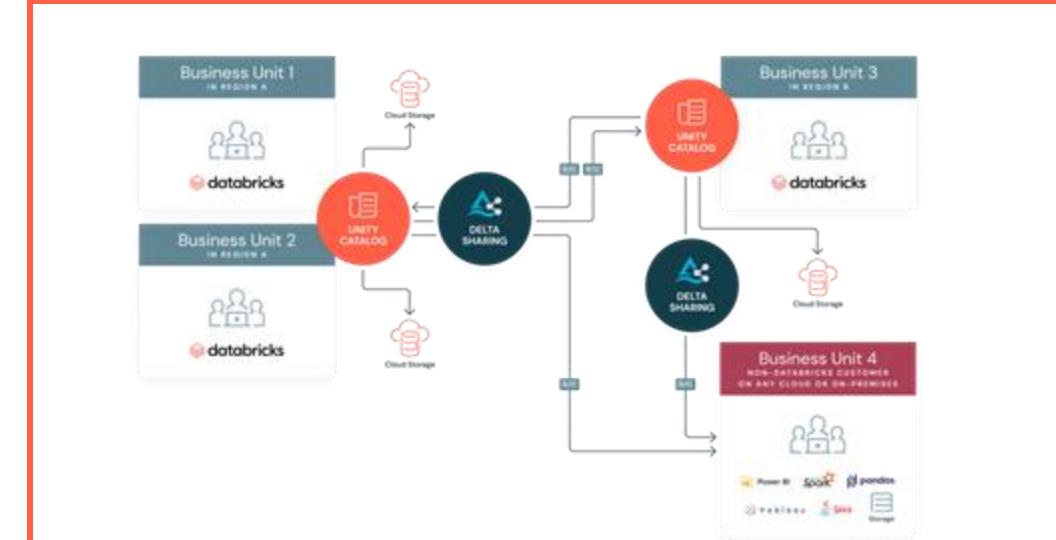
- Ad hoc reporting
- Proof-of-concept work
- The exploratory phase of new ETL pipelines or reports
- Supporting workloads during incremental migration



### Zerobus (coming soon)

A **Lakeflow Connect API** that allows developers to **write event data directly to their lakehouse** at very high throughput (100 MB/s) with near real-time latency (<5 seconds)

**Simplify ingestion** for IOT, clickstreams, telemetry, and more.



### Delta Sharing

Allows you to securely share data across platforms, clouds, and regions



# Additional Ingestion Features Overview

## Ingesting Data with Databricks Marketplace



**Open exchange for all data products**

- **Datasets**
- **Notebooks**
- **Dashboards**
- **ML models**
- **Solutions Accelerators**

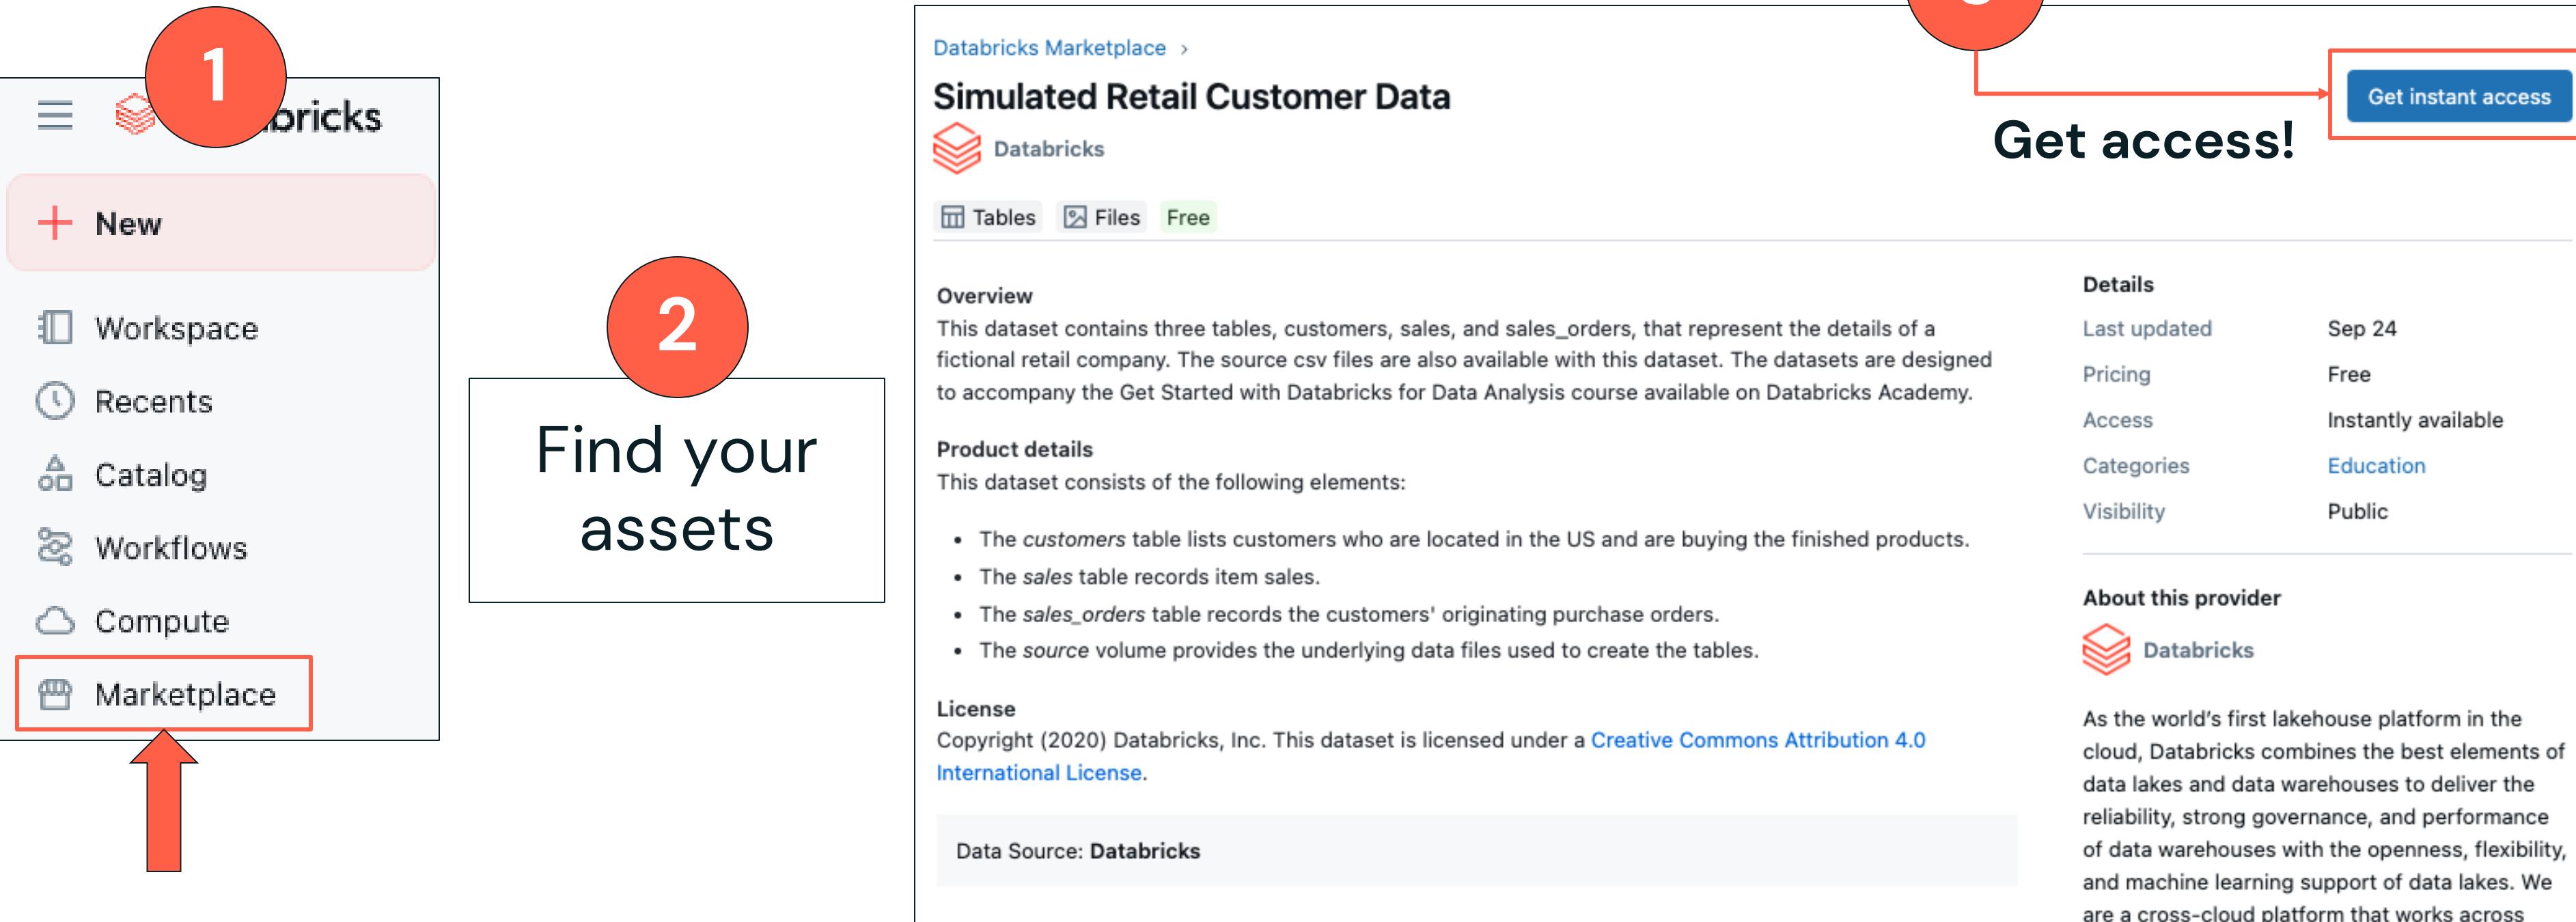


**Powered by**  **Delta Sharing**



# Ingesting Data with Databricks Marketplace

## Delta Sharing with Databricks Marketplace





Ingestion Alternatives

LECTURE

# Ingesting into Existing Delta Tables



# Ingesting into Existing Delta Tables

## Updates, Inserts, and Deletes on Delta Tables with MERGE INTO

- Merges a set of **updates, insertions, and deletions** from a **source** table into a **target** Delta table.
- MERGE INTO supports **schema enforcement** or **schema evolution** and allows different actions depending on whether a row is matched between a source and target table:
  - Matched rows: **UPDATE** or **DELETE**
  - Unmatched rows by target: **INSERT**
  - Unmatched rows by source: **UPDATE** or **DELETE**
- This functionality makes MERGE INTO ideal for handling **slowly changing dimensions (SCDs)**, **incremental loads**, and complex **change data capture (CDC)** scenarios.



# Ingesting into Existing Delta Tables

Updates, Inserts, and Deletes on Delta Tables with MERGE INTO

**target\_table**

users	email	status	
peter	peter@email.com	current	
zebi	zebi@email.com	current	
...	...	...	



**source\_table**

users	email	status	
peter	peter@email.com	delete	
zebi	zebi@other.com	update	
samarth	samarth@other.com	new	

Update **target\_table** with the  
**source\_table**



**target\_table**

users	email	status	
zebi	zebi@other.com	update	
samarth	samarth@other.com	new	
...	...	...	



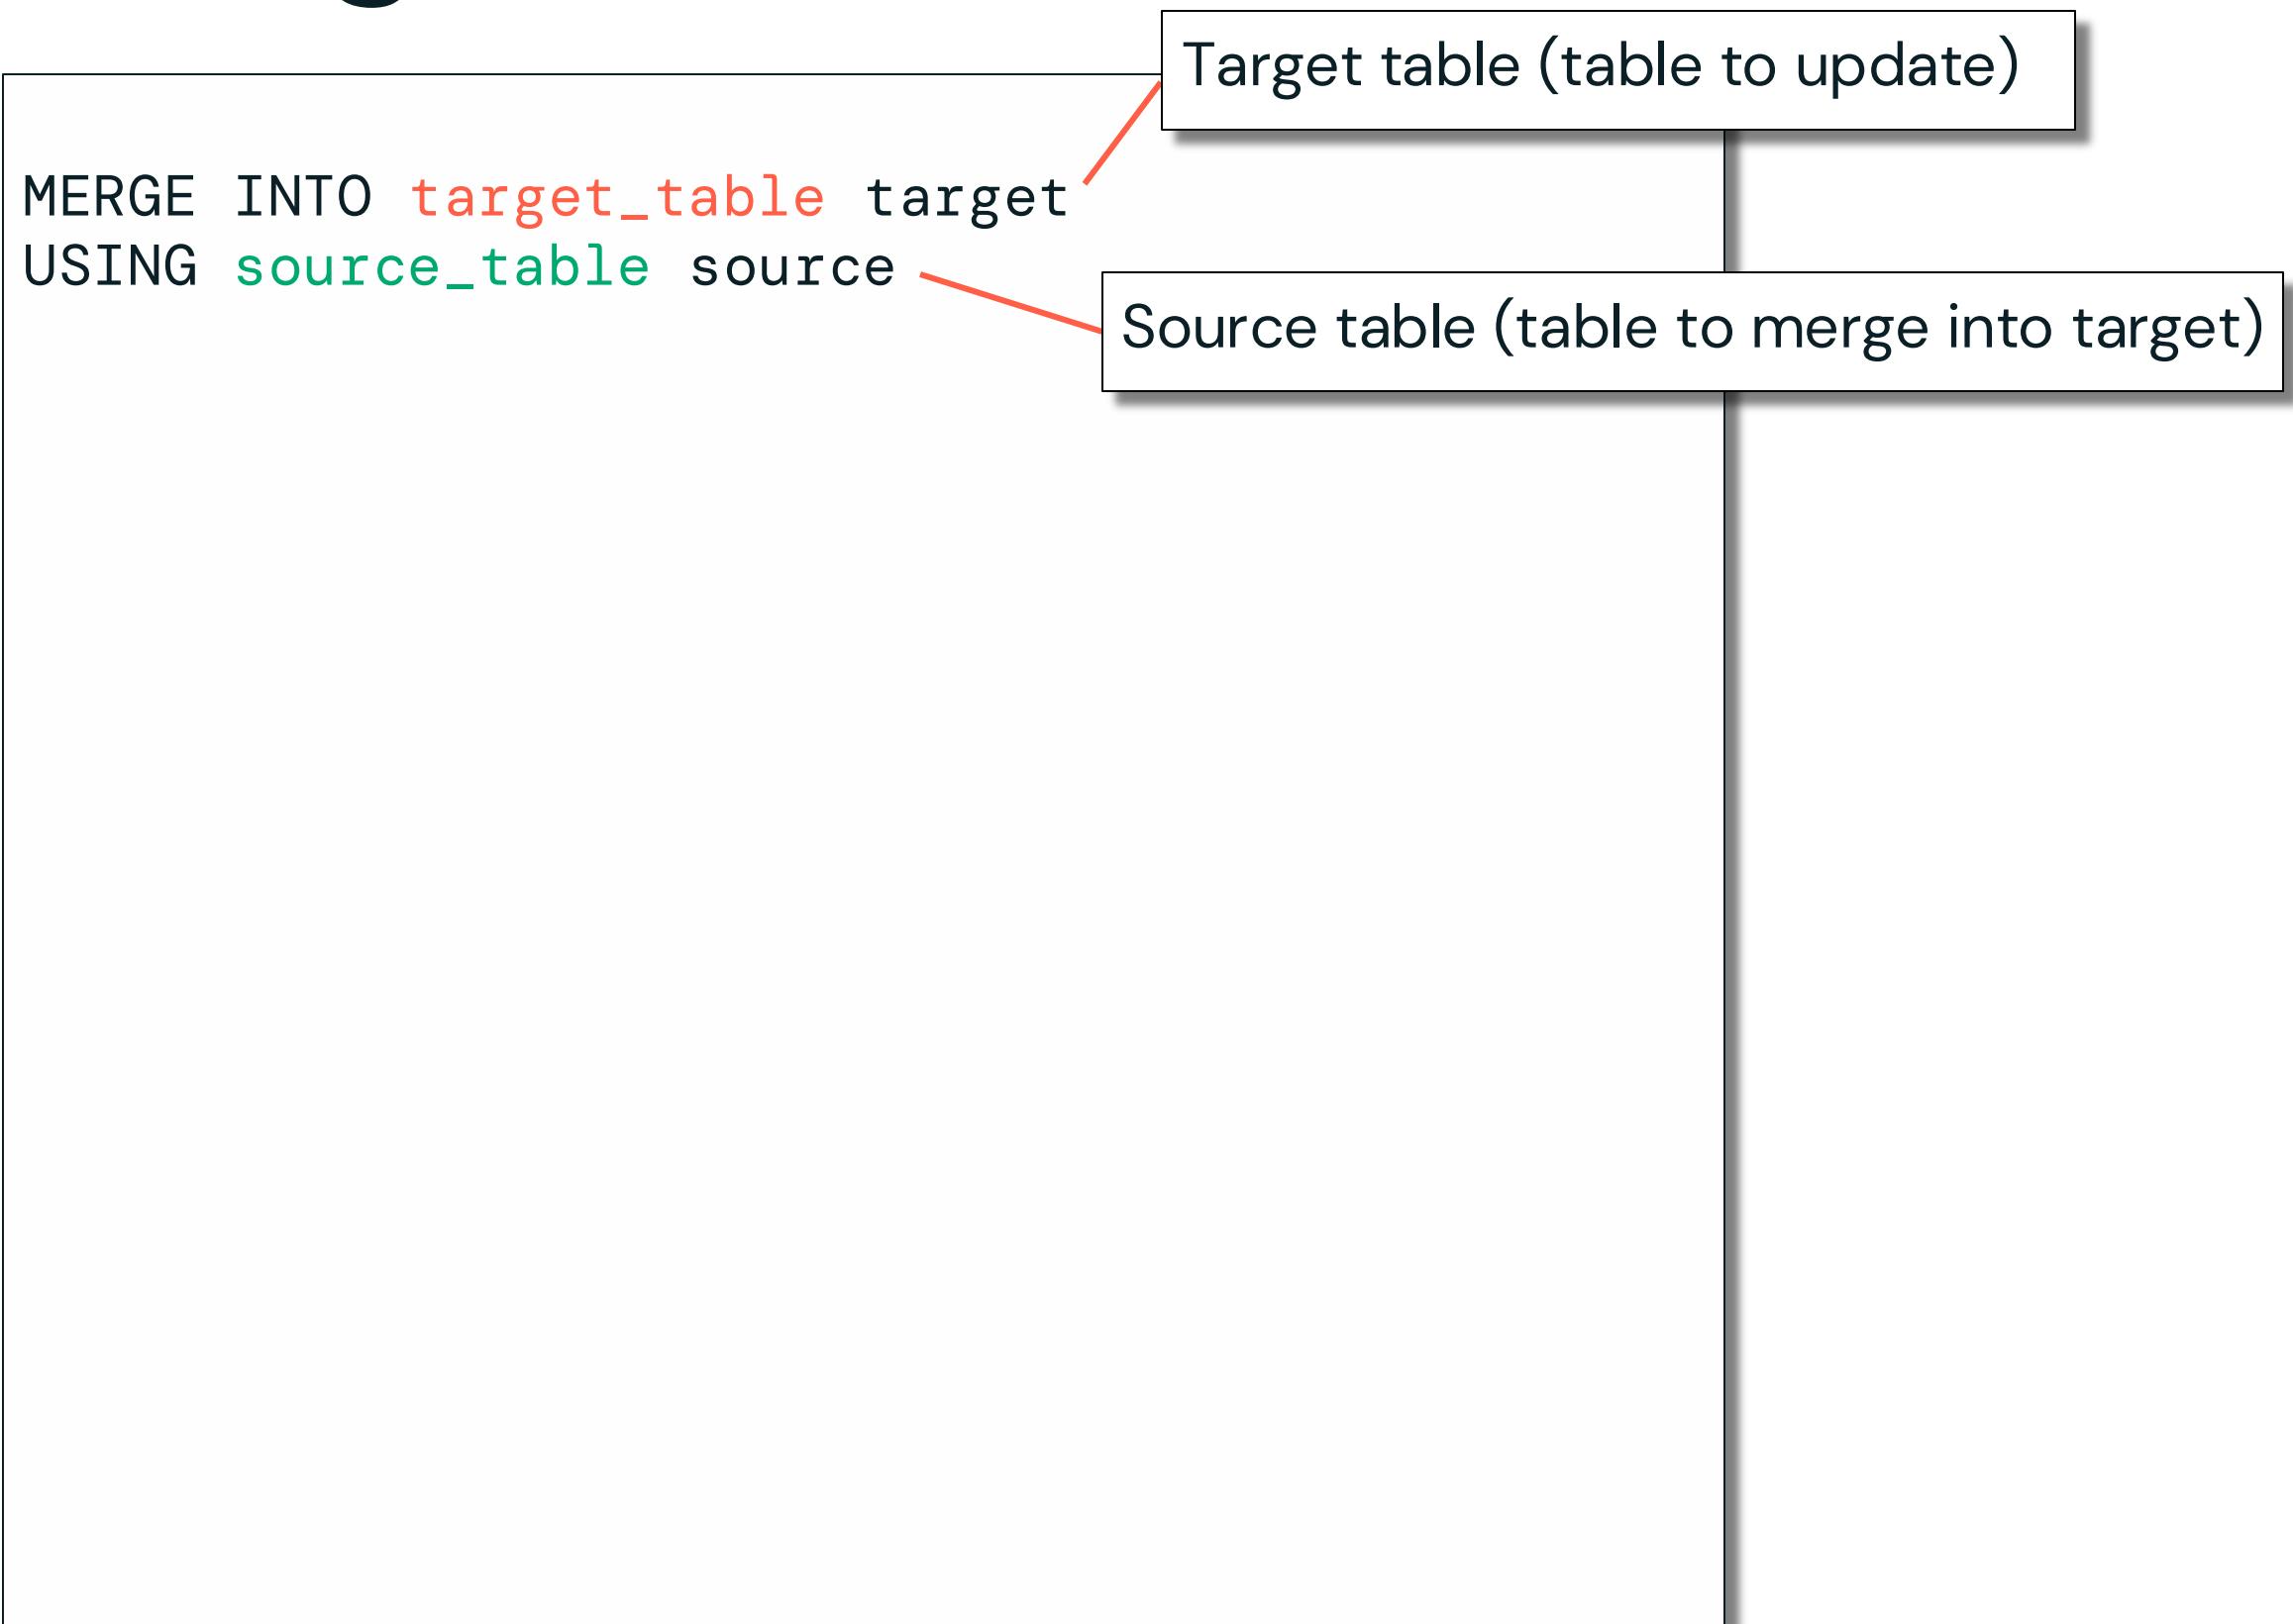
# Ingesting into Existing Delta Tables

MERGE INTO  
`source_table`

<code>id</code>	<code>users</code>	<code>email</code>	<code>status</code>
1	peter	peter@email	delete
2	zebi	zebi@email	update
3	samarth	samarth@email	new

Original  
`target_table`

<code>id</code>	<code>users</code>	<code>email</code>	<code>status</code>
1	peter	peter@email	current
2	zebi	zebi@email	current
4	matt	matt@email	current



# Ingesting into Existing Delta Tables

**source\_table**

<b>id</b>	<b>users</b>	<b>email</b>	<b>status</b>
1	peter	peter@email	delete
2	zebi	zebi@email	update
3	samarth	samarth@email	new

**Updated  
target\_table**

<b>id</b>	<b>users</b>	<b>email</b>	<b>status</b>
1	peter	peter@email	current
2	zebi	zebi@email	current
4	matt	matt@email	current

```
MERGE INTO target_table target
USING source_table source
ON target.id = source.id
```

Specify the condition for merging



# Ingesting into Existing Delta Tables

`source_table`

<code>id</code>	<code>users</code>	<code>email</code>	<code>status</code>
1	peter	peter@email	delete
2	zebi	zebi@email	update
3	samarth	samarth@email	new

Updated  
`target_table`

<code>id</code>	<code>users</code>	<code>email</code>	<code>status</code>
1	peter	peter@email	current
2	zebi	zebi@email	update
4	matt	matt@email	current

```
MERGE INTO target_table target
USING source_table source
ON target.id = source.id
WHEN MATCHED AND source.status = 'update' THEN
    UPDATE SET
        target.email = source.email,
        target.status = source.status
```

First WHEN clause  
for merge



# Ingesting into Existing Delta Tables

**source\_table**

<b>id</b>	<b>users</b>	<b>email</b>	<b>status</b>
1	peter	peter@email	delete
2	zebi	zebi@email	update
3	samarth	samarth@email	new

**Updated target\_table**



<b>id</b>	<b>users</b>	<b>email</b>	<b>status</b>
1	peter	peter@email	current
2	zebi	zebi@email	update
4	matt	matt@email	current

```
MERGE INTO target_table target
USING source_table source
ON target.id = source.id
WHEN MATCHED AND source.status = 'update' THEN
    UPDATE SET
        target.email = source.email,
        target.status = source.status
WHEN MATCHED AND source.status = 'delete' THEN
    DELETE
```

Second WHEN clause for merge



# Ingesting into Existing Delta Tables

**source\_table**

<b>id</b>	<b>users</b>	<b>email</b>	<b>status</b>
1	peter	peter@email	delete
2	zebi	zebi@email	update
3	samarth	samarth@email	new

**Updated target\_table**

<b>id</b>	<b>users</b>	<b>email</b>	<b>status</b>
2	zebi	zebi@email	update
4	matt	matt@email	current
3	samarth	samarth@email	new

```
MERGE INTO target_table target
USING source_table source
ON target.id = source.id
WHEN MATCHED AND source.status = 'update' THEN
    UPDATE SET
        target.email = source.email,
        target.status = source.status
WHEN MATCHED AND source.status = 'delete' THEN
    DELETE
WHEN NOT MATCHED THEN
    INSERT (id, first_name, email, sign_up_date,
            status)
    VALUES (source.id, source.first_name,
            source.email, source.sign_up_date,
            source.status);
```

WHEN NOT  
MATCHED clause  
for merge



# Ingesting into Existing Delta Tables

`source_table`

<code>id</code>	<code>users</code>	<code>email</code>	<code>status</code>
1	peter	peter@email	delete
2	zebi	zebi@email	update
3	samarth	samarth@email	new

`Fully Updated target_table`

<code>id</code>	<code>users</code>	<code>email</code>	<code>status</code>
2	zebi	zebi@email	current
4	matt	matt@email	current
3	samarth	samarth@email	new

```
MERGE INTO target_table target
USING source_table source
ON target.id = source.id

WHEN MATCHED AND source.status = 'update' THEN
    UPDATE SET
        target.email = source.email,
        target.status = source.status

WHEN MATCHED AND source.status = 'delete' THEN
    DELETE

WHEN NOT MATCHED THEN
    INSERT (id, first_name, email, sign_up_date,
status)
    VALUES (source.id, source.first_name,
source.email, source.sign_up_date,
source.status);
```



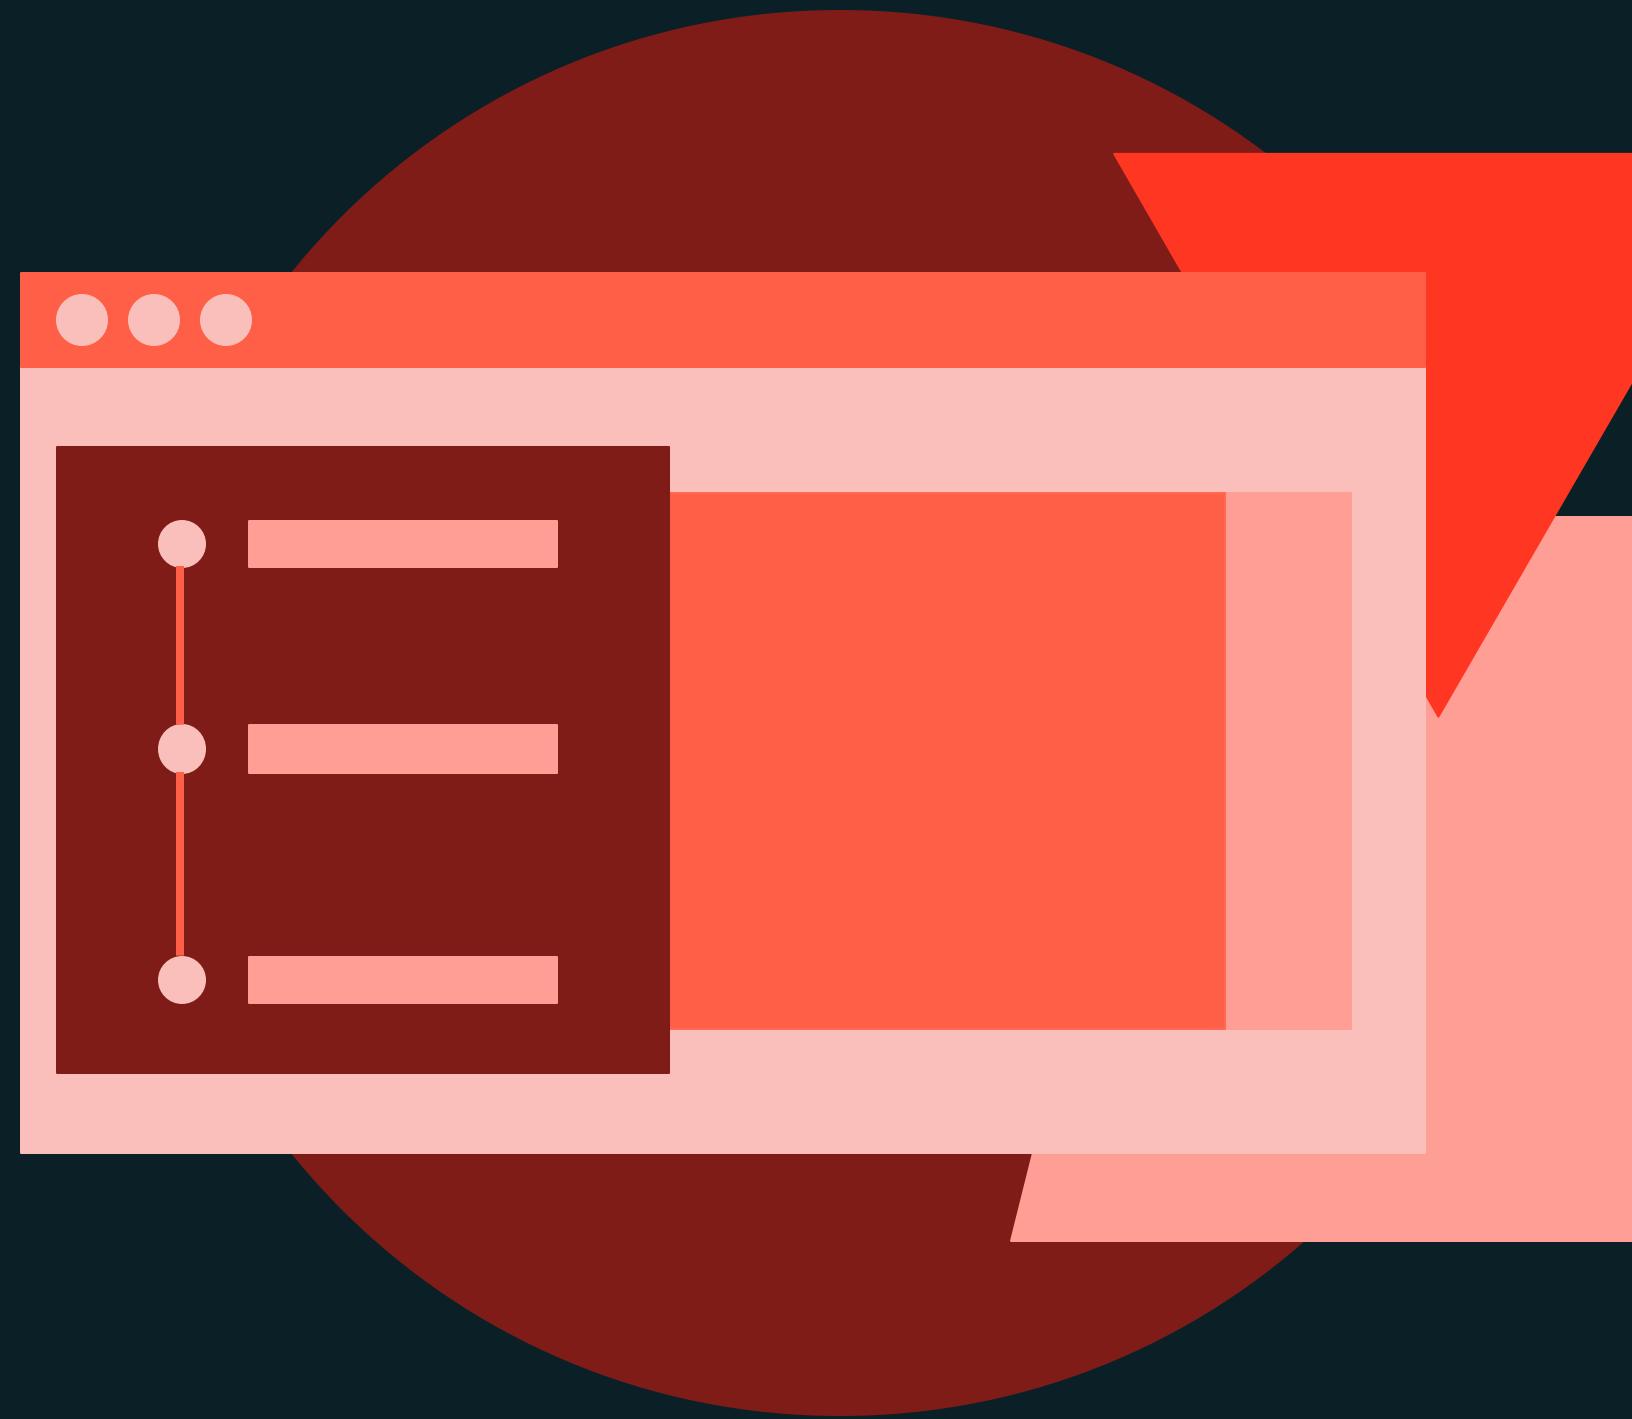


## Ingestion Alternatives

### DEMONSTRATION

# BONUS – Data Ingestion with MERGE INTO

If time permits, take a look at the fundamentals of MERGE INTO within Databricks for upserting data into another table.



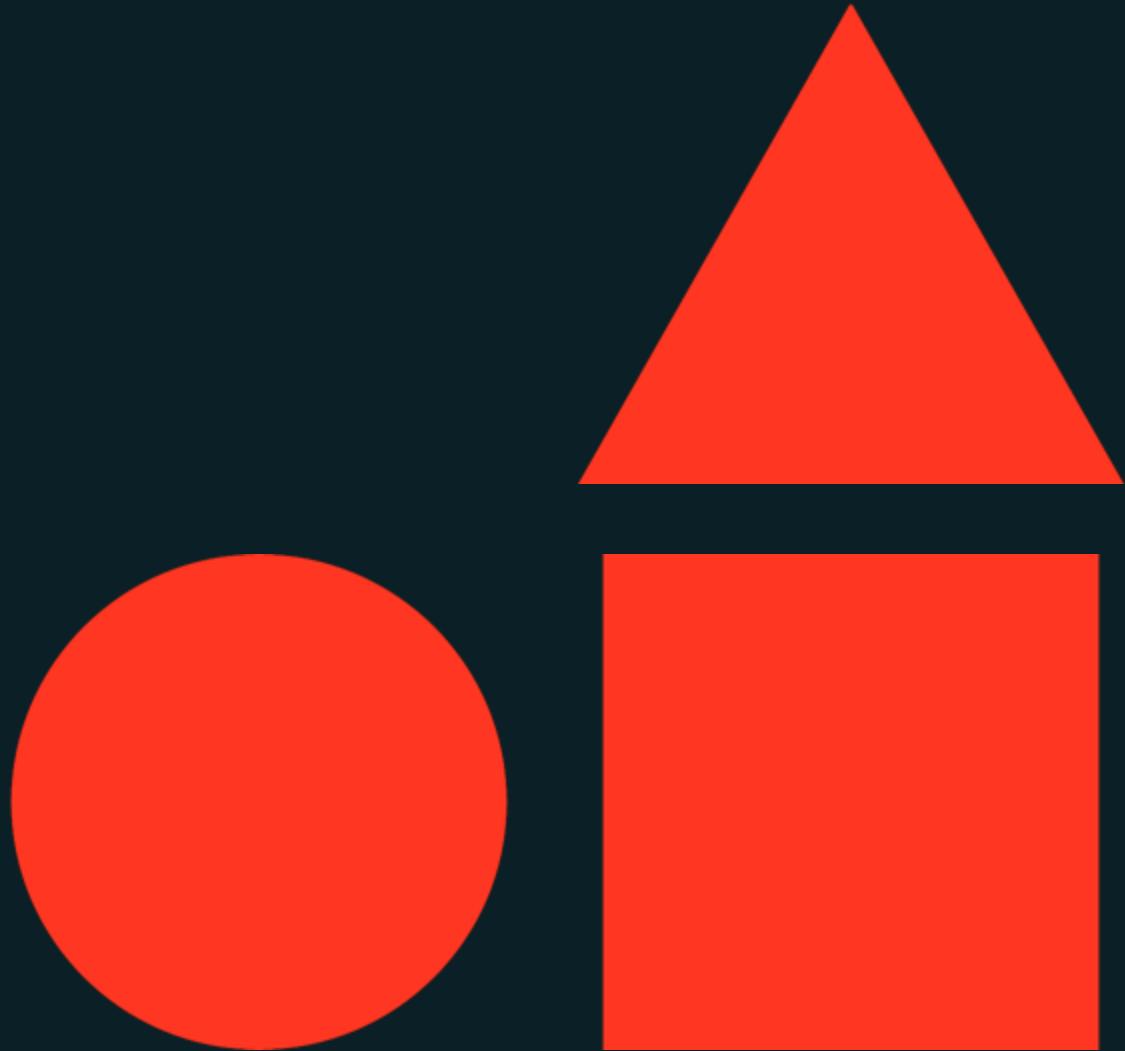
Notebook: 09 – BONUS – Data Ingestion with Merge Into



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



# Summary and Next Steps



## Data Ingestion with Lakeflow Connect



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Course Learning Objective Recap

- Describe Lakeflow Connect as a scalable and simplified solution for data ingestion into Databricks from a variety of sources.
- Review the benefits of Delta tables and the Medallion architecture.
- Demonstrate how to ingest data from cloud object storage into Delta tables using CREATE TABLE AS, COPY INTO, and Auto Loader, including capturing input file metadata in Bronze layer tables.
- Explain how rescued columns are used during ingestion to manage malformed records.
- Illustrate techniques for ingesting and flattening semi structured JSON data from cloud storage.
- Describe available options for ingesting data from enterprise systems using Lakeflow Connect Managed Connectors.
- Discuss alternative ingestion methods such as MERGE INTO, Delta Sharing and Databricks Marketplace.



# Next Steps

Additional resources for continuing the learning journey.

## Data Engineering with Databricks

- Continue your learning through [self-paced](#) or [instructor-led](#) offerings
- The courses offer hands-on instruction in:
  - Databricks Data Science & Engineering Workspace
  - Databricks SQL
  - Declarative Pipelines
  - Databricks Repos
  - Databricks Task Orchestration
  - Unity Catalog

## Data Engineer Associate Certification

- Validate your data and AI skills on Databricks by earning a Databricks credential
- [Exam information](#) and [exam guide](#)
- The exam covers:
  - Data Intelligence Platform
  - ELT With Spark SQL and Python
  - Incremental Data Processing
  - Production Pipelines
  - Data Governance





# databricks



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).